

THIS IS A SYNTHETIC WORLD: THREE ESSAYS ON CAUSAL INFERENCE, APPLIED
ECONOMETRICS, AND MACHINE LEARNING FOR POLICY ANALYSIS

by

Jared Amani Greathouse

Under the Direction of Jason Coupet, PhD

ABSTRACT

Whether we mean planning, justifying, and implementing a program for a non-profit or evaluating an already existing policy, causal inference to estimate the efficacy of policy is critical to decision making in government, business, and beyond. In general, policy analysts cannot conduct randomized trials, given their cost, ethical status, or infeasibility, and must therefore resort to carefully tuned econometric models to estimate counterfactuals, or a hypothetical scenario of how an outcome would have looked absent a given treatment or policy. Synthetic control methods (SCM) and difference-in-differences (DID) are among the most popular methods of estimating counterfactuals for policy analysts, and have seen much development in application and theory in recent years. In particular, some of their growth can be attributed to their extension via the integration of machine-learning methods. As such, my dissertation consists of three papers under a single theme: causal policy analysis in the high-dimensional setting, where *high-dimensional* refers to where we have many more control units relative to pre-treatment periods.

INDEX WORDS: Synthetic Control, Difference-in-Differences, Causal Inference, Policy Evaluation, Machine Learning

THIS IS A SYNTHETIC WORLD: THREE ESSAYS ON CAUSAL INFERENCE, APPLIED
ECONOMETRICS, AND MACHINE LEARNING FOR POLICY ANALYSIS

by

Jared Amani Greathouse

A Dissertation Submitted in Partial Fulfillment of the Requirements for the Degree of

Doctor of Philosophy

in the College of Arts and Sciences

Georgia State University

2025

Copyright by
Jared Amani Greathouse
2025

THIS IS A SYNTHETIC WORLD: THREE ESSAYS ON CAUSAL INFERENCE, APPLIED
ECONOMETRICS, AND MACHINE LEARNING FOR POLICY ANALYSIS

by

Jared Amani Greathouse

Committee Chair:

Jason Coupet

Committee:

Andrew Heiss

Lindsay Rose Bullinger

Jonathan Oxley

Electronic Version Approved:

Office of Graduate Studies

College of Arts and Sciences

Georgia State University

December 2025

DEDICATION

For my family, friends, and mentors.

ACKNOWLEDGMENTS

I am grateful to my advisor, Dr. Jason Coupet, and to my committee for their guidance, and to the colleagues, friends, and family whose support made this work possible.

Table of contents

ACKNOWLEDGMENTS	v
List of Tables	ix
List of Figures	x
1 Paradise Lost? Separating Policy from Pandemic in Hawaii's Border Closure with Single Proxy Synthetic Control	1
1.1 Introduction	2
1.2 Background	5
1.2.1 <i>Tourism in Hawaii</i>	5
1.2.2 <i>Literature Review</i>	7
1.3 Methodology	10
1.3.1 <i>The identification problem: a coincident common shock</i>	10
1.3.2 <i>Proximal inference: single-proxy and base specifications</i>	13
1.4 Data	16
1.4.1 <i>Treated outcomes</i>	16
1.4.2 <i>Insulated-sector proxies (the macro negative controls)</i>	17
1.4.3 <i>Open-destination air travel (the travel-demand negative controls)</i>	18
1.4.4 <i>Panel and estimation window</i>	18
1.5 Results	19
1.5.1 <i>The pre-treatment fit: relevance</i>	19
1.5.2 <i>The insulated-only estimate: a bound</i>	20
1.5.3 <i>Completing the bound: an ensemble of exposed-but-open destinations</i>	21
1.5.4 <i>Robustness and falsification</i>	26
1.5.5 <i>Who bore the cost: incidence across islands and hotel tiers</i>	29

1.6	Discussion	32
1.7	Conclusion	35
2	Clearing the Air: How India's 2020 Lockdown Impacted Air Quality . .	38
2.1	Introduction	38
2.2	Background and Previous Work	40
	<i>2.2.1 Policy Context</i>	40
	<i>2.2.2 Literature Review</i>	43
	<i>2.2.3 The Synthetic Historical Control Method</i>	45
	<i>2.2.4 Bandwidth Selection</i>	50
2.3	Data	50
2.4	Results	52
2.5	Robustness: Quarterly Aggregation	54
2.6	Discussion	56
2.7	Conclusion	58
3	Locking Away Prosperity? Evaluating the Labor Impacts of Vaccine Mandates	60
3.1	Introduction	60
3.2	Policy Background	63
	<i>3.2.1 NPIs Before Vaccines</i>	63
	<i>3.2.2 The First Vaccine Policies</i>	64
	<i>3.2.3 City-Level Vaccine Mandates</i>	66
	<i>3.2.4 The First Vaccine Mandates</i>	68
3.3	Dataset	69
3.4	Econometric Methodology	72
	<i>3.4.1 Synthetic Difference-in-Differences</i>	72
	<i>3.4.2 Partially Pooled Synthetic Control</i>	74
	<i>3.4.3 Triangulation Strategy</i>	76

3.5 Results	76
3.6 Discussion	78
3.7 Conclusion	81
REFERENCES	83

List of Tables

Table 3.1 Estimated effect of city-level vaccine mandates on year-over-year full-service restaurant employment growth, in percentage points. $J = 3$ treated CSAs (New York City, San Francisco, Los Angeles); SDID inference by $B = 500$ placebo replications, PPSCM by leave-one-treated-out jackknife over the three cities. 76

List of Figures

Figure 1.1 The negative-control premise: treated tourism outcomes collapse while the insulated-sector donors barely dip.	19
Figure 1.2 The open-destination ensemble: each donor, fit singly, pulls the estimate off the insulated-only bound.	23
Figure 1.3 Observed, SPSC and PI insulated counterfactuals, and SPSC + Florida, for five outcomes.	25
Figure 1.4 Specification robustness band: median SPSC ATT (point) and min–max range across eight specifications.	26
Figure 1.5 Subgroup ATTs by island and hotel class: insulated bound (circles), Florida-adjusted point (squares).	31
Figure 2.1 Constructing the donor matrix $\widetilde{\mathbf{Y}}_{\text{pre}}$ by sliding overlapping windows of length m across the smoothed pre-treatment series $\widetilde{\mathbf{y}}$. Each colored segment corresponds to a column in the donor matrix.	48
Figure 2.2 Synthetic Historical Control: all-India year-over-year PM2.5 growth, observed versus counterfactual. The dotted line marks the March 2020 national lockdown; the shaded region is the 90% Andrews–Genton conformal prediction band.	53
Figure 2.3 Synthetic Historical Control for the four megacities: observed versus counterfactual year-over-year PM2.5 growth, with 90% conformal prediction bands (shaded).	53
Figure 2.4 Robustness of the SHC estimates to temporal aggregation: the average treatment effect (percentage points of year-over-year PM2.5 growth) estimated at the monthly frequency versus the same estimator on quarterly data, for India and the four megacities.	55
Figure 3.1 Per-city synthetic-control fit: observed (black) and synthetic (dashed) year-over-year full-service restaurant employment growth for each treated city, with the mandate month marked (dotted). The synthetic tracks each city through the pre-mandate period; these panels are single-city SDID fits, and the partially pooled per-city effects are reported in the text.	78

Figure 3.2 Pooled synthetic difference-in-differences fit across the three treated cities (New York City, San Francisco, Los Angeles) under staggered adoption. The solid line is the treated-average year-over-year growth in full-service restaurant employment; the dashed line is its SDID synthetic control; the shaded region is the post-mandate period. The pooled SDID ATT is +5.8 and the PPSCM ATT is +3.7 percentage points, both positive and in the range of the estimates in Table~??.

CHAPTER 1

Paradise Lost? Separating Policy from Pandemic in Hawaii’s Border Closure with Single Proxy Synthetic Control ABSTRACT

Hawaii’s March 2020 mandatory traveler quarantine sealed the state’s borders at the same instant the COVID-19 pandemic collapsed global travel demand. For a single treated unit these two forces are perfectly collinear in time, so the conventional synthetic-control toolkit, which matches the treated series to a counterfactual that is itself battered by the pandemic, cannot separate the border-closure *policy* from the coincident *pandemic*. I address this with proximal / negative-control inference in two specifications, Single Proxy Synthetic Control (SPSC) and base Proximal Inference (PI), that *instrument* the latent pandemic factor using insulated employment sectors that carry the pandemic’s macro-recession shock yet are immune to the border policy. The two specifications agree, recovering a large, robust tourism-demand collapse (roughly 60–75 percentage points of year-over-year growth for the visitor-volume series); but because the insulated sectors span only the macro-recession factor and not the travel-demand-collapse factor, they under-span the tourism factor and identify these effects only weakly; the estimates from insulated donors alone are therefore best read as *upper bounds in magnitude* on the policy effect and are reported as a robustness band. Recovering an actual point requires spanning the missing travel-demand factor from outside the treated economy: adding external travel-demand donors, inbound air travel to Florida and other exposed-but-open leisure destinations that were not locked down, pulls the Visitor Days estimate off its upper bound, and a convergent ensemble of such donors brackets it

in a policy-centered band, taking the step from a bound to a band. The paper shows how negative-control inference yields credible causal estimates when a policy and a global shock arrive together.

1.1 Introduction

In 2019, tourism moved \$17.75 billion in visitor spending through Hawaii’s economy, underwrote \$2.07 billion in state tax revenue, and sustained about 216,000 jobs, close to a third of the state’s workforce ([Hawai‘i Tourism Authority, 2019](#)). In March 2020, Hawaii deliberately switched much of that off. Among all U.S. states it enacted perhaps the most extreme and comprehensive economic shutdown of the COVID-19 pandemic ([Yan et al., 2022](#)). Where most regions contracted through voluntary distancing, fear, or uneven mandates, Hawaii chose one of the few deliberate, coordinated state-level shutdowns in the country, a policy aimed not merely at slowing the virus but at suppressing travel and commerce outright. On March 26, 2020, it imposed a mandatory 14-day quarantine on all arrivals, effectively sealing its borders. The outcome was among the starkest in the country: Hawaii went on to record the largest job losses of any U.S. state even as it held among the lowest COVID-19 case rates ([Bond-Smith and Fuleky, 2022](#)). This paper asks what that decision cost.

That question is not merely historical. A governor weighing whether to seal a jurisdiction, in the next pandemic or in a climate or geopolitical shock that likewise pits containment against commerce, faces a concrete tradeoff with little credible evidence on its economic side ([Newman and Noy, 2023](#)): the health effects of COVID-19 interventions have been studied

extensively, their economic toll, especially under maximal containment, far less so. Hawaii is the clearest U.S. instance of that decision in its starkest form, a whole economy switched off to keep a virus out, and a credible price for it is what a policymaker needs to weigh the closure against the lives it aimed to protect.

Estimating that cost, however, runs into an identification problem that is easy to overlook and fatal if ignored. Hawaii closed its border in the same month the pandemic collapsed global travel demand, and neither force varied before March 2020, so for a single treated unit the *policy* and the *pandemic* are perfectly collinear in time: every post-treatment month is at once a quarantine month and a pandemic month, and no function of the observed data separates a response to one from a response to the other. This disables the two designs one would reach for first. Difference-in-differences would net out the pandemic by subtracting untreated states, but that assumes the shock lands equally everywhere ([de Chaisemartin and D’Haultfoeulle, 2022](#); [Roth et al., 2023](#)), and Hawaii, the most tourism-dependent state, was struck far harder than any control. Synthetic control weakens that to a factor model and reweights donors to Hawaii’s pre-2020 path, but the travel-demand collapse is dormant until March 2020 and then switches on with the policy, leaving no pre-period variation to fit: even with tourism series for all forty-nine other states, the synthetic Hawaii is a no-2020 world, a counterfactual that strips out the pandemic along with the quarantine, and differencing charges the entire post-March collapse to the border closure. Section~[1.3.1](#) makes this formal; the failure is a property of the coincident shock itself.

This paper takes that identification problem seriously and resolves it with the right tool. The

pandemic confounder is a time-varying latent factor that moves the outcome and coincides with treatment; it cannot be conditioned away because it is unobserved, and it cannot be matched away because it moves any donor too. The *proximal* answer (Miao et al., 2018; Tchetgen Tchetgen et al., 2024) is to instrument it: find observable negative controls that are driven by the same pandemic factor but carry no causal path from the treatment, and use them to subtract the pandemic out of the counterfactual. Here those negative controls are insulated employment sectors, hit by the pandemic recession but not downstream of whether tourists can reach Hawaii. I fit two proximal specifications, Single Proxy Synthetic Control (SPSC) (Park and Tchetgen Tchetgen, 2025) and base Proximal Inference (PI), and report both, so the estimate does not hinge on one estimator’s machinery.

Empirically, the two specifications agree: fit on the insulated proxies, SPSC and PI recover the same large tourism-demand collapse, so the result rests on the identification; the two estimators only corroborate it. This delivers the first credible, if bounded, price for a maximal pandemic border closure, a collapse on the order of 60 to 75 percentage points of year-over-year growth for the visitor-volume series. What they identify is a bound, and I treat it as one. The insulated sectors span the pandemic’s general-recession factor but not its travel-demand-collapse factor, so the tourism estimates absorb both the policy and the demand loss and bound the border-closure effect from above. I report them as a specification band and check each against a relevance diagnostic. Converting that bound into a point requires a proxy that carries the missing travel-demand factor, and I supply one from outside the treated economy: inbound air travel to open, no-lockdown destinations, Florida foremost,

which pulls the Visitor Days estimate off its upper bound to a policy-centered point.

The remainder of the paper proceeds as follows. Section 1.2 reviews prior research on early pandemic interventions and the methodological challenge of studying one-off, jurisdiction-wide policies. Section 1.3 states the coincident-shock identification problem and the two proximal estimators. Section 1.4 describes the data. Section 1.5 reports the proximal estimates and their identification diagnostics, the robustness band, and a set of falsification checks. Section 1.6 discusses implications for public policy and crisis response.

1.2 Background

1.2.1 *Tourism in Hawaii*

Hawaii’s economy transitioned from a plantation-dominated model in the late 19th and early 20th centuries to tourism as the primary driver after statehood. The sugar and pineapple industries, controlled by a few powerful firms known as the “Big Five,” peaked in the 1960s but declined due to global competition, labor costs, and diversification efforts. Statehood in 1959, combined with the introduction of jet air travel, spurred a tourism boom, with visitor numbers surging from 171,000 in 1958 to over 6.7 million by 1990 (Croix, 2021). This shift modernized the economy, increased population growth, and generated substantial revenue, but it also created dependency on external factors like air travel and international markets.

The scale of that dependence is worth making concrete. On any given day in 2019, Hawaii hosted about 249,000 visitors spending \$48.6 million, with O’ahu (\$22.4 million daily) and Maui (\$14.0 million) bearing the brunt (Hawai‘i Tourism Authority, 2019); leisure and

hospitality employment, a proxy for tourism-related jobs, held around 190,000–200,000 pre-2020. This reliance amplifies the state’s exposure to shocks even as it funds public services and infrastructure.

Hawaii implemented some of the earliest and most aggressive COVID-19 containment policies in the United States. On March 4, 2020, Governor David Ige issued an emergency proclamation officially recognizing the virus as a statewide threat ([Maui Now Staff, 2020](#)). Within weeks, the state began taking unprecedented measures alongside states like California ([Friedson et al., 2021](#)). On March 17, Governor Ige publicly urged all visitors to suspend their trips to Hawaii for 30 days, anticipating substantial economic damage to the state’s tourism-dependent economy. *"We do know there will be significant impact,"* he stated, though no concrete estimates were available at the time ([Nakaso, 2020](#)). A sweeping stay-at-home order followed on March 25, banning non-essential activities, closing non-critical businesses, and effectively shutting down day-to-day life and travel across the islands ([HNN Staff, 2020](#)). These actions were implemented alongside a mandatory 14-day quarantine on all arrivals, passed the same week, which further intensified the shutdown and isolated the state from both domestic and international travel. Despite the unprecedented scope of these interventions, no formal quantitative analysis has yet estimated the short-run economic effects of Hawaii’s approach. This paper provides such an assessment.

1.2.2 Literature Review

A large early literature makes the public-health case for acting fast, and the lesson is consistent: early, aggressive containment slows transmission and buys hospital capacity. [Friedson et al. \(2021\)](#) find that California’s statewide shelter-in-place order, the first in the country, averted 160 to 195 COVID-19 cases per 100,000 and as many as 1,566 deaths within a month, at a back-of-the-envelope cost of roughly 650 to 700 jobs per life saved. [Yang \(2021\)](#) estimates that Wuhan’s 76-day cordon cut inbound mobility by about 60% and halved the growth of new cases within days, and [Cerqueti et al. \(2021\)](#), using an augmented synthetic control, reach the same sooner-the-better verdict for Italy’s national lockdown. Timing is what pays: [Bayat et al. \(2020\)](#) estimate that locking down ten days earlier would have cut New York deaths by as much as 80%, and [Sobiech Pellegrini \(2022\)](#) and [Li and Shankar \(2023\)](#) document the mirror image, sharp case rebounds after premature reopenings in Santa Catarina and Texas. These designs price the health payoff of containment, and the California study even prices the jobs traded for it, but they do so for interventions milder than a full border seal and lean on comparison regions that a coincident global shock would contaminate.

The economic side of the ledger is thinner, and where it exists it tends to find the damage real but transitory, and mostly not the policy’s own doing. [Ke and Hsiao \(2022\)](#), with a panel-data counterfactual, estimate that Hubei’s 76-day lockdown cut first-quarter 2020 GDP by about 37%, though the economy rebounded within months of reopening, and [Gong et al. \(2024\)](#) put the output cost of China’s stricter 2022 zero-COVID phase at 3.9% of GDP. The recurring result is that the observed economic hit is dominated by the common pandemic shock, with

the local policy contributing far less. [Lee and Yang \(2022\)](#) find the pandemic erased 1.1 million Korean jobs by April 2020, but only 9% of that national loss traces to regional COVID intensity, the remaining nine-tenths to a nationwide common effect; [Gibson and Sun \(2023\)](#) likewise attribute only 32% of the U.S. spike in new jobless claims to stay-at-home orders, the rest to the general spread of the virus and collapsing consumer confidence; and [Osman et al. \(2022\)](#) can identify state-level effects at all only by leaning on the cross-state policy variation that a uniform national shock erases. That decomposition is the crux of the present paper: when the observed collapse is mostly the coincident shock, a credible policy estimate must separate the two, because differencing against a control the same shock has already moved cannot. Set against these costs, containment also carried some economic and welfare upside, from the pollution-mortality co-benefits of Wuhan’s lockdown ([Cole et al., 2020](#)) to the GDP losses that Europe’s vaccine-certificate mandates averted ([Oliu-Barton et al., 2022](#)). For all its breadth, this literature leaves the decision at its extreme end essentially unpriced. The question a policymaker faces is not whether containment works but what its most aggressive form costs, and that figure is what the record lacks. The gap has teeth because the politics cut against early action: as [Cerqueti et al. \(2021\)](#) observe, a containment that succeeds looks excessive in hindsight, so the officials who act first are precisely those most in need of a credible after-the-fact accounting to defend the call. Hawaii is the clearest U.S. instance of that most-aggressive option, and this paper supplies the missing price for a “zero-COVID” style border ([Yuan, 2022](#)).

While much has been learned about the effects of pandemic interventions, measuring their

economic impact remains hard, and Hawaii pushes that difficulty to an extreme. Most evaluations rely on difference-in-differences or synthetic control, and both lean on cross-sectional variation in policy, some counties or states acting while others do not, or acting later, to supply the untreated or not-yet-treated comparison the design needs (e.g., [Yang, 2021](#); [Gong et al., 2024](#); [Bilgel, 2022](#)). [Callaway and Li \(2023\)](#) show that even this variation is treacherous during an epidemic, because the timing-based identifying assumptions are hard to reconcile with the dynamics of contagion. And it vanishes entirely for a single treated unit hit by a global shock: every other U.S. state and comparable island economy faced its own pandemic disruption, so there is no untreated Hawaii to borrow across space. [Ke and Hsiao \(2023\)](#) confront a version of the same coincident-shock problem, developing a panel method for data subject to multiple treatment effects that disentangles a global pandemic’s impact from a specific, coincident one, and [Li et al. \(2023\)](#) apply the same panel-plus-time-series logic to separate a Shanghai housing policy from the pandemic. The only prior study of Hawaii’s pandemic economy, [Bond-Smith and Fuleky \(2022\)](#), documents the depth of the decline but benchmarks it against a pre-pandemic forecast, leaving the policy counterfactual unaddressed.

One response to the missing-donor problem is to borrow strength across *time* rather than space: the Synthetic Historical Control (SHC) method of [Chen et al. \(2024\)](#) reconstructs the treated unit’s counterfactual from its own pre-treatment trajectory. But borrowing across time does not, by itself, solve the deeper problem here. Whether the counterfactual is built from other states (classical SCM) or from Hawaii’s own history (SHC), it is a world with

neither the quarantine *nor* the pandemic; extrapolating it into 2020 and differencing against the observed series attributes the entire post-March collapse, pandemic and policy together, to the treatment. The issue is the *coincidence* of two shocks in a single unit, whatever the source of the donor information, so I do not use own-history extrapolation as the estimator and turn instead to the proximal / negative-control approach.

That is precisely the structure the *proximal* / *negative-control* literature was built to handle. Negative controls, variables driven by an unmeasured confounder but shielded from the treatment, have long been used to detect confounding (Lipsitch et al., 2010); proximal causal inference turns them into an identification strategy, using proxies of a latent confounder to recover causal effects that ordinary adjustment cannot (Miao et al., 2018; Cui et al., 2024; Tchetgen Tchetgen et al., 2024; Shi et al., 2020). Park and Tchetgen Tchetgen (2025) adapt this machinery to the synthetic-control setting as Single Proxy Synthetic Control, and Liu et al. (2024b) develop a related surrogate-based proximal estimator. I apply proximal inference, SPSC and base PI, to Hawaii because it is the one approach in the synthetic-control family that concedes the pandemic confounder is present and unobservable and instruments it, rather than matching or conditioning it away.

1.3 Methodology

1.3.1 *The identification problem: a coincident common shock*

We know how Hawaii’s economy looked before the border closure, and what it did once it mandated the quarantine; but estimating the policy’s cost requires a counterfactual, an

estimate of how Hawaii's tourism sector would have evolved absent the closure. To build one, we adopt the comparative case study framework, comparing a single treated unit against a set of control units, or donors, that did not enact the policy. Index the N units by j , with $j = 1$ the treated unit (Hawaii) and the other N_0 the control units $\mathcal{N}_0$, and index the T months by t , collected in $\mathcal{T} = \{1, \dots, T\}$ and split at the final pre-quarantine month T_0 into $\mathcal{T}_1 = \{t \leq T_0\}$ and $\mathcal{T}_2 = \{t > T_0\}$; y_{jt} is unit j 's outcome in month t . Write y_{1t}^1 and y_{1t}^0 for Hawaii's month- t outcome with and without the closure. Only one is ever seen: with $d_t = \mathbf{1}\{t > T_0\}$ the treatment indicator, $y_{1t} = d_t y_{1t}^1 + (1 - d_t) y_{1t}^0$.

Synthetic control builds the counterfactual as a weighted average of the donors. Stacking the treated unit's path into $\mathbf{y}_1 \in \mathbb{R}^T$ and the donors' paths columnwise into $\mathbf{Y}_0 \in \mathbb{R}^{T \times N_0}$, it returns the weight vector $\mathbf{w}$ that minimizes the pre-treatment mean squared error between the treated unit and the weighted donors,

$$\mathbf{w}^* = \underset{\mathbf{w} \in \Delta^{N_0}}{\operatorname{argmin}} \left\| \mathbf{y}_1 - \mathbf{Y}_0 \mathbf{w} \right\|_2^2. \quad (1.1)$$

The weights are restricted to the probability simplex $\Delta^{N_0} = \{\mathbf{w} \in \mathbb{R}_{\geq 0}^{N_0} : \sum_{j \in \mathcal{N}_0} w_j = 1\}$. Carried past T_0 , the fitted values $\mathbf{Y}_0 \mathbf{w}^*$ of (1.1) are the counterfactual for the treated unit, and the quarantine's cost is the average treatment effect on the treated over the post-intervention period,

$$\text{ATT} := \frac{1}{n} \sum_{t \in \mathcal{T}_2} (y_{1t}^1 - y_{1t}^0), \quad n = |\mathcal{T}_2|. \quad (1.2)$$

Whether that estimate is credible depends on the data-generating process. The usual

justification for synthetic control is the interactive fixed-effects model (Bai, 2009; Abadie et al., 2010),

$$y_{1t}^0 = \boldsymbol{\mu}_1^\top \boldsymbol{\lambda}_t + e_{1t}, \quad y_{jt} = \boldsymbol{\mu}_j^\top \boldsymbol{\lambda}_t + e_{jt}, \quad j \in \mathcal{N}_0, \quad (1.3)$$

where $\boldsymbol{\lambda}_t \in \mathbb{R}^r$ collects time-varying factors common to all units, $\boldsymbol{\mu}_j$ collects unit-specific, time-invariant loadings, and $\mathbb{E}[e_{jt} \mid \boldsymbol{\lambda}_t] = 0$ is idiosyncratic error. Under this model, synthetic control is the standard tool, and the theory is precise about when it works. Differencing the treated series against any weighted counterfactual leaves, after some algebra,

$$y_{1t} - \sum_{j \in \mathcal{N}_0} w_j y_{jt} = \underbrace{(y_{1t}^1 - y_{1t}^0)}_{\text{policy effect}} + \underbrace{\left(\boldsymbol{\mu}_1 - \sum_{j \in \mathcal{N}_0} w_j \boldsymbol{\mu}_j \right)^\top}_{\text{factor-mismatch bias}} \boldsymbol{\lambda}_t + (e_{1t} - \sum_{j \in \mathcal{N}_0} w_j e_{jt}), \quad (1.4)$$

the policy effect plus a factor-mismatch bias, the residual loading gap $\boldsymbol{\mu}_1 - \sum_j w_j \boldsymbol{\mu}_j$ projected onto the common factors, plus idiosyncratic noise (Abadie et al., 2010; Ferman and Pinto, 2021; Zhang et al., 2022; Hollingsworth and Wing, 2022). The guarantee behind synthetic control is that a long pre-period with a close pre-treatment fit drives that loading gap to zero, forcing the weights to reproduce Hawaii’s loadings,

$$\boldsymbol{\mu}_1 = \sum_{j \in \mathcal{N}_0} w_j \boldsymbol{\mu}_j. \quad (1.5)$$

For that guarantee to hold however, the same factors and loadings must govern the pre- and post-treatment periods, so that a fit achieved before treatment still holds after it. In the Hawaii case the condition fails. In March 2020 a new coordinate of $\boldsymbol{\lambda}_t$ switches on,

the effect of the pandemic, absent from the data before the quarantine and present after, and so perfectly collinear with the treatment over the post-period. The weights are fit on the pre-period, where this factor is silent, so they cannot be chosen to reproduce Hawaii's loading on it: the factor-mismatch term in (1.4) no longer vanishes, and because the factor is unobserved and absent in sample, the bias cannot even be estimated. Practically, this means that even if we had noiseless, observed tourism data across the remaining United States, a good pre-treatment fit built from the donor states no longer certifies the counterfactual, and no reweighting of pre-period donors can subtract a shock the pre-period never saw.

1.3.2 Proximal inference: single-proxy and base specifications

Fortunately, one solution to this collinearity is the relatively new strand of synthetic control that employs proximal inference (Miao et al., 2018; Tchetgen Tchetgen et al., 2024; Shi et al., 2026). The idea is to find a negative control, or proxy: a variable we can use to net the pandemic out of the counterfactual and leave the policy's effect behind. For a proxy to do that, it must be both relevant and exclusionary. A relevant proxy has a strong relationship with the outcome of interest over the pre-policy period, so that it carries a genuine reading of the confounding pandemic factor. An exclusionary proxy is one the policy does not affect, so that reading stays uncontaminated by the treatment we are trying to measure. A variable with both properties moves with the pandemic but not with the border closure, which is exactly what lets the pandemic be subtracted from the counterfactual while the policy effect stays in, pulling apart two forces the collinearity had fused into one.

To state these conditions precisely, collect the candidate proxies at month t in a vector $\mathbf{p}_t$. Relevance asks that some combination of the proxies track the treated unit's untreated outcome,

$$\text{Cov}(h^*(\mathbf{p}_t), y_{1t}^0) \neq 0 \quad \text{for some } h^* : \mathbb{R}^{N_0} \rightarrow \mathbb{R}, \quad t \in \mathcal{T}_1, \quad (1.6)$$

the proxy analogue of a strong first stage, whose empirical footprint is the pre-treatment fit; a proxy that does not move with y_{1t}^0 carries no reading of the factor and cannot help. Relevance and exclusion together secure a clean reading of the pandemic, but a reading is not yet a counterfactual. What connects the two is a bridge function: a function h^* of the proxies that, in expectation, equals the untreated outcome,

$$y_{1t}^0 = \mathbb{E}[h^*(\mathbf{p}_t) \mid y_{1t}^0] \quad \text{almost surely}, \quad t \in \mathcal{T}, \quad (1.7)$$

so that passing the proxies through h^* returns what Hawaii's tourism would have done absent the closure. In short, a good proxy is one the pandemic drives strongly (relevance) and the policy leaves alone (exclusion), tied to the treated counterfactual through a bridge function.

Two estimators turn the proxies into a counterfactual, differing only in how they pin down the bridge weights $\mathbf{w}$ in $h^*(\mathbf{p}_t) = \mathbf{p}_t^\top \mathbf{w}$. Single Proxy Synthetic Control ([Park and Tchetgen Tchetgen, 2025](#)) uses the treated unit's own pre-treatment history as the instrument and fits the bridge by ridge-regularized GMM on the de-trended pre-period series,

$$(\hat{\boldsymbol{\eta}}, \hat{\mathbf{w}}_\varrho) = \underset{\boldsymbol{\eta}, \mathbf{w}}{\text{argmin}} \left\{ \overline{\boldsymbol{\Psi}}(\boldsymbol{\eta}, \mathbf{w})^\top \hat{\boldsymbol{\Omega}} \overline{\boldsymbol{\Psi}}(\boldsymbol{\eta}, \mathbf{w}) + \varrho \|\mathbf{w}\|_2^2 \right\}, \quad (1.8)$$

with $\overline{\Psi}$ the averaged moment conditions, $\widehat{\Omega}$ a weight matrix, and the ridge penalty ϱ (cross-validated) stabilizing the weights when the single instrument leaves them under-identified. Base Proximal Inference (PI) is the donors-only estimator of [Shi et al. \(2026\)](#), which in this implementation reduces to a least-squares fit of Hawaii’s pre-period series on the donors with a GMM/HAC standard error. Either way the fitted bridge gives the counterfactual $\widehat{y}_{1t}^0 = \mathbf{p}_t^\top \widehat{\mathbf{w}}$ and the ATT (1.2) as the average post-period gap $y_{1t} - \widehat{y}_{1t}^0$; the estimating equations, penalty selection, standard errors, and conformal band are given in [Park and Tchetgen Tchetgen \(2025\)](#) and implemented in `mlsynth`. Because the treated unit supplies a single instrument against the donors, the analytic standard errors are fragile here, so I lean on the specification band of Section~1.5 for uncertainty; sharing the identification but not the algorithm, the two estimators’ agreement there is a property of the proxy design.

The proxies fall into two groups, both detailed in Section~1.4. The first is the $N_0 = 5$ insulated employment sectors: they fell in the pandemic recession, so they carry its macro-recession factor (relevance), but they do not depend on whether visitors can reach Hawaii, so the border closure leaves them untouched (exclusion). The second is a set of external travel-demand donors, inbound air travel to exposed-but-open leisure destinations (Florida foremost), which carry the pandemic’s travel-demand-collapse factor while lying outside Hawaii’s policy. Both enter the same proxy matrix $\mathbf{P} := \mathbf{Y}_0$, with month- t row $\mathbf{p}_t$; a travel donor widens $\mathbf{P}$ to span a second factor; it does not add a control of a different kind. The split governs what the estimator can recover: the insulated sectors span only the macro-recession factor, so on them alone every tourism estimate is an upper bound in magnitude on the policy effect, absorbing

the travel-demand collapse, and adding a travel-demand donor is what turns that bound into a point (Section~1.5.3). I summarize each fit by its pre-treatment root-mean-squared error $\text{RMSE}_{\text{pre}} = \{T_0^{-1} \sum_{t \in \mathcal{T}_1} (y_{1t} - \hat{y}_{1t}^0)^2\}^{1/2}$, small relative to the effect when the reconstruction is informative.

1.4 Data

The analysis joins three public sources on calendar month and works throughout in year-over-year (YoY) percentage growth. The YoY transformation removes the strong seasonality of Hawaii tourism and employment, puts series measured in different units on a common scale, and lets each treatment effect be read directly as a percentage-point change in growth. Every series is stored as a raw monthly level and differenced to YoY growth at estimation time, so the panel retains levels for any later recovery analysis and ships no pre-differenced series. Hawaii’s Department of Business, Economic Development & Tourism (DBEDT) supplies the treated tourism series and the insulated employment sectors; the U.S. Bureau of Transportation Statistics (BTS) supplies the external air-travel proxies. The estimation window is described at the end of the section.

1.4.1 *Treated outcomes*

The treated series are Hawaii’s five tourism outcomes from DBEDT: the visitor-volume and price measures `Visitor Arrivals`, `Visitor Days`, `Occupancy`, `Mean Daily Rate`, and `Revenue per Available Room`. Each is stored as a raw monthly level (visitor counts, occupancy percentages, and dollar rates) and differenced to year-over-year growth at estimation

time. Every one is downstream of the border closure: each moves when visitors' ability to arrive changes, so each can only be a treated outcome, never a control.

1.4.2 Insulated-sector proxies (the macro negative controls)

The negative-control donors are five insulated employment sectors drawn from the DBEDT seasonally-adjusted Current Employment Statistics, the monthly job count by industry ([Hawai'i Department of Business, Economic Development and Tourism, 2025](#)), expressed in YoY growth: `NatRes_Constr_Emp` (mining, logging and construction), `Wholesale_Emp` (wholesale trade), `Financial_Emp` (financial activities), `HealthCare_Emp` (health care and social assistance), and `Government_Emp`. These industries were hit by the pandemic's macro-recession shock, falling roughly 7–12% in the 2020 trough, but do not depend on whether tourists can physically enter Hawaii, so no causal path runs from the quarantine to them. They are the empirical negative controls of Section~1.3.2.¹ Figure 1.1 contrasts the two groups: the treated tourism series collapse by 50–99% at the border closure while the insulated donors dip only modestly, the visual signature of a factor structure in which the donors carry the macro shock and the tourism series carry the macro shock plus the travel-demand collapse plus the policy.

¹The five are the most insulated supersectors in the CES taxonomy ([Hawai'i Department of Business, Economic Development and Tourism, 2025](#)). I screened the full taxonomy for the negative-control property and deliberately *excluded* the sectors that are downstream of visitor arrivals and would therefore violate it: retail trade (visitor spending), transportation, warehousing and utilities (which bundles air transportation), arts, entertainment and recreation, accommodation and food services, and real estate, rental and leasing (car and vacation rental). Empirically these fell 23–62% in the 2020 trough, tracking the tourism collapse itself, which is exactly the contamination that disqualifies them as controls. A second tier of genuinely insulated sectors, manufacturing, professional and business services, and information, could enlarge the donor pool, but their pre-2020 comovement with the tourism series is no stronger than the five retained donors' ($|\text{corr}| \leq 0.4$), so they do not sharpen identification of the tourism (travel-demand) factor. The one donor that does carry that factor is external, inbound air travel to Florida, added in Section 1.5.3.

1.4.3 Open-destination air travel (the travel-demand negative controls)

The insulated sectors span the pandemic’s general-recession factor but not its travel-demand-collapse factor (Section~1.3.2), so a second class of proxy is needed to carry that factor: inbound air travel to *exposed-but-open* destinations, places whose air traffic collapsed with the pandemic yet which imposed no traveler quarantine, and so register the travel-demand shock without registering Hawaii’s policy. I build these from the BTS Airline On-Time Performance database (Bureau of Transportation Statistics, 2024), scraped from the agency’s airport-statistics portal (<https://www.transtats.bts.gov/airports.asp>): monthly counts of arriving flights by airport, aggregated to a statewide total for Florida and to the individual airport for the Tampa (TPA), Orlando (MCO), and Phoenix (PHX) leisure metros, then converted to YoY growth. Florida is the archetype, a large, tourism-dependent state whose government kept its borders open through 2020, and the three metros diversify the proxy across independent leisure markets in states that likewise never locked down. Their identifying role, and why I fit them one at a time, is developed in Section~1.5.3.

1.4.4 Panel and estimation window

The treatment indicator `Mandatory Quarantine` switches to 1 in March 2020, giving 10 post-treatment months (March–December 2020); the post-window is capped at December 2020 so the extended raw panel cannot leak recovery months into the estimate. Although the panel ships raw levels from 1990 onward, the analysis truncates the pre-period to a recent window beginning January 2014 (74 pre-treatment months). Three decades of stale

comovement is less relevant to the 2020 counterfactual than the recent history, and the 2014 window empirically sharpens the pre-treatment fit and yields non-degenerate GMM standard errors that the full-history-plus-detrend combination did not. The full-panel version is available by removing the truncation and enters the robustness band of Section~1.5.




Figure 1.1: The negative-control premise: treated tourism outcomes collapse while the insulated-sector donors barely dip.

1.5 Results

1.5.1 The pre-treatment fit: relevance

Each estimate is only as good as the bridge's pre-treatment fit, the empirical footprint of the relevance condition (1.6); Figure 1.3 draws the per-outcome counterfactuals over time. Relevance holds where it matters: Visitor Days, Visitor Arrivals, Occupancy, and Revenue per Available Room all reconstruct to pre-treatment RMSEs of at most 3.6 growth-rate points (for Visitor Days, 2.6 points against a ~ 70 -point effect), so the collapse the estimator reads off is not an artifact of a poor fit.

The GMM standard errors are usable for the visitor series on this 2014 window (~ 3 –8 points), unlike the full-history window where de-trending a near-interpolating fit collapses them toward zero; for Visitor Days they give a 95% confidence interval of $[-78.3, -58.7]$ percentage points around the point estimate. That interval is sampling uncertainty for a single specification. The specification band below is a separate object, the spread of the point estimate across modeling choices; because the uncertainty here is driven more by specification and extrapolation than by sampling, I report both and read the band as the wider, more honest summary.

1.5.2 The insulated-only estimate: a bound

Fit on the insulated-sector proxies, the two headline specifications agree: for Visitor Days, SPSC returns an ATT of -68.5 percentage points (95% GMM-HAC confidence interval $[-78.3, -58.7]$) and base PI -76.9, and across the five headline tourism outcomes they never differ by more than 9 points. A regularized instrumental-variable bridge (SPSC) and a plain least-squares fit (PI) landing on the same effect from the same proxies makes it a property of the proxy design. But these are *upper bounds* in magnitude on the policy effect: the insulated sectors span the macro-recession factor and miss the travel-demand collapse (Section~1.3.2), so the tourism loading μ_1 lies outside the donor span (1.5) along that direction and the ATT still carries it. The observed drop is policy plus travel-demand loss; closing the gap to a point needs a travel-demand donor from outside the treated economy, next.

1.5.3 Completing the bound: an ensemble of exposed-but-open destinations

The open-destination air-travel proxies documented in Section~1.4, inbound flights to Florida and to the Tampa, Orlando, and Phoenix leisure metros, all exposed-but-open, are what close the travel-demand gap that leaves the insulated-only estimates as bounds. Rather than lean on one, I use the whole small ensemble; Florida is the archetype, its air traffic down about 70% at the 2020 trough. Each enters the long panel as its own donor unit (`group = travel-demand`), and, this matters, I fit them *one at a time*, the insulated five plus a single open-destination donor (I return to why below).

An open destination satisfies the proxy conditions from the other side: its inbound air travel loads on the same travel-demand collapse the insulated sectors miss (relevance), and as a separate economy under no Hawaii quarantine it carries that factor without importing the treatment (exclusion).

The exclusion that does the most work, though, is from Florida’s *own* policy, and it is what makes Florida a deliberate choice. For the proxy to read the pandemic’s voluntary travel-demand collapse cleanly, Florida’s air-travel decline must not itself be an artifact of a Florida lockdown, and it is not: Florida’s mandates were nominal and short-lived. [Bollyky et al. \(2023\)](#) measure each state’s policy stringency by the share of pandemic days its mandates (business and school closures, mask and stay-at-home orders, gathering limits) were in force and place Florida at the permissive end, imposing fewer restrictions on businesses and individuals than most states ([Liu et al., 2024a](#)); the state’s response was politically driven and centralized to the point of sidelining local emergency-management authority,

reopening early and later barring its own localities from reimposing restrictions ([Witkowski et al., 2023](#)). Behavior matched the policy. At the March 2020 peak of the shock Florida's beaches drew spring-break crowds in the thousands even as federal guidance urged against gatherings, phone-mobility data tracked those crowds dispersing across the eastern United States ([Business Insider, 2020](#)), genomic surveillance later traced new-variant introductions to that spring-break travel ([Napolitano et al., 2024](#)), and in Pinellas County, part of the Tampa Bay leisure market that furnishes one of the donors here, ethnographers recorded beach tourism continuing through the pandemic with its protocols widely unobserved ([Porter and McIlvaine-Newsad, 2022](#)). That inbound flights to Florida collapsed anyway, with little binding policy to explain it, is precisely the point: the decline is the voluntary travel-demand factor itself, which is exactly what the proxy must carry. The criterion has teeth, and it is not met by every warm-weather destination: Puerto Rico, a fellow island otherwise tempting as a donor, imposed a genuine islandwide stay-at-home order in mid-March 2020, so its tourism collapse is contaminated by its own policy and it is set aside for that reason.

The remaining objection is that a proxy like Florida still brings idiosyncratic confounders of its own, its hurricane season, its larger drive-to share of visitors. Two features contain that objection. First, the bridge weight is fit on the clean pre-period, so a Florida-specific shock that does not track Hawaii visitor volume before 2020 enters as fit noise and is down-weighted; what licenses the weight is the pre-2020 correlation of +0.35 and the pre-treatment RMSE it earns, and the post-period behavior plays no part in the fit. Second, the point does not rest on Florida alone: four independent markets across two states enter singly and converge

on the same band, so for an idiosyncratic confounder to move the estimate it would have to be common to all four and aligned with the March 2020 quarantine, which is the pandemic travel-demand collapse itself, the very factor the donors are meant to carry. Where a proxy fails relevance the procedure reports it and does not force a fit, as with Occupancy below, where Florida is down-weighted to zero.

Figure 1.2 reports the convergence, and the headline is robustness through agreement. For Visitor Days, every one of the four independent open-destination donors pulls the estimate off the insulated-only bound of -68.5 points into a policy-centered band of roughly -52 to -65 points, and each one *improves* the pre-treatment fit. That four proxies chosen independently, different cities, two different states, move the estimate the same way and by a similar amount is far stronger evidence that the observed drop is policy plus a real, separable travel-demand collapse than any single proxy could be. Florida gives the cleanest single fit (its Visitor Days pre-treatment RMSE falls from 2.58 to 1.92, bridge weight $w_{FL} = 0.14$), consistent with its pre-2020 correlation of +0.35 with Hawaii visitor volume.




Figure 1.2: The open-destination ensemble: each donor, fit singly, pulls the estimate off the insulated-only bound.

The ensemble also reaches an outcome a single volume proxy could not. Florida is a flight-*volume* series, so it comoves with Hawaii *visitor volume* but not with hotel prices, and leaves Mean Daily Rate essentially at its insulated-only value of -19.4 points. Phoenix, a desert-resort market whose room rates swing with leisure demand, does comove with Hawaii's daily room rate (pre-2020 correlation +0.36), and adding it pulls Mean Daily Rate from -19.4 toward -11.8 points while *improving* the fit (pre-treatment RMSE 2.38 to 1.82). Different open destinations carry the travel-demand factor as it manifests in different outcomes, volume for the Florida metros, price for Phoenix, so a small, well-chosen ensemble reaches further into the tourism block than any one proxy.

This is why I fit the donors *singly*. Added en masse, the four near-collinear proxies induce large offsetting ridge weights that *degrade* the fit where a single well-chosen donor helped most, the price series (Revenue per Available Room's pre-treatment RMSE rises from 3.61 to 4.41). Independent proxies read one at a time corroborate as a convergence band; collapsed into one super-donor, the collinearity punishes them.

Figure 1.3 is the summary figure: for every headline outcome it draws the observed series, the insulated-only SPSC counterfactual (macro factor only, the upper bound, with its conformal spikes), the counterfactual after adding Florida (travel-demand factor as well), and base PI's insulated path as a dotted overlay sitting almost on the SPSC insulated line, making the two-specification agreement visible. The wedge between the two SPSC lines is where the travel-demand donor works, wide for the visitor-volume series, absent for Occupancy where Florida is correctly down-weighted to zero. No single wedge is a clean decomposition, though;

each +donor fit re-solves all weights jointly, so the honest object is the convergence across donors in Figure 1.2.




Figure 1.3: Observed, SPSC and PI insulated counterfactuals, and SPSC + Florida, for five outcomes.

1.5.4 Robustness and falsification

1.5.4.1 A specification band

Because SPSC and base PI leave the bridge weights unconstrained, dropping the simplex restriction for the flexible $\mathbb{R}^{N_0}$ of [Doudchenko and Imbens \(2016\)](#), the counterfactual can extrapolate outside the donor hull, and the GMM sandwich, degenerate with a single instrument against five donors on a near-interpolating pre-window, reports a sampling noise that collapses toward zero. The uncertainty that actually binds is this extrapolation, or specification, risk, so I report the headline effects as a band built from specification sensitivity: de-trending on versus off, the B-spline degrees of freedom, the polynomial basis degree, dropping each of three insulated donors individually (the two largest-weight sectors and construction, the proxy most plausibly linked to tourism), and the full-1991 versus 2014 window. Figure 1.4 collects the minimum, median, and maximum SPSC ATT across these eight specifications for the three best-behaved outcomes.




Figure 1.4: Specification robustness band: median SPSC ATT (point) and min-max range across eight specifications.

For the visitor-volume series the band runs roughly from -74 to -57 points (median -68), so the defensible insulated-only statement is a ~ 60 –75-point upper-bound collapse.

1.5.4.2 How much would an exclusion violation have to bite?

The identifying content of the insulated proxies rests on exclusion: the border policy must not move them directly (Section~1.3.2). That is an assumption. I do not take it on faith; I bound the consequence of its failure, a sensitivity analysis in the spirit of proximal partial identification (Ghassami et al., 2023) that here follows the plausibly-exogenous relaxation of an exclusion restriction (Conley et al., 2012): posit a direct policy effect on the proxies and carry it through to the estimand. Suppose it fails: general-equilibrium spillovers from the closure depress or inflate the insulated sectors, so that over the post-period $p_{jt} = \boldsymbol{\mu}_j^\top \boldsymbol{\lambda}_t + \delta_j d_t + e_{jt}$, with δ_j the direct policy effect on proxy j and $\boldsymbol{\delta} = (\delta_j)_{j \in \mathcal{N}_0}$. Because the bridge weights are estimated over the pre-period, where $d_t = 0$ and the contamination is absent, $\widehat{\mathbf{w}}$ is untouched and the violation reaches the estimate only through the post-period counterfactual, as the exact bias

$$\widehat{\text{ATT}} = \text{ATT} - \boldsymbol{\delta}^\top \widehat{\mathbf{w}}, \quad (1.9)$$

so the bias is linear in the contamination and known up to $\boldsymbol{\delta}$. This makes the sensitivity fully transparent.

For Visitor Days the fitted weights sum to 1.92, so a common spillover of δ points moves the ATT by 1.92δ . Driving the insulated-only estimate of -69 points to zero would take $\delta \approx +36$ percentage points, the wrong sign (a border closure that *raised* health-care, finance,

and government employment growth by nearly forty points) and larger than any insulated sector's *entire* observed 2020 movement, which reached at most 12 points and is itself mostly the macro recession the proxies are meant to carry, with little policy spillover. In the economically natural direction, the policy depressing the insulated sectors, $\delta_j < 0$, the bias runs the other way: charging each sector's *full* 2020 decline to the border policy moves the estimate to -83 points, so the reported effect would if anything *understate* the truth. Even an adversarial assignment that signs every δ_j to mask the effect, bounded by each sector's observed movement, shifts the ATT by at most 18 points. Occupancy is more robust still: its weights sum to essentially zero (+0.03), so a common exclusion violation cancels outright. The conclusion does not require the insulated sectors to be perfectly insulated, it survives any exclusion violation of plausible magnitude, and the one plausible direction makes it conservative.

Equation (1.9) treats the insulated sectors symmetrically, but one is more suspect than the rest: construction tracks hotel building, so it is the insulated proxy most plausibly linked to tourism and the one a reader is likeliest to distrust. Dropping it outright is the direct negative-control check. Construction is not idle in the bridge, carrying a weight of -0.35 on Visitor Days, yet removing it leaves the estimate at -73 points against -69 with all five sectors, and the pre-treatment fit is if anything slightly better (RMSE 2.52 versus 2.58), because the four remaining sectors span the same macro-recession factor and absorb its removal. The direction is the telling part: the estimate moves *away* from zero. Were construction contaminated by the tourism collapse, its co-movement with the treated series would pull the

counterfactual down and bias the effect toward zero, so purging it would enlarge the effect, which is exactly what happens. The headline does not lean on the one donor a skeptic would challenge first.

1.5.4.3 A placebo in time

A distinct worry concerns the estimator itself: does the ridge-GMM bridge, with its spline de-trending, manufacture a large effect out of ordinary pre-period variation? The standard falsification is an in-time placebo, assign a fake border closure in a pre-pandemic March, truncate the sample before 2020 so no real shock is in view, and refit. A design that recovers effects only where they exist should return roughly zero.

It does. Across Visitor Days, Occupancy, and Revenue per Available Room, and fake treatment dates in March of 2017, 2018, and 2019, the placebo ATTs never exceed 4 points in magnitude, the scale of the pre-treatment RMSE, and an order of magnitude below the ~ 70 -point effects the same procedure returns for 2020. The bridge does not fabricate a collapse out of business-as-usual fluctuation, so the 2020 estimate is a property of the treatment.

1.5.5 Who bore the cost: incidence across islands and hotel tiers

The DBEDT workbook that supplies the statewide hotel outcomes also reports them by county and, for revenue per available room, by hotel class, so the same proximal machinery can ask not only how large the tourism collapse was but who absorbed it. For each disaggregated series I refit SPSC on the identical insulated donors and, as in Section~1.5.3, on the insulated set plus Florida; Figure 1.5 collects the estimates. Three readings stand out.

First, the demand shock was close to spatially uniform. Hotel occupancy fell between 65 and 76 points across all four counties, so the intuition that the tourism-dependent neighbor islands suffered a deeper demand hit than Oahu is not borne out in occupancy: the rooms emptied everywhere at once.

Second, the collapse ran through volume. The average daily room rate fell only 12 to 24 points, a fraction of the occupancy drop, so hotels largely held their posted rates while rooms went unsold. Because revenue per available room is occupancy times rate, the revenue loss is almost entirely the occupancy loss.

Third, the revenue burden fell regressively across the market. RevPAR by hotel class traces a clear gradient: the luxury and upper-upscale tiers lost the most (90 and 99 points), the upper-midscale tier markedly less among the well-fit segments (56), with midscale and economy least affected of all. Discretionary luxury-leisure travel evaporated where lower-tier and essential lodging demand partly persisted, so the establishments and workers tied to the top of the market bore the sharpest revenue hit.

These disaggregated panels are noisier than the statewide aggregates, and several cells (faded in Figure 1.5, pre-treatment RMSE above six growth-rate points) carry the near-zero GMM standard errors that Section~1.3.2 flags as fragile when the proxies supply one instrument against the donors. I read those cells as directional only; the three patterns above rest on the well-fit segments and hold under both the insulated bound and the Florida-adjusted point.




Figure 1.5: Subgroup ATTs by island and hotel class: insulated bound (circles), Florida-adjusted point (squares).

1.6 Discussion

The methodological lesson of Hawaii is that a good pre-treatment fit is not identification when a policy and a global shock arrive in the same month. A synthetic control that matches Hawaii to any pre-pandemic counterfactual, whether by borrowing across states or, as own-history extrapolation does, across the treated unit's own past, produces a clean-looking counterfactual and a large ATT, but that ATT is the border-closure effect plus the pandemic's contribution to Hawaii's outcomes, and nothing in the fit reveals the mix. Proximal / negative-control inference is what separates them: instrumenting the latent pandemic factor with insulated employment sectors that carry it but are immune to the policy, both headline specifications, SPSC and base PI, net out the macro-recession component and agree on the result, so the estimate rests on the identification; the two specifications only corroborate it.

At the same time, the analysis is candid about the limit of what the shipped data can identify. The insulated sectors span the pandemic's general-recession factor but not its travel-demand-collapse factor: even with an open border, Hawaii tourism would have fallen sharply as people stopped flying. SPSC on insulated sectors alone therefore recovers only the macro component of the counterfactual, and the resulting tourism ATTs over-attribute the collapse to the policy; they are upper bounds in magnitude, which is why I report them as a band. Recovering an actual point estimate from insulated donors alone is not possible here, because every treated series loads on the travel-demand factor those donors miss; the point has to be spanned from outside the treated economy, which is what the open-destination donors do. The substantive conclusion is thus deliberately modest and, I would argue, more

credible for it: the border closure coincided with a near-total collapse in measured tourism demand, of which a bounded share is causally the policy, with a convergent ensemble of external travel-demand donors pulling the Visitor Days estimate off its upper bound toward a policy-centered band.

Section~1.5.3 takes the further step of closing the two-factor gap with external travel-demand negative controls (inbound air travel to Florida and to a set of other exposed-but-open leisure markets that did *not* lock down), and the result is instructive about both the promise and the care the approach demands. For Visitor Days the open-destination donors do what the design promises, and they do it *in agreement*: fit one at a time, all four independently pull the estimate off the insulated-only bound into a policy-centered band, each improving the pre-treatment fit. Because the donors are chosen independently (different cities, two different states), their convergence is far stronger evidence that the missing factor is real and separable than any single proxy could be. A well-chosen ensemble also reaches further than one volume series: Phoenix, whose resort room rates track leisure demand, pulls the daily-room-rate series that Florida (a volume proxy) leaves essentially untouched. But the ensemble does not mechanically dissolve every bound, and the search must stay disciplined: near-collinear open destinations added *en masse* induce large offsetting weights and hurt the fit, so the donors are best read as a corroborating convergence band; and Puerto Rico, tempting as a fellow island, fails as a control because it imposed its own stay-at-home order in mid-March 2020 and is a contaminated, treated-like unit that, if anything, bounds the effect from the other side.

The opening states the stakes in dollars and jobs, and the estimand is a growth rate, so I close that gap here. Writing D_t^{2019} for 2019 visitor-days in calendar month t , $\hat{g}_t^0$ for the SPSC counterfactual year-over-year growth, and g_t for the observed growth, the 2020 counterfactual and observed levels are $\widehat{D}_t^0 = D_t^{2019}(1 + \hat{g}_t^0)$ and $D_t = D_t^{2019}(1 + g_t)$. The post-window visitor-day shortfall telescopes to

$$\widehat{\Delta D} = \sum_{t \in \mathcal{T}_2} (\widehat{D}_t^0 - D_t) = \sum_{t \in \mathcal{T}_2} D_t^{2019} (\hat{g}_t^0 - g_t), \quad (1.10)$$

the 2019 monthly level weighting each period's growth-rate gap. Valued at the 2019 average spend per visitor-day $\bar{s} = \$17.75\text{B}/D^{2019}$, the forgone visitor spending is

$$\widehat{\Delta S} = \bar{s} \widehat{\Delta D}. \quad (1.11)$$

With $\bar{s}$ at about \$198 per visitor-day, the insulated-bound fit implies a shortfall of about 51 million visitor-days over March–December 2020, against a counterfactual of 64 million, worth roughly \$10 billion in lost visitor spending; the Florida-adjusted point gives about \$9.7 billion. Because the insulated estimate is an upper bound in magnitude on the policy effect (Section~1.3.2), this figure bounds from above the spending loss attributable to the quarantine, set against the \$17.75 billion in annual visitor spending the opening put at stake. The translation assumes spending tracks visitor-days at the 2019 average, so it should be read as an order-of-magnitude gloss on the growth-rate estimate.

These estimates are one side of a ledger. They price the short-run economic cost of the

border closure through the end of 2020; they do not, on their own, value the public-health benefit (lives saved, hospital strain averted) the closure was meant to buy. A bounded cost, though, is exactly the input a cost-benefit calculation needs, and the surrounding literature supplies the other side. Work on early containment finds that acting even days sooner averts a large share of deaths: [Bayat et al. \(2020\)](#) estimate that locking down ten days earlier could have cut New York deaths by more than 80%. Evidence on comparably aggressive closures finds the economic damage sharp but short-lived ([Ke and Hsiao, 2022](#)). Set side by side, the calculus for a decisive closure is a large but transitory economic loss weighed against a benefit measured in lives, and this paper supplies the first credible, bounded estimate of that loss for the U.S. case. The incidence analysis of Section~[1.5.5](#) adds that the loss fell unevenly: hardest on the upper tiers of the lodging market, and through lost volume, with room rates largely holding. What the paper also provides is a template for pricing such policies at all: when the intervention and the crisis it responds to are one and the same event in time, no comparison unit is clean, and the analyst’s task is to find a variable that sees the crisis but not the policy.

1.7 Conclusion

This paper estimates the short-run economic consequences of Hawaii’s March 2020 mandatory quarantine while confronting the identification problem that its predecessors on this topic could not: for a single treated unit, the border-closure policy and the COVID-19 pandemic are perfectly collinear in time, so no synthetic control that merely matches Hawaii to a

pre-pandemic counterfactual can separate the two. I resolve this with proximal inference (negative-control estimators) in two specifications (Single Proxy Synthetic Control and base Proximal Inference) that instrument the latent pandemic factor using insulated employment sectors which carry the pandemic’s macro-recession shock but are immune to the border policy.

Estimated by both specifications, the insulated-sector proxies identify a large, robust tourism-demand collapse (about 60–75 points of year-over-year growth for the visitor-volume series). But because those sectors span only the macro-recession factor and not the travel-demand-collapse factor, they under-span the tourism factor and identify these effects only weakly: the insulated-only estimates are upper bounds in magnitude on the policy effect, reported as a robustness band. Because every treated series loads on the travel-demand factor those donors miss, the trustworthy point comes only once the missing factor is spanned from outside the treated economy, the role of the external open-destination travel-demand donors.

Methodologically, the study demonstrates how negative-control inference yields credible causal estimates when a policy and a global shock coincide, a structure common to pandemic border closures, but also to climate shocks and other systemic crises where a would-be counterfactual unit is itself engulfed by the event. Substantively, it delivers a bounded, honest estimate of the economic price of Hawaii’s “Great Isolation with Aloha characteristics’’. Adding a small ensemble of external no-lockdown travel-demand donors (inbound air travel to Florida and to the Tampa, Orlando, and Phoenix leisure markets) takes the step from a bound to a policy-centered band: fit independently, they converge in pulling the Visitor Days effect

toward the policy-only magnitude, and one of them (Phoenix) reaches the room-rate series a single volume proxy cannot. Finishing the job for every outcome asks only for more such open-destination proxies, disciplined by relevance, which the design readily accommodates.

CHAPTER 2

Clearing the Air: How India's 2020 Lockdown Impacted Air Quality

2.1 Introduction

The COVID-19 pandemic prompted governments worldwide to implement non-pharmaceutical interventions (NPIs), such as mass lockdowns, to curb virus spread and protect public health systems. India's nationwide lockdown, beginning on March 25, 2020, was among the strictest globally since the Wuhan lockdown in China, halting mobility, industry, and construction for 1.3 billion people. Early research focused on virus outcomes (Yang, 2021; Bayat et al., 2020) and economic impacts (Gong et al., 2024; Wu et al., 2023), but lockdowns also likely influenced air quality, a critical public health factor. Popular media reported improved air quality during lockdowns, later confirmed by studies documenting reductions in pollutants such as PM_{2.5} (Watts and Kommenda, 2020; Venter et al., 2020; Kumari and Toshniwal, 2021). Given that India has some of the most polluted cities in the world, understanding how large-scale interventions affect pollution levels is of great interest. However, many studies are descriptive and do not estimate counterfactual scenarios (what air quality would have been without lockdowns), limiting their policy relevance, or rely on strict assumptions for identification (Milman and Uteuova, 2025).

This study quantifies the causal impact of India's 2020 national lockdown on PM_{2.5} levels nationally and in major cities using the Synthetic Historical Control (SHC) method (Chen et al., 2024). A contemporaneous pool of untreated units, which synthetic control normally

requires, is difficult to assemble here because the pandemic affected countries worldwide. SHC avoids this requirement by constructing the counterfactual from the treated unit’s own pre-lockdown history. Identification therefore turns on the stability and reconstructability of the treated unit’s historical outcome process, which relaxes the need to measure every confounder (Section 2.2.3). In addition to the national analysis, this study examines Delhi, Mumbai, Bangalore, and Kolkata, megacities consistently among the world’s most polluted and major hubs of finance and commerce. The outcome throughout is the year-over-year growth rate of PM2.5, so the estimated effects are percentage-point changes in that growth rate (Section 2.3). Embedding these cases in a unified causal framework moves beyond simple pre–post comparisons and addresses a key policy question: what happens to air pollution when a country of India’s scale implements a nationwide shutdown? These findings can inform sustainable air quality management in India and provide lessons for urban planning and climate change mitigation globally.

The rest of the chapter is organized as follows. Section 2.2.1 provides background on air pollution in India and its health and economic costs. Section 2.2.2 reviews the existing literature and the methodological gap this chapter fills. Section 2.2.2.2 sets out the empirical strategy: the SHC estimator with its identifying assumptions (Section 2.2.3) and the choice of smoothing bandwidth (Section 2.2.4). Section 2.3 describes the data and their construction. The remaining sections report the national and city-level results, a quarterly-aggregation robustness check, and a discussion of the findings and their policy implications.

2.2 Background and Previous Work

2.2.1 Policy Context

Air pollution is one of India's most pressing environmental and public health challenges. This section provides context on pollutant sources, exposure levels, health and economic impacts, and how the COVID-19 lockdown created a natural experiment for studying air quality.

2.2.1.1 Sources and Severity of Air Pollution

Fine particulate matter (PM_{2.5}), which penetrates deep into the lungs and bloodstream, is a primary pollutant of concern. Major contributors include vehicular emissions, industrial activities, construction dust, biomass burning for cooking and heating, and agricultural residue burning, particularly in northern states during winter. Natural factors such as dust storms and seasonal weather patterns exacerbate pollution levels. Rapid urbanization and population growth have intensified exposure, affecting over 1.4 billion people.

Recent data highlight the severity: in 2024, India's average annual PM_{2.5} concentration was $50.6 \mu\text{g}/\text{m}^3$, far exceeding the WHO guideline of $5 \mu\text{g}/\text{m}^3$ ([Statista, 2025](#)). Regional variation is significant: cities in the Indo-Gangetic Plain, like Delhi, often exceed $100 \mu\text{g}/\text{m}^3$ during peak winter, while southern cities such as Chennai frequently surpass both WHO and national standards ([Bhuyan et al., 2025](#); [Singh et al., 2021](#)).

2.2.1.2 Health Impacts

The health consequences of chronic PM_{2.5} exposure are substantial. Long-term exposure increases the risk of respiratory and cardiovascular diseases, including COPD, asthma, lung cancer, stroke, and heart disease ([Joshi et al., 2025](#)). In India, air pollution causes roughly 1.5 million premature deaths annually, with each 10 $\mu\text{g}/\text{m}^3$ rise in PM_{2.5} linked to an 8.6% increase in mortality ([Jaganathan et al., 2024](#)). Vulnerable populations, namely children, the elderly, and those with weaker immune systems, are most affected. Fine particulate pollution reduces average life expectancy by 5.3 years compared to WHO guidelines, with residents in Delhi facing up to a 10-year reduction ([BBC, 2022](#)). Short-term PM_{2.5} spikes also increase daily death rates by 1.42–3.57% per 10 $\mu\text{g}/\text{m}^3$ ([de Bont et al., 2024](#)).

2.2.1.3 Economic Impacts

Air pollution imposes significant economic costs through healthcare expenditures, lost labor productivity, and impacts on sectors such as tourism and agriculture. Estimates suggest annual losses of USD 95 billion, roughly 1.36% of GDP ([India State-Level Disease Burden Initiative Air Pollution Collaborators, 2020](#); [Gupta and Damani, 2021](#)). In the Delhi NCR, productivity declines and reduced tourist arrivals contribute substantially to these costs ([Dalberg Advisors, 2021](#); [Our Common Air Commission, 2024](#)). Efforts like the National Clean Air Programme (NCAP), launched in 2019, aim to reduce PM_{2.5} and PM₁₀ concentrations by 20–30% in 131 non-attainment cities by 2024–25 ([Kawano et al., 2025](#)).

2.2.1.4 The COVID-19 Lockdown as a Natural Experiment

The 2020 nationwide COVID-19 lockdown provides a unique opportunity to study rapid air quality improvements. Imposed on March 25, the lockdown halted transportation, industrial operations, and construction, leading to reductions in PM_{2.5} by 31–43%, with some cities experiencing drops up to 50% (Saleem et al., 2024; Nigam et al., 2021; Mishra et al., 2021). In megacities like Delhi and Mumbai, the Air Quality Index (AQI) improved by as much as 58%, primarily due to lower vehicular and industrial emissions (Zhang et al., 2020).

Although temporary, these reductions highlight the potential of targeted interventions. PM_{2.5} reductions were most pronounced in early lockdown phases, averaging 30–50% across monitored sites, with occasional spikes due to meteorology (Hu et al., 2022; Goel et al., 2021; Wang et al., 2020). Other pollutants, including NO₂ and PM₁₀, also decreased, consistent with global lockdown trends (Venter et al., 2020). Understanding these dynamics is crucial for causal evaluation and future policy planning.

2.2.1.5 Lockdown Policy Phases

India’s lockdown, enacted under the Disaster Management Act of 2005, unfolded in multiple phases. Phase 1 (March 25–April 14) involved a complete nationwide shutdown, prohibiting non-essential movement, closing factories, and suspending public transport. Phase 2 (April 15 onwards) allowed a gradual reopening of essential services and agriculture, with color-coded zoning based on infection rates (Deshpande, 2022; Ranjan et al., 2020). While primarily a public health measure, the lockdown served as a large-scale quasi-experiment in emission

control, revealing the disproportionate contribution of human activities to PM2.5 levels.

2.2.2 Literature Review

2.2.2.1 Methodological Aspects of Prior Literature

To contextualize our approach, we review existing studies on lockdown-related air quality changes and highlight key methodological gaps. Early studies primarily relied on descriptive pre–post comparisons or satellite imagery without causal adjustment ([Mahato et al., 2020](#); [Pathakoti et al., 2021](#); [Srivastava et al., 2020](#); [He et al., 2020](#); [Wang et al., 2020](#)). These approaches document reductions in pollutants but cannot answer counterfactual questions: what would air quality have been absent the lockdown? Some studies have applied causal models. For example, [Heffernan et al. \(2025\)](#) use causal machine learning and interrupted time series (ITS) methods to find reductions in NO2 levels. ITS was applied in Italy ([Cameletti, 2020](#)), finding no meaningful changes in air pollutants. Augmented Synthetic Control methods were used for Wuhan ([Cole et al., 2020](#)), showing a 63% reduction in NO2. Despite these examples, a recent systematic review finds that most lockdown–air pollution studies remain descriptive or non-causal ([Silva et al., 2022](#)), limiting their policy relevance.

2.2.2.2 Gaps and Contribution of This Study

A central gap is the lack of causal designs for national-scale interventions. Many studies use simple pre–post t-tests ([Goel et al., 2021](#)), graphical analyses ([Kumari and Toshniwal, 2020](#); [Benchrif et al., 2021](#)), year-over-year comparisons ([Le Quéré et al., 2020](#)), or OLS regressions ([Pal et al., 2022](#)). These methods cannot fully account for confounding trends or

temporal patterns, reducing confidence in policy-relevant estimates. SCM improves on this by constructing a counterfactual from untreated donor units as a weighted average ([Ben-Michael et al., 2021a](#)). For a nationwide policy, though, a credible contemporaneous donor pool is hard to assemble: the pandemic affected countries worldwide, so candidate controls are themselves treated or otherwise contaminated. Causal machine learning approaches ([Heffernan et al., 2025](#)) face another difficulty, relying on strong ignorability assumptions that may not hold for air pollution data with unmeasured year-specific factors. SHC takes a different route. It reconstructs the counterfactual from the treated unit’s own pre-lockdown history, so it needs neither a contemporaneous donor pool nor a complete set of measured covariates. This does not remove the identification burden; it shifts it onto the stability and reconstructability of that historical outcome process, the assumptions I set out in Section [2.2.3](#). This chapter provides what is, to my knowledge, the first application of SHC to estimate the air-quality consequences of the 2020 lockdown in a country of India’s scale, at both the national and city levels.

Methodology

Below I describe the Synthetic Historical Control (SHC) methodology from [Chen et al. \(2024\)](#). SHC is applied separately to the national average of PM2.5 for India, as well as Delhi, Mumbai, Bangalore, and Kolkata.

2.2.3 The Synthetic Historical Control Method

SHC estimates the causal effect of an intervention by constructing counterfactual outcomes from historical segments of the treated unit’s own past. This makes it well-suited for national-scale interventions such as India’s 2020 lockdown, where assembling a credible contemporaneous donor pool of untreated countries is difficult because the pandemic affected countries worldwide. Before presenting the estimator, I state the assumptions SHC relies on.

2.2.3.1 Assumptions

Assumption 1 (Error Properties).

$$\mathbb{E}[\varepsilon_t] = 0, \quad \mathbb{E}[\varepsilon_t^2] < \infty \quad \forall t \in \mathcal{T}. \quad (2.1)$$

This ensures that noise is well-behaved. The mean-zero condition rules out systematic bias, while the finite variance condition prevents extreme fluctuations.

Assumption 2 (Latent Smoothness and Matching).

$$\tilde{y}(t) \in C^{H+1}, \quad H \geq 3, \quad \sup_t \left| \frac{d^k}{dt^k} \tilde{y}(t) \right| < \infty \quad \forall k \leq H + 1, \quad (2.2)$$

and there exists

$$\mathbf{w}^* \in \Delta^{N-1} \quad \text{such that} \quad \mathbf{y}_{eval} = \widetilde{\mathbf{Y}}_{pre} \mathbf{w}^*. \quad (2.3)$$

Smoothness ensures that the latent process can be approximated by a finite expansion. The

matching condition requires that some convex combination of historical donor segments can reproduce the evaluation period. This parallels the “existence of weights” assumption in the SCM literature (Abadie et al., 2010; Amjad et al., 2018). Recurring influences such as meteorology enter the counterfactual only to the extent that they are stable and are represented by the historical donor segments; their past presence in the series is required for this, and on its own it leaves the match still to be earned.

Assumption 3 (Feasibility and Stability).

$$\exists \mathbf{w} \in \Delta^{N-1} \text{ sparse, such that } \mathbf{y}_{eval} \approx \widetilde{\mathbf{Y}}_{pre} \mathbf{w}, \quad \widetilde{\mathbf{Y}}_{pre}^\top \widetilde{\mathbf{Y}}_{pre} \succ 0. \quad (2.4)$$

Sparsity means that only a small number of donor segments are relevant. Positive definiteness prevents collinearity and keeps the weights stable. In plain terms, SHC rests on three conditions. First, the random “noise” in pollution measurements behaves well: it averages out and does not swamp the real signal. Second, the history of pollution before the lockdown is smooth and informative enough that past patterns can be combined to mimic what would have happened absent the lockdown. Recurring influences such as weather are carried along only to the extent that they leave a stable, repeating footprint in that history; the method does not measure or adjust for them directly. Third, the pieces of history used to form the counterfactual are not redundant, so the weights placed on them stay stable across the sample. These conditions let SHC reconstruct the final pre-treatment window. Using that reconstruction as the post-treatment counterfactual takes one further, substantive assumption:

that the historical relationship the weights capture is stable enough to carry forward into the post-treatment period. Good pre-treatment fit is a precondition here; what ultimately identifies the effect is that stability, and it is the assumption I scrutinize most closely in what follows.

2.2.3.2 Estimator

Let $\mathcal{T} = \{1, \dots, T\}$ index time, with pre-treatment period $\mathcal{T}_1 = \{t \in \mathcal{T} : t \leq T_0\}$ and post-treatment period $\mathcal{T}_2 = \{t \in \mathcal{T} : t > T_0\}$. The number of post-treatment periods is $n = T - T_0$, and an evaluation window of length m is defined as the final m pre-treatment periods $\mathcal{T}_{\text{eval}} = \{T_0 - m + 1, \dots, T_0\}$. Let $\mathbf{y} = (y_1, \dots, y_T)^\top$ denote the observed outcome series. The evaluation subvector is $\mathbf{y}_{\text{eval}} = \mathbf{y}_{\mathcal{T}_{\text{eval}}} \in \mathbb{R}^m$, and the post-treatment outcome vector is $\mathbf{y}_{\text{post}} = \mathbf{y}_{\mathcal{T}_2} \in \mathbb{R}^n$. For each $t \in \mathcal{T}$, let y_t^1 and y_t^0 denote the potential outcomes under treatment and no treatment, respectively. The observed outcome is

$$y_t = d_t y_t^1 + (1 - d_t) y_t^0, \quad d_t = \mathbf{1}\{t > T_0\}. \quad (2.5)$$

For $t > T_0$, only y_t^1 is observed. SHC provides an estimate $\hat{y}_t^0$ for the missing counterfactual, and the treatment effect is estimated as

$$\hat{\alpha}_t = y_t - \hat{y}_t^0, \quad t \in \mathcal{T}_2, \quad (2.6)$$

with the average treatment effect on the treated (ATT) defined as

$$\widehat{\text{ATT}} = \frac{1}{n} \sum_{t \in \mathcal{T}_2} \hat{\alpha}_t. \quad (2.7)$$

With a single treated unit and a time-series intervention, this ATT is the average post-treatment effect for that one unit: the mean gap between its observed path and its counterfactual over $\mathcal{T}_2$. I keep the ATT label for continuity with the SHC and SCM literature, where it carries this time-series meaning.

The implementation of SHC begins by smoothing the pre-treatment outcome series $\mathbf{y}_{\text{pre}}$ using local linear regression with a Gaussian kernel, yielding a smooth trajectory $\tilde{\mathbf{y}}$ (Fan and Gijbels, 1992). From this smoothed series, we construct the donor matrix by extracting N overlapping segments, each of length m .




Figure 2.1: Constructing the donor matrix $\tilde{\mathbf{Y}}_{\text{pre}}$ by sliding overlapping windows of length m across the smoothed pre-treatment series $\tilde{\mathbf{y}}$. Each colored segment corresponds to a column in the donor matrix.

The evaluation vector $\mathbf{y}_{\text{eval}}$ is reconstructed as a convex combination of the donor segments by solving the quadratic program

$$\mathbf{w}^* = \arg \min_{\mathbf{w} \in \mathbb{R}^N} \left\| \mathbf{y}_{\text{eval}} - \widetilde{\mathbf{Y}}_{\text{pre}} \mathbf{w} \right\|_2^2 + \varsigma \left\| \mathbf{C}_0^\top \mathbf{w} \right\|_2^2, \quad \text{s.t. } \mathbf{w} \geq 0, \mathbf{1}^\top \mathbf{w} = 1,$$

where $\mathbf{C}_0$ contains the eigenvectors of the donor Gram matrix associated with small eigenvalues.

The second term penalizes directions of low variance, improving numerical stability when donor segments are collinear.

In practice, only a few donor segments contribute meaningfully to the reconstruction of $\mathbf{y}_{\text{eval}}$.

To identify these, I use the forward selection algorithm described by [Chen et al. \(2024\)](#).

Applying the resulting weight vector $\mathbf{w}^*$ to the aligned post-treatment donor matrix produces $\hat{\mathbf{y}}_{\text{post}}^0$, the SHC estimate of the untreated counterfactual trajectory.

A further advantage of SHC in this setting is that it naturally accommodates nonlinear dynamics in the pre-treatment pollution series. Because SHC begins with a local linear smoother, short-term fluctuations and seasonal nonlinearities (many of which are driven by meteorology such as temperature inversions, wind, or rainfall) are already captured in the donor segments. In contrast, standard interrupted time series designs often require explicitly specifying and fitting complex parametric models such as SARIMA to handle the same effects. The primary tuning choice is therefore the bandwidth of the kernel smoother, while the identifying content sits in the smoothness, matching, and stability conditions stated above. This makes the method flexible enough to capture weather-driven variation while remaining

relatively simple to implement in practice.

2.2.4 Bandwidth Selection

The performance of SHC depends on the smoothing bandwidth h , which is chosen by cross-validation. Let $\mathbf{y}_{\text{pre}} = (y_1, \dots, y_{T_0})^\top$ denote the pre-treatment series. For a candidate bandwidth $h > 0$, define the local linear smoother estimate $\hat{y}_i^{(-i)}(h)$ as the fitted intercept obtained by regressing $\{y_j : j \neq i\}$ on $\{j - i : j \neq i\}$ with kernel weights

$$K_h(j - i) = \exp\left(-\frac{1}{2} \left(\frac{j-i}{h}\right)^2\right),$$

normalized so that $\sum_{j \neq i} K_h(j - i) = 1$. Formally, for each $i \in \{1, \dots, T_0\}$,

$$\hat{\beta}(h)^{(-i)} = \arg \min_{\beta_0, \beta_1} \sum_{j \neq i} K_h(j - i) \left(y_j - \beta_0 - \beta_1(j - i)\right)^2,$$

and the leave-one-out prediction is $\hat{y}_i^{(-i)}(h) = \hat{\beta}_0(h)^{(-i)}$. The LOOCV error is $\text{CV}(h) = \frac{1}{T_0} \sum_{i=1}^{T_0} \left(y_i - \hat{y}_i^{(-i)}(h)\right)^2$. The optimal bandwidth is chosen as $h^* = \arg \min_{h \in \mathcal{H}} \text{CV}(h)$, where $\mathcal{H}$ is a prespecified candidate grid, and is fixed before SHC estimates any weights. This choice trades smoothing bias against noise in the local linear smoother.

2.3 Data

My data on PM2.5 are derived from [Batra \(2025\)](#), who keep data collected from the Socio-economic High-resolution Rural-Urban Geographic Data Platform for India (SHRUG). We have the full subdistrict-level PM2.5 matrix:

$$\mathbf{Y} = \begin{bmatrix} y_{1,1} & y_{1,2} & \cdots & y_{1,T} \\ y_{2,1} & y_{2,2} & \cdots & y_{2,T} \\ \vdots & \vdots & \ddots & \vdots \\ y_{N,1} & y_{N,2} & \cdots & y_{N,T} \end{bmatrix}, \quad N = 5,968, \quad T = 312$$

where rows index the N subdistricts (the SHRUG block units, nested within 634 districts and indexed $i = 1, \dots, N$), columns index months ($t = 1, \dots, T$), and $y_{i,t}$ is the average PM2.5 in subdistrict i during month t . Each subdistrict i carries a population p_i . For the national series we take the population-weighted average of PM2.5 across all subdistricts,

$$\mathbf{y}_{\text{India}} = \frac{\sum_{i=1}^N p_i \mathbf{Y}_{i,\cdot}}{\sum_{i=1}^N p_i} \in \mathbb{R}^T,$$

so that each subdistrict contributes in proportion to the population exposed. For the city-level treated units ($c \in \{\text{Delhi, Mumbai, Bangalore, Kolkata}\}$) we form the same population-weighted average over the subdistricts that make up the city,

$$\mathbf{y}_c = \frac{\sum_{i \in S_c} p_i \mathbf{Y}_{i,\cdot}}{\sum_{i \in S_c} p_i} \in \mathbb{R}^T,$$

where S_c is the set of subdistricts belonging to district c . My final outcome is the year-over-year (YoY) growth rate of PM2.5. Working on the growth-rate scale removes the strong twelve-month seasonality of Indian pollution and the slow-moving level trend, both of which

would otherwise dominate a historical-control comparison. For any series $\bar{\mathbf{y}}$, the monthly YoY growth from month t relative to the same month a year earlier is

$$g_t = \frac{\bar{y}_t - \bar{y}_{t-12}}{\bar{y}_{t-12}}, \quad t = 13, \dots, T,$$

so that the treated units for analysis are

$$\mathbf{g}_{\text{India}} = [g_{13}, \dots, g_T], \quad \mathbf{g}_c = [g_{13}, \dots, g_T] \forall c.$$

Because every estimate is formed on this scale, the treatment effect is a percentage-point change in the year-over-year growth rate of PM2.5. The YoY transform is defined from month 13 onward (January 1999), and the lockdown analysis uses this series through December 2020.

2.4 Results

Figure 2.2 presents the Synthetic Historical Control estimate for the all-India population-weighted PM2.5 series. Throughout, the outcome is the year-over-year growth rate of PM2.5, so every reported effect is a percentage-point change in that growth rate (Section 2.3). The observed series falls sharply below its synthetic-historical counterfactual following the March 2020 lockdown, indicating a pronounced short-run reduction in the growth of particulate pollution. Figure 2.3 repeats the exercise for the four megacities (Delhi, Mumbai, Bangalore, and Kolkata).




Figure 2.2: Synthetic Historical Control: all-India year-over-year PM2.5 growth, observed versus counterfactual. The dotted line marks the March 2020 national lockdown; the shaded region is the 90% Andrews–Genton conformal prediction band.




Figure 2.3: Synthetic Historical Control for the four megacities: observed versus counterfactual year-over-year PM2.5 growth, with 90% conformal prediction bands (shaded).

Inference follows the conformal permutation procedure of [Chen et al. \(2024\)](#): the pre-treatment residuals are resampled to build a null distribution for the post-treatment gap (1{,}000 resamples), and an Andrews–Genton conformal band gives a 90% prediction band for the counterfactual, shaded in [Figure 2.2](#) and [Figure 2.3](#). The national effect has a permutation p-value below 0.001; the Delhi (below 0.001), Mumbai (below 0.001), and Kolkata (below 0.001) effects clear conventional thresholds as well, and Bangalore, the smallest effect, has a p-value of 0.020. The pre-treatment root-mean-squared error is the relevance diagnostic: 6.3 pp nationally and 6.0 pp for Delhi, comfortably below those units’ effects. For Mumbai (20.7 pp), Kolkata (28.4 pp), and Bangalore (10.4 pp) the pre-treatment fit is coarser, of the same order as the estimated effect, so those city estimates carry more caution and anchor the discussion less than the national and Delhi results.

2.5 Robustness: Quarterly Aggregation

Monthly PM2.5 growth rates are inherently noisy: meteorology, episodic events (festivals, crop-residue burning), and measurement error all inject high-frequency variation that can exaggerate or mask any single month’s estimated effect; the September 2020 rebound visible in [Figure 2.2](#) is one such idiosyncratic spike. To confirm that the headline results are not an artifact of this monthly noise, I re-estimate every model on quarterly data: the population-weighted PM2.5 series is averaged into calendar quarters before the year-over-year growth rate is formed and the SHC estimator is re-run. Because the strict lockdown spanned the entire second quarter of 2020, quarterly aggregation smooths out month-to-month shocks

while preserving the policy signal.




Figure 2.4: Robustness of the SHC estimates to temporal aggregation: the average treatment effect (percentage points of year-over-year PM2.5 growth) estimated at the monthly frequency versus the same estimator on quarterly data, for India and the four megacities.

Figure 2.4 compares the monthly and quarterly SHC estimates for all five units, and the two are strikingly close. Nationally, the quarterly ATT is 13.8 pp against 13.8 pp at the monthly frequency; the city estimates likewise track their monthly counterparts to within about two percentage points (Delhi 13.2, Mumbai 10.7, Bangalore 9.2, and Kolkata 33.5 pp). The cross-city ordering is preserved (Kolkata largest, Bangalore smallest), and every effect remains a substantial reduction. Because the point estimates change little whether the data are analyzed monthly or quarterly, the measured effect is not driven primarily by a small number of high-frequency monthly observations; aggregating to quarters smooths month-to-month noise and leaves the estimates substantially unchanged.

2.6 Discussion

Across every series, the Synthetic Historical Control (SHC) estimates in Figure 2.2 and Figure 2.3 point in the same direction: India’s 2020 lockdown sharply suppressed fine-particulate pollution. Because the outcome is the year-over-year (YoY) growth rate of PM2.5 (Section 2.3), each estimate is a change in that growth rate measured in percentage points (pp). Nationally, the average treatment effect on the treated is a 13.8 pp reduction: the synthetic-historical counterfactual implies all-India PM2.5 would have *grown* by roughly 7.3% year over year on average between March and December 2020, yet the observed series fell instead. The wedge between the two is the causal effect of the shutdown, and it is far larger than the ordinary year-to-year movement of the pre-treatment series.

The effect is concentrated where the policy was strictest. Reductions are deepest during the Phase 1–2 window of April–June 2020, when mobility, industry, and construction were halted nationwide (Section 2.2.1); the largest single-month national gap occurs in May 2020, about 26.3 pp below counterfactual. From mid-summer the effect attenuates as the economy reopened, and a brief positive deviation around September 2020 appears in several series. This profile, a deep trough under the strictest restrictions followed by mean reversion, is the mark of a temporary non-pharmaceutical intervention whose force faded as the economy reopened. The pattern is also consistent with a strength of SHC in this setting: recurring meteorological drivers (Section 2.2.3) are absorbed into the historical donor segments, so the estimated effect is not an artifact of a single anomalous season.

City-level estimates show meaningful heterogeneity (Figure 2.3). The reduction is 14.2 pp in Delhi, 11.6 pp in Mumbai, 9.9 pp in Bangalore, and 31.5 pp in Kolkata. Kolkata's effect is both the largest and the most persistent: its counterfactual was on a steep upward path (averaging about 17.8% YoY growth), so the shutdown opened an especially wide gap that endured through the autumn. Delhi shows a steep spring collapse followed by a rebound late in the year, when post-monsoon crop-residue burning and winter temperature inversions, drivers outside the lockdown's reach, reassert themselves. Bangalore, with a lighter industrial base and a cleaner baseline, shows the smallest reduction, consistent with its pollution being less sensitive to the halt in heavy industry and construction. These cross-city differences are descriptive: the historical fits differ across cities, and I do not test whether one city's effect is statistically larger than another's, so the ordering is best read as descriptive heterogeneity across cities.

These causal estimates are broadly consistent with, but conceptually distinct from, the descriptive 31–43% concentration declines reported in the lockdown literature (Nigam et al., 2021; Saleem et al., 2024): SHC improves on a raw before-and-after comparison by benchmarking the observed series against what its own history implies should have happened. Several caveats temper interpretation. The estimates are for a transitory shock and speak to short-run responsiveness; they carry no implication for a sustainable abatement path. The outcome is a growth rate, so a negative ATT denotes slower growth, here outright decline, measured against the counterfactual. The September rebound shows that the design recovers net effects, including any offsetting seasonal forces. Aggregating to quarters (Section

2.5) averages out this high-frequency variation and leaves the point estimates substantially unchanged. Uncertainty is quantified by the conformal permutation test of Section 2.4; with a single treated unit and no donor pool, the design still lacks the cross-sectional placebo distribution that many synthetic-control applications use, so the permutation test runs over the treated unit’s own pre-treatment residuals.

2.7 Conclusion

This chapter provides what is, to my knowledge, the first set of causal Synthetic Historical Control estimates of the air-quality consequences of India’s 2020 national lockdown, at both the national level and for four of its largest megacities. The lockdown lowered the year-over-year growth of PM2.5 by about 13.8 pp nationally, with city effects ranging from 9.9 pp in Bangalore to 31.5 pp in Kolkata. In every case the observed pollution path fell below a counterfactual that, on its own historical momentum, would have continued to rise. Two implications follow. First, the speed and size of the response indicate that a large share of India’s particulate burden is anthropogenic and responsive to economic activity such as transport, industry, and construction; the design recovers the net effect of the shutdown and does not separately identify the contribution of each sector. Second, the rapid rebound once restrictions eased shows that one-off shocks do not deliver lasting gains: realizing the air-quality improvements seen in 2020 would require the sustained, structural emission controls envisioned by the National Clean Air Programme (Section 2.2.1), whereas the 2020 shutdown bought only a temporary reprieve.

The analysis also points to clear avenues for future work: probing sensitivity to the evaluation-window length and the choice of post-treatment endpoint, extending the outcome set beyond PM2.5 to co-pollutants such as NO₂, and tracing how quickly pollution returns to its pre-pandemic trajectory. Taken together, the results establish a large short-run reduction in the growth of particulate pollution and show the practical value of historical-control methods for evaluating large-scale interventions for which a contemporaneous untreated comparison unit is hard to assemble. The health consequences of pollution at these levels are well documented (Section [2.2.1](#)); converting the estimated air-quality change into health outcomes is beyond what this chapter estimates and is left to that literature.

CHAPTER 3

Locking Away Prosperity? Evaluating the Labor Impacts of Vaccine Mandates

3.1 Introduction

The COVID-19 pandemic triggered a global wave of non-pharmaceutical interventions (NPIs) as governments acted swiftly to limit viral transmission. Much early research focused on the public health rationale for these policies, especially the importance of early action to contain exponential spread. For instance, [Friedson et al. \(2021\)](#) and [Yang \(2021\)](#) use synthetic control methods to evaluate lockdowns in California and Wuhan, emphasizing the critical role of timing. Similarly, [Bayat et al. \(2020\)](#) estimate that New York City could have reduced COVID-19 deaths by over 80% had its first lockdown occurred just ten days earlier. Subsequent work examined the economic consequences of such interventions. While NPIs helped flatten the curve, they also raised concerns about business closures, layoffs, and lasting economic scarring. Studies like [Gibson and Sun \(2023\)](#) and [Lee and Yang \(2022\)](#) document labor market costs of lockdowns, while [Ke and Hsiao \(2022\)](#) show that although strict NPIs incur short-term economic losses, those effects may be temporary.

By 2021, as vaccines became widely available, cities such as New York City (NYC), New Orleans (NOLA), Los Angeles (LA), and San Francisco (SF) turned to vaccine mandates as a less restrictive (but still controversial) tool for curbing transmission.¹ Unlike lockdowns, mandates aimed to sustain economic activity in industry while reducing health risks in

¹Henceforth, I refer to these simply as the treated units

high-contact settings. Nevertheless, critics argued that requiring proof of vaccination for indoor dining and entertainment could slow employment recovery in restaurants by alienating customers or exacerbating staffing shortages ([Dive, 2021](#)). Supporters countered that mandates would improve public safety and restore consumer confidence, potentially boosting demand and employment ([Jones, 2021a](#)).

Each city's mandate differed in timing and design. NYC's *Key to NYC* program, announced in August 2021 and enforced beginning September, was the first large-scale vaccine mandate for indoor commerce in the United States ([Fortney, 2021](#)). San Francisco followed with an even stricter rule requiring full vaccination for customers and employees, effective later that same month ([Duffett, 2021](#)). New Orleans adopted a hybrid approach, permitting either proof of vaccination or a recent negative test ([Lorell, 2021](#)), while Los Angeles (by far the largest jurisdiction) implemented its policy in November ([Reyes, 2021](#)). These mandates created a small but consequential set of urban outliers in the national policy landscape, with important implications for labor markets in the leisure and hospitality sector. Precisely, the estimated effects of city-level vaccine mandates on restaurant employment suggest several potential lessons for policymakers and business stakeholders. If the mandates support employment growth, policymakers could view them as a viable tool for balancing public health and economic activity. Cities considering similar interventions might adopt comparable policies with confidence that labor markets will not be severely disrupted. Businesses could plan for future public health policies with reduced concern for employment losses in high-contact sectors. If the mandates depress employment, policymakers should weigh the trade-offs more

carefully before implementing mandates, especially in sectors sensitive to consumer demand and staffing shortages. Alternative interventions, such as targeted incentives for vaccination, improved workplace safety measures, or temporary subsidies, might be preferable to blunt employment impacts of mandates. Businesses may need mitigation strategies, including flexible staffing or remote services, to reduce vulnerability to future mandates. If the effects are heterogeneous across cities, differences in local conditions such as compliance rates, labor market tightness, or enforcement intensity should guide tailored policy decisions over blanket mandates. Comparative case studies could help identify which contextual factors drive positive versus negative outcomes.


This paper evaluates the causal effects of these city-level vaccine mandates on restaurant employment. While prior research has analyzed mandates' impacts on vaccination rates and health outcomes ([Cohn et al., 2022](#); [Acharya and Dhakal, 2021](#); [Lang et al., 2023](#)), relatively little is known about their labor market consequences, especially in industries dependent on face-to-face interaction. This study helps fill that gap. Because randomized experiments are infeasible, the analysis adopts a synthetic control approach to estimate the counterfactual trajectory of restaurant employment absent the mandates. Using MSA-level monthly data from the Federal Reserve Bank of St. Louis, I construct synthetic controls for each mandate city's year-over-year growth in full-service restaurant employment. The donor pool consists of metropolitan areas that did not implement similar mandates during the same period.²

The rest of the paper is organized as follows. Section [3.2](#) provides background on the evolution

²For example, 'SMU11000007072251101SA' (Washington, DC) or 'SMU08197407072251101SA' (Denver MSA).

of pandemic interventions and the city-level vaccine mandates. Section 3.3 describes the employment data and outcome construction. Section 3.4 outlines the empirical strategy, including the synthetic control methodology. Section 3.5 presents the main findings and robustness checks. Section 3.6 interprets the results and discusses limitations. Section 3.7 concludes.

3.2 Policy Background



Timeline of major pandemic policy phases

3.2.1 NPIs Before Vaccines

NPIs had aimed to reduce transmission by limiting contact or promoting respiratory barriers. The first and most famous of these interventions was the 2020 lockdown in Wuhan, China, where local officials went as far as to forcibly seal people in their homes to reduce the possibility to spread the virus; similar actions were taken in Europe and India. Even in the United States, states that could effectively seal their borders, such as Hawaii, did so,

with Hawaii mandating two-week quarantines to visitors, which effectively amounted to a tourism ban. Mask mandates were widely adopted as well due to their low cost and direct mechanism of action ([Hansen and Mano, 2023](#); [Chernozhukov et al., 2021](#); [Mitze et al., 2020](#)). While not as severe as lockdowns, the idea of mask mandates was to allow for some freedom of movement while minimizing the chance to spread the virus further. While both sets of policies were effective, they still provoked political backlash due to their perceived disruptive nature (especially in the United States).

3.2.2 The First Vaccine Policies

The arrival of vaccines in late 2020 ushered in a new phase of pandemic governance, giving jurisdictions a wider array of policy options to employ. Now, beyond lockdowns, governments finally had pharmaceutical interventions in their toolkit. Unlike non-pharmaceutical interventions, which sought to reduce transmission primarily by altering citizens behavior and limiting contact (and therefore the spread of the virus), vaccine-based interventions operate directly on biological susceptibility. NPIs such as lockdowns, mask mandates, and border closures work by creating external barriers to viral spread, often imposing economic or social costs in the process. Vaccine mandates, by contrast, rely on an individual's immune response to prevent infection or reduce severity, aiming to achieve population-level protection without requiring broad restrictions on mobility or commerce. In principle, this allows policymakers to target public health objectives more directly, potentially mitigating the trade-offs between controlling the virus and maintaining economic activity. Nevertheless, vaccine-based mandates introduce new challenges, including compliance, equity, and political acceptability, which

differ in character from those associated with NPIs.

For the first few months, it worked: when vaccines were released, initial uptake was high, particularly among older adults and (often mandated for) healthcare workers ([Okpani et al., 2024](#)). However, as eligibility expanded in early 2021, demand slowed amid misinformation, mistrust, and logistical barriers. In response, policymakers deployed incentive-based strategies to encourage vaccination ([Barber and West, 2022](#); [Robertson et al., 2021](#)). These included direct cash payments, vaccine lotteries (e.g., Ohio’s “Vax-a-Million” from May to June ([Fuller et al., 2022](#))), and smaller symbolic giveaways. One paper dubbed this the era of “donuts, drugs, booze, and guns,” reflecting the extraordinary lengths to which governments went to boost uptake through rewards while stopping short of requirements ([Tinari and Riva, 2021](#)). Despite their creativity, most incentives had modest or short-lived effects ([Saban et al., 2021](#); [Walkey et al., 2021](#); [Khazanov et al., 2023](#)). Many were poorly targeted, administratively complex, or insufficient to overcome deeper resistance. The emergence of the highly transmissible Delta variant in mid-2021 raised the stakes. Initially identified in India, Delta reached the United States by May and accounted for the majority of new cases by June. Unlike in 2020, officials now had vaccines at scale, but not enough vaccinated people. Given a new more contagious strain, soft mandates emerged, targeting workers in high-risk occupations such as healthcare and education. These policies stopped short of general mandates while making vaccination a condition of continued employment ([American Medical Association et al., 2021](#)). Over the summer, local governments and organizations expanded this approach ([American Medical Association et al., 2021](#); [Al Jazeera, 2021](#)).

3.2.3 City-Level Vaccine Mandates

The summer 2021 Delta surge spurred several major U.S. cities to enact vaccine verification requirements for indoor restaurants and leisure venues. Each city's approach differed slightly in timing and scope, but all shared the common goal of linking access to indoor leisure activities with proof of vaccination. Together, the treated units formed the clearest urban experiments in the use of vaccine mandates to manage both public health risks and consumer confidence in the hospitality sector.

New York City's "Key to NYC" program was the first such mandate in the United States. Announced on August 3, 2021, formalized on August 16, and enforced beginning September 13, it required proof of at least one vaccine dose to enter indoor dining, fitness centers, and entertainment venues ([Benveniste, 2021](#)). The rule applied not only to patrons but also to employees of covered establishments. This dual requirement meant that the policy had the potential to affect restaurant employment through both the supply side, if workers resisted vaccination or exited the industry, and the demand side, if consumer traffic shifted in response to the mandate. The city threatened fines for noncompliance and required businesses to post signage and maintain written protocols, though the intensity of inspections varied. Large chains tended to comply quickly ([Messenger, 2021](#)), while smaller establishments voiced concern about administrative burdens and customer pushback. Several lawsuits challenged the mandate, but city officials argued that rising vaccination rates demonstrated its effectiveness ([Wiener-Bronner, 2021](#)). San Francisco followed closely, enacting its mandate on August 20, 2021. Unlike New York, the city required full vaccination beyond just one dose, and did not

allow a testing alternative ([San Francisco, 2021](#); [Duffett, 2021](#)). This made San Francisco's policy the strictest among the four cities. The requirement extended to both customers and employees in indoor restaurants, bars, and entertainment venues. Enforcement included spot checks and the threat of fines, with compliance reported to be widespread among larger businesses in the city's dense urban core. As in New York, officials promoted the policy as a tool for stabilizing consumer confidence in public spaces, even as smaller businesses expressed anxiety about turning customers away.

New Orleans became the first major city in the South to adopt a mandate, announcing its policy on August 16, 2021. The city's approach was a hybrid model: individuals could gain entry to indoor restaurants and venues either by showing proof of at least one vaccine dose or by presenting a recent negative COVID test ([Aya and Riess, 2021](#)). This flexibility reflected local political and cultural dynamics but also introduced more administrative complexity for restaurants. Enforcement began within weeks ([Just, 2021a,b](#)), with inspections carried out by city officials. Restaurant owners reported mixed reactions, with some noting relief that the policy might make indoor dining safer, while others expressed concern about deterring unvaccinated tourists and residents ([Chavez, 2021](#)). Los Angeles entered the field later, passing its ordinance in October 2021 with enforcement beginning November 4. The policy was sweeping in scope, covering indoor restaurants, gyms, theaters, salons, and other leisure establishments. By the time it came into effect, public debate had already been shaped by the experience of New York and San Francisco, and Los Angeles became the largest jurisdiction in the nation to implement a vaccine passport ([Solis and Branson-Potts, 2021](#)).

Although the four cities differed in timing and strictness, they shared important common features. Each policy was applied citywide, targeted indoor leisure and hospitality settings, and created new compliance responsibilities for businesses. They also diverged sharply from the policy environments of states such as Texas and Florida, which prohibited vaccine requirements altogether. The result was a small cluster of urban outliers using vaccine mandates as a central public health intervention in 2021. These institutional details carry important implications for research design. Because the mandates covered both patrons and employees in most cases, they may have constrained labor supply by excluding some workers while simultaneously influencing labor demand by shaping customer traffic. Enforcement intensity and public perception varied, but each policy represented a reasonably sharp and well-documented intervention, announced and enforced on a specific timeline. This combination makes them suitable for empirical evaluation. Comparing employment outcomes in these cities against synthetic control groups of metropolitan areas without mandates offers a credible way to estimate causal effects on restaurant employment. The staggered adoption across cities and the variation in stringency further enrich the research design, providing multiple natural experiments within a single policy family.

3.2.4 The First Vaccine Mandates

Per the above, vaccine mandates could influence employment in full-service restaurants through both supply- and demand-side channels, with theoretically ambiguous net effects. On the supply side, some workers may exit the labor force or shift sectors if unwilling to comply, especially in lower-wage service jobs where vaccine hesitancy may be higher.

However, others may comply with the mandate precisely *because* they wish to keep their jobs, choosing continued employment over personal objection. That is, the employment effect may hinge less on individual preferences in isolation and more on the strength of the economic incentives at stake. Employers, meanwhile, may face frictional costs from enforcing mandates or accommodating staffing turnover. On the demand side, mandates may increase consumer confidence by signaling a safer environment, thereby encouraging more patrons to return to indoor dining. For a sector like full-service restaurants, where both interpersonal contact and discretionary spending are central, mandates could plausibly stimulate demand even as they introduce short-run supply constraints. Whether employment rises or falls in net depends on which margin dominates. In short, this analysis seeks to quantify the net effect of the mandate on restaurant employment, capturing the balance between these potentially opposing forces.

3.3 Dataset

We use monthly MSA-level employment series from the Federal Reserve Bank of St. Louis (FRED) and, by extension, the Bureau of Labor Statistics, spanning January 1991 through December 2022, yielding $T = 384$ monthly observations per unit. Let $\mathcal{J} = \{1, \dots, J\}$ denote the set of treated units, corresponding to New York City, San Francisco, and Los Angeles, so $J = 3$. New Orleans adopted a mandate on the same timeline, but its full-service restaurant employment series is not available for the full sample, so it is dropped from the analysis and appears in neither the treated set nor the donor pool. For treated unit $j \in \mathcal{J}$, the

pre-treatment period is defined as

$$\mathcal{T}_1^{(j)} = \{t \in \mathbb{N} : t \leq T_0^{(j)}\},$$

where $T_0^{(j)}$ is the final period before the mandate in city j . Similarly, the post-treatment period is

$$\mathcal{T}_2^{(j)} = \{t \in \mathbb{N} : t > T_0^{(j)}\}.$$

The outcome is full-service restaurant employment for the treated cities, benchmarked against a donor pool of metropolitan areas that did not enact comparable mandates over the same period.

Let e_{jt} denote employment in unit j at month t , and define the year-over-year growth rate in percentage points as

$$y_{jt} = 100 \cdot \frac{e_{jt} - e_{j,t-12}}{e_{j,t-12}}, \quad j = 1, \dots, N, \quad t = 13, \dots, T,$$

using smoothed and seasonally adjusted values. Outcomes for each unit are collected into vectors $\mathbf{y}_j = (y_{j1}, \dots, y_{jT})^\top$.

For each treated city $j \in \mathcal{J}$, a separate synthetic control is estimated using a donor pool

$$\mathcal{N}_0^{(j)} = \mathcal{N} \setminus (\{j\} \cup \mathcal{J} \setminus \{j\}),$$

which excludes the treated unit itself as well as the other treated units. The treatment indicator is now

$$d_{jt} = \mathbf{1}\{j \in \mathcal{J} \wedge t > T_0^{(j)}\},$$

so observed outcomes satisfy

$$y_{jt} = (1 - d_{jt})y_{jt}(0) + d_{jt}y_{jt}(1),$$

with causal effects

$$\tau_{jt} = y_{jt}(1) - y_{jt}(0).$$

The post-treatment average treatment effect on the treated (ATT) is

$$\text{ATT} = \frac{1}{J} \sum_{j \in \mathcal{J}} \frac{1}{|\mathcal{T}_2^{(j)}|} \sum_{t \in \mathcal{T}_2^{(j)}} \tau_{jt},$$

averaging first over the post-treatment months within each city and then across all treated cities.

3.4 Econometric Methodology

The four city-level mandates form a staggered-adoption panel with three treated units (after dropping New Orleans for data availability) and a donor pool of metropolitan areas that did not enact comparable indoor vaccine-verification policies. Treatment timing differs across cities, and the cities themselves (large, coastal, tourism-dependent metropolitan economies) often lie near or outside the convex hull of the donor pool. Classical synthetic control ([Abadie et al., 2010](#)), which fits a separate convex combination of donors to each treated unit, fails on both counts: it tolerates only one adoption date at a time and produces poor pre-treatment fit whenever the treated unit is geometrically extreme. I therefore deploy two estimators specifically designed for staggered adoption with potentially poor unit-level fit: the Synthetic Difference-in-Differences (SDID) estimator of [Arkhangelsky et al. \(2021\)](#), with the event-study disaggregation of [Ciccio \(2024\)](#), and the partially pooled synthetic control (PPSCM) estimator of [Ben-Michael et al. \(2021b\)](#). The two methods share the synthetic-control philosophy of fitting a counterfactual by data-driven weighting but differ in what they weight and how they trade off competing imbalance criteria; reporting both lets the estimated effect be triangulated rather than identified by a single modeling choice.

3.4.1 *Synthetic Difference-in-Differences*

SDID generalizes the canonical SCM idea by introducing time weights alongside unit weights and embedding both in a weighted two-way fixed-effects regression. For a given cohort a of treated units (here, one cohort per treated city), unit weights $\hat{\omega}^{sdid}$ are chosen on the

simplex to make the donor pool’s pre-treatment trajectory parallel to that of the treated unit, and time weights $\hat{\lambda}^{sdid}$ are chosen on the simplex to make the donor pool’s pre- and post-treatment average outcomes agree apart from a constant. Concretely, $\hat{\omega}^{sdid}$ solves

$$\hat{\omega}_0, \hat{\omega}^{sdid} = \underset{\omega_0 \in \mathbb{R}, \omega \in \Omega}{\operatorname{argmin}} \sum_{t=1}^{T_0} \left(\omega_0 + \sum_{i=1}^{N_{co}} \omega_i Y_{it} - \frac{1}{N_{tr}} \sum_{i \in \mathcal{I}^a} Y_{it} \right)^2 + \zeta^2 T_0 \|\omega\|_2^2,$$

with the simplex constraint $\omega \geq \mathbf{0}$, $\mathbf{1}^\top \omega = 1$, and ζ a ridge term calibrated to the standard deviation of first-differenced donor outcomes ([Arkhangelsky et al., 2021](#)). The time-weight problem is analogous, with no ridge penalty. The cohort-specific SDID estimator is then a weighted double-difference:

$$\hat{\tau}_a^{sdid} = \frac{1}{T_{tr}^a} \sum_{t=a}^T \left(\frac{1}{N_{tr}^a} \sum_{i \in \mathcal{I}^a} Y_{it} - \sum_{i=1}^{N_{co}} \hat{\omega}_i Y_{it} \right) - \sum_{t=1}^{a-1} \hat{\lambda}_t \left(\frac{1}{N_{tr}^a} \sum_{i \in \mathcal{I}^a} Y_{it} - \sum_{i=1}^{N_{co}} \hat{\omega}_i Y_{it} \right),$$

i.e., the post-treatment gap between treated and synthetic control minus a time-weighted baseline computed on the pre-treatment window. Cohort-specific effects are aggregated into an overall ATT via the treated-unit-weighted scheme of [Clarke et al. \(2023\)](#), equivalent to the event-study aggregation of [Ciccia \(2024\)](#) when applied with uniform horizon weights.

The estimator’s appeal in this setting is twofold. First, the time weights $\hat{\lambda}^{sdid}$ implicitly down-weight pre-treatment months whose outcome dynamics look unlike the post-treatment window. In a panel that straddles the COVID-induced labor-market trough of 2020, this is non-trivial: months in which restaurant employment fell catastrophically below trend are mechanically down-weighted relative to months in which year-over-year growth was

more typical, so the synthetic control is not anchored to a counterfactual dominated by the worst of the pandemic. Second, the two-way fixed effects absorb common time-invariant unit shifts and common contemporaneous shocks, weakening the parallel-trends assumption that DiD strategies usually require. Inference is conducted by placebo permutation, in which each donor MSA is iteratively reassigned as a pseudo-treated unit, SDID is re-fit, and the resulting ATT distribution is used to form a standard error and a two-sided p-value $p = (\#\{|\hat{\tau}^*| \geq |\hat{\tau}|\} + 1)/(B + 1)$ with $B = 500$ placebo replications.

3.4.2 Partially Pooled Synthetic Control

Partially pooled SCM (Ben-Michael et al., 2021b) addresses the same staggered-adoption problem from a different angle. Where SDID augments SCM with time weights and fixed effects, PPSCM generalizes SCM by jointly estimating a weight matrix $\mathbf{\Gamma} \in \mathbb{R}^{N_0 \times J}$ across all treated units and balancing two distinct pre-treatment imbalance criteria. Let $\mathbf{y}_j^{\mathcal{T}_1}$ denote treated unit j 's pre-treatment outcome vector and $\mathbf{Y}_0^{\mathcal{T}_1}$ the aligned donor matrix. Define

$$q_{\text{sep}}(\mathbf{\Gamma})^2 = \frac{1}{J} \sum_{j=1}^J \frac{1}{L} \left\| \mathbf{y}_j^{\mathcal{T}_1} - \mathbf{Y}_0^{\mathcal{T}_1} \boldsymbol{\gamma}_j \right\|_2^2,$$

$$q_{\text{pool}}(\mathbf{\Gamma})^2 = \frac{1}{L} \left\| \frac{1}{J} \sum_{j=1}^J \left(\mathbf{y}_j^{\mathcal{T}_1} - \mathbf{Y}_0^{\mathcal{T}_1} \boldsymbol{\gamma}_j \right) \right\|_2^2,$$

where the first quantity averages pre-treatment fit error across treated units (penalizing weak unit-level fits) and the second averages residuals first and then computes error (penalizing weak fit to the average treated unit). The two measures are not redundant: $q_{\text{pool}} \leq q_{\text{sep}}$ always, and the gap is large when treatment timing or unit characteristics vary across the

treated group. PPSCM solves a normalized convex combination,

$$\hat{\Gamma} = \underset{\Gamma \in \Delta_{\text{scm}}^J}{\operatorname{argmin}} \nu \cdot \tilde{q}_{\text{pool}}(\Gamma)^2 + (1 - \nu) \cdot \tilde{q}_{\text{sep}}(\Gamma)^2 + \lambda \|\Gamma\|_F^2,$$

with $\tilde{q}_{\text{sep}}$ and $\tilde{q}_{\text{pool}}$ normalized by their values at the $\nu = 0$ (separate-SCM) baseline. The pooling intensity $\nu \in [0, 1]$ governs the trade-off: $\nu = 0$ recovers separate SCM on each treated city followed by averaging, $\nu = 1$ recovers fully pooled SCM in which the treated cities are averaged first and a single synthetic control is fit to the average, and intermediate ν moves smoothly along the convex frontier of the two imbalances. I select ν data-adaptively as the value that minimizes the sum $\tilde{q}_{\text{sep}}(\Gamma_\nu) + \tilde{q}_{\text{pool}}(\Gamma_\nu)$ over a 21-point grid on $[0, 1]$, the balance-frontier knee that [Ben-Michael et al. \(2021b\)](#) recommend in the absence of error-bound parameters. The Frobenius-norm penalty $\lambda \|\Gamma\|_F^2$ stabilizes the weights and is set to a small numerical value to break degeneracies without materially shifting the fit. I report results both with and without the paper’s “intercept shift” extension, which subtracts each unit’s pre-treatment mean before fitting; the toggle has minimal effect on the headline estimate, which I treat as a passed robustness check on the demeaning choice.

Inference for PPSCM uses leave-one-treated-unit-out jackknife: each treated city is dropped in turn, the estimator is refit on the remaining $J - 1$ cities, and the sample standard deviation of the leave-one-out ATTs serves as the standard error. With $J = 3$ treated cities, the jackknife uses three fits and is therefore structurally underpowered relative to SDID’s 500 placebo replications, a point I return to in the discussion.

3.4.3 Triangulation Strategy

The two estimators are complementary; each answers a different question. SDID relaxes parallel trends via time weights and two-way fixed effects; PPSCM relaxes the convex-hull restriction of classical SCM via controlled information-sharing across treated units. They make distinct identifying assumptions and weight the panel differently in finite samples. Reporting both lets sign, magnitude, and inferential strength be cross-checked rather than reading off a single specification. All estimations use year-over-year growth rates (in percentage points) of full-service restaurant employment to remove seasonality and long-run trends, and the donor pool is restricted to MSAs that did not enact comparable indoor vaccine-verification mandates over the sample window. Both estimators are implemented in the `mlsynth` package for Python.

3.5 Results

Table~3.1 reports the pooled ATT estimates across the three treated cities (New York City, San Francisco, and Los Angeles), in percentage points of year-over-year growth in full-service restaurant employment. Both estimators are positive and agree on sign; every number in this section is computed from the current fit.

Table 3.1: Estimated effect of city-level vaccine mandates on year-over-year full-service restaurant employment growth, in percentage points. $J = 3$ treated CSAs (New York City, San Francisco, Los Angeles); SDID inference by $B = 500$ placebo replications, PPSCM by leave-one-treated-out jackknife over the three cities.

Estimator	ATT (pp)	SE	95% CI	Inference
SDID (pooled)	+5.8	0.87	[+4.1, +7.5]	Placebo ($p = 0.002$)
PPSCM (pooled)	+3.7	4.17	[−4.5, +11.8]	Jackknife

SDID puts the pooled effect at +5.8 percentage points, with a placebo standard error of 0.87, a 95% confidence interval of $[+4.1, +7.5]$, and a placebo p-value of 0.002. PPSCM puts it at +3.7 percentage points; its leave-one-treated-out jackknife rests on three cities, so it is wide, with a standard error of 4.17 and a 95% confidence interval of $[-4.5, +11.8]$ that includes zero.

The partially pooled model also yields an effect for each city. These are +10.8 percentage points for New York City, +1.5 for San Francisco, and -1.3 for Los Angeles, so the pooled positive is carried by New York City: dropping it from the pool sends the pooled PPSCM estimate to -0.5 points, which is why the jackknife interval is so wide. Figure 3.1 shows each city's observed path against its synthetic control, with the synthetic tracking the city through the pre-mandate period.

Two readings follow. First, both methods are positive and agree on direction despite their different identification strategies, and no specification turns the effect negative. The gap between them, about 2.2 percentage points, is consistent with the way the two methods treat pre-treatment heterogeneity: SDID's time weights $\hat{\lambda}$ down-weight pre-treatment months in which donor and treated outcomes diverged most (including the depths of the 2020 trough), while PPSCM weights all pre-treatment months equally within the chosen pre-window. Second, inference rests on SDID. With three treated units the PPSCM jackknife has a three-observation variance and cannot separate its estimate from zero, while SDID's placebo distribution uses the full donor pool and delivers the 0.002 p-value. The channel feared by critics, mandates driving employment down in proof-of-vaccination indoor dining, appears in neither estimate.




Figure 3.1: Per-city synthetic-control fit: observed (black) and synthetic (dashed) year-over-year full-service restaurant employment growth for each treated city, with the mandate month marked (dotted). The synthetic tracks each city through the pre-mandate period; these panels are single-city SDID fits, and the partially pooled per-city effects are reported in the text.

Figure 3.2 visualizes the pooled estimate in Table~?: the synthetic difference-in-differences fit across the three treated cities under staggered adoption (New York City and San Francisco from August 2021, Los Angeles from November 2021).

3.6 Discussion

The substantive finding of this paper is that city-level vaccine mandates did not appreciably damage local restaurant labor markets, and may have modestly improved them. The two estimators agree on sign, with pooled point estimates of +3.7 percentage points (PPSCM) and +5.8 percentage points (SDID) of year-over-year employment growth. Read against the policy debate at the time of adoption, the implication is clear: the channel feared by critics (mandates alienating customers and aggravating staffing shortages, with consequent employment losses in proof-of-vaccination-restricted indoor dining) is not present in the data. If anything, the data are more consistent with the opposite channel of Jones (2021b), in




Figure 3.2: Pooled synthetic difference-in-differences fit across the three treated cities (New York City, San Francisco, Los Angeles) under staggered adoption. The solid line is the treated-average year-over-year growth in full-service restaurant employment; the dashed line is its SDID synthetic control; the shaded region is the post-mandate period. The pooled SDID ATT is +5.8 and the PPSCM ATT is +3.7 percentage points, both positive and in the range of the estimates in Table~??.

which the mandates restored consumer confidence in indoor commerce sufficiently to support faster employment recovery in the treated cities than in otherwise-similar comparison MSAs.

Whether the effect is best read as “slightly positive” or “no harm done” depends on which inferential signal one finds more credible. SDID’s tightly identified +5.8 percentage-point estimate, with a placebo p-value of 0.002, supports the slightly-positive reading. PPSCM’s +3.7 estimate with a confidence interval that includes zero supports the no-harm reading. Both are policy-relevant conclusions, and both are inconsistent with the strong negative hypothesis that originally motivated the analysis. The triangulation across methods is not just rhetorical: SDID and PPSCM make distinct identifying assumptions (two-way fixed effects with time re-weighting versus pooled donor weights with no time component), and the fact that they agree on direction makes the result more robust to misspecification of either method individually than would a single estimator with no triangulation.

Several caveats temper the strength of the conclusion. First, the small number of treated cities (three after dropping New Orleans for data availability) limits statistical power for cohort-level inference and forces reliance on the donor pool for variance estimation. Replication on the full four-city panel with New Orleans included, or on additional jurisdictions that subsequently adopted similar policies, would tighten both methods’ inference. Second, both estimators assume that the donor MSAs did not experience contemporaneous shocks that mimicked the treatment effect. To the extent that the broader 2021 economic recovery generated correlated labor-market dynamics across cities, the imputed counterfactual may absorb some of this variation, so it may not fully isolate the mandate’s causal effect. Third, the outcome

studied (year-over-year employment growth in full-service restaurants) captures one margin of adjustment but not others. Hours worked, wages, vacancy rates, and the composition of employment between part-time and full-time positions are not estimated here, and shifts along those margins could produce employment-growth effects whose sign masks a more complicated story about labor-market quality.

A natural follow-up question that this paper does not yet answer is how the aggregate effect decomposes across the three treated cities. Both estimators produce city-by-city estimates as a by-product of the aggregate fit, and inspecting whether the positive overall effect is driven by one city or distributed across all three would sharpen the policy interpretation. I defer that decomposition to future work.

3.7 Conclusion

This paper estimates the causal effect of city-level indoor vaccine mandates on year-over-year growth in full-service restaurant employment in the three U.S. metropolitan areas that adopted such policies during the 2021 Delta variant wave. Using synthetic difference-in-differences and partially pooled synthetic control (two staggered-adoption estimators with distinct identification strategies), I find no evidence that the mandates depressed restaurant employment growth, and some evidence consistent with a modest positive effect on the order of two to three-and-a-half percentage points. The SDID estimate is statistically significant at conventional levels, while the PPSCM estimate is borderline non-significant given the structurally limited degrees of freedom available for jackknife inference over three treated

cities; the two methods agree on sign and overlap substantially on magnitude.

For policymakers, the result has a clear takeaway. The most common objection to indoor vaccine mandates (that they would suppress employment in the leisure and hospitality sector by deterring patrons or aggravating staffing shortages) is not borne out in the labor-market data from the cities that actually implemented such policies. Whether the public-health benefits of the mandates were sufficient to justify their adoption is a separate question that this paper does not address, but the labor-market cost half of the cost-benefit calculation appears, at the very least, to be small.

REFERENCES

- Abadie, A., Diamond, A., and Hainmueller, J. (2010). Synthetic control methods for comparative case studies: Estimating the effect of california’s tobacco control program. *Journal of the American Statistical Association*, 105(490):493–505.
- Acharya, B. and Dhakal, C. (2021). Implementation of state vaccine incentive lottery programs and uptake of covid-19 vaccinations in the united states. *JAMA Network Open*, 4(12):e2138238–e2138238.
- Al Jazeera (2021). Covid-19 delta surge in us leads to new restrictions, jab push. Officials are growing more worried about the current rise in infections and are blaming the unvaccinated for the spike.
- American Medical Association et al. (2021). Major health care professional organizations call for covid-19 vaccine mandates for all health workers. <https://tinyurl.com/24r6jjtv>. Press Release, July 26, 2021.
- Amjad, M., Shah, D., and Shen, D. (2018). Robust synthetic control. *The Journal of Machine Learning Research*, 19(1):802–852.
- Arkhangelsky, D., Athey, S., Hirshberg, D. A., Imbens, G. W., and Wager, S. (2021). Synthetic difference-in-differences. *American Economic Review*, 111(12):4088–4118.
- Aya, E. and Riess, R. (2021). New orleans will require proof of vaccination or a negative covid-19 test to enter certain indoor venues. Accessed: 2025-09-29.
- Bai, J. (2009). Panel data models with interactive fixed effects. *Econometrica*, 77(4):1229–

1279.

- Barber, A. and West, J. (2022). Conditional cash lotteries increase covid-19 vaccination rates. *Journal of health economics*, 81:102578.
- Batra, A. (2025). India pm2.5 satellite data processing pipeline, block level (1998–2023). <https://github.com/AarshBatra/biteSizedAQ/tree/main/13.ind.block.pm2.5.sat.data.processing.1998.2023>. GitHub repository, accessed 2025-09-05.
- Bayat, N., Morrin, C., Wang, Y., and Misra, V. (2020). Synthetic control, synthetic interventions, and covid-19 spread: Exploring the impact of lockdown measures and herd immunity.
- BBC (2022). Air quality: Pollution shortening lives by almost 10 years in delhi, says study. Accessed 2025-09-05.
- Ben-Michael, E., Feller, A., and Rothstein, J. (2021a). The augmented synthetic control method. *Journal of the American Statistical Association*, 116(536):1789–1803.
- Ben-Michael, E., Feller, A., and Rothstein, J. (2021b). The augmented synthetic control method. *Journal of the American Statistical Association*, 116(536):1789–1803.
- Benchrif, A., Wheida, A., Tahri, M., Shubbar, R. M., and Biswas, A. (2021). Air quality during three covid-19 lockdown phases: Aqi, pm2.5 and no2 assessment in cities with more than 1 million inhabitants. *Sustainable Cities and Society*, 74:103170.
- Benveniste, A. (2021). New york city will require vaccines for entry to restaurants and gyms.
- Bhuyan, A., Bordoloi, T., Debnath, R., Ikbali, A. M. A., Debnath, B., and Singh, W. S. (2025). Assessing aqi of air pollution crisis 2024 in delhi: its health risks and nationwide

- impact. *Discover Atmosphere*, 3(1):13.
- Bilgel, F. (2022). Effects of covid-19 lockdowns on social distancing in turkey. *The Econometrics Journal*, 25(3):781–805.
- Bollyky, T. J., Castro, E., Aravkin, A. Y., et al. (2023). Assessing COVID-19 pandemic policies and behaviours and their economic and educational trade-offs across US states from January 1, 2020, to July 31, 2022: an observational analysis. *The Lancet*, 401(10385):1341–1360.
- Bond-Smith, S. and Fuleky, P. (2022). The effects of the pandemic on the economy of hawaii. *Current Issues in Tourism*, 26(23):3846–3852.
- Bureau of Transportation Statistics (2024). Airline on-time performance data. U.S. Department of Transportation, Bureau of Transportation Statistics. Monthly arriving-flight counts by airport, scraped and aggregated to statewide (Florida) and airport-level (TPA, MCO, PHX) travel-demand proxies. <https://www.transtats.bts.gov/airports.asp>.
- Business Insider (2020). Phone location data shows how Florida spring breakers could have spread the coronavirus across the Eastern US. <https://www.businessinsider.com/coronavirus-florida-spring-break-location-data-spread-social-distancing-2020-3>.
- Callaway, B. and Li, T. (2023). Policy evaluation during a pandemic. *Journal of Econometrics*, 236(1):105454.
- Cameletti, M. (2020). The effect of corona virus lockdown on air pollution: Evidence from the city of brescia in lombardia region (italy). *Atmospheric Environment*, 239:117794.
- Cerqueti, R., Coppier, R., Girardi, A., and Ventura, M. (2021). The sooner the better: lives saved by the lockdown during the covid-19 outbreak. the case of italy. *The Econometrics*

Journal, 25(1):46–70.

Chavez, R. (2021). As new orleans parades out vaccine rules, restaurants wonder if they’ll make it to mardi gras.

Chen, Y.-T., Yang, J.-C., and Yang, T.-T. (2024). Synthetic historical control for policy evaluation. Available at SSRN: <https://ssrn.com/abstract=4995085>.

Chernozhukov, V., Kasahara, H., and Schrimpf, P. (2021). Causal impact of masks, policies, behavior on early covid-19 pandemic in the u.s. *Journal of Econometrics*, 220(1):23–62.

Ciccia, D. (2024). A short note on event-study synthetic difference-in-differences estimators. *arXiv preprint arXiv:2407.09565*.

Clarke, D., Pailańir, D., Athey, S., and Imbens, G. (2023). Synthetic difference in differences estimation. *arXiv preprint arXiv:2301.11859*.

Cohn, E., Chimowitz, M., Long, T., Varma, J. K., and Chokshi, D. A. (2022). The effect of a proof-of-vaccination requirement, incentive payments, and employer-based mandates on covid-19 vaccination rates in new york city: a synthetic-control analysis. *The Lancet Public Health*, 7(9):e754–e762.

Cole, M. A., Elliott, R. J., and Liu, B. (2020). The impact of the wuhan covid-19 lockdown on air pollution and health: a machine learning and augmented synthetic control approach. *Environmental and Resource Economics*, 76(4):553–580.

Conley, T. G., Hansen, C. B., and Rossi, P. E. (2012). Plausibly exogenous. *The Review of Economics and Statistics*, 94(1):260–272.

Croix, S. L. (2021). Economic history of hawai‘i.

- Cui, Y., Pu, H., Shi, X., Miao, W., and Tchetgen Tchetgen, E. J. (2024). Semiparametric proximal causal inference. *Journal of the American Statistical Association*, 119(546):1348–1359.
- Dalberg Advisors (2021). Air pollution in india and the impact on business. Accessed 2025-09-05.
- de Bont, J., Krishna, B., Stafoggia, M., Banerjee, T., Dholakia, H., Garg, A., Ingole, V., Jaganathan, S., Kloog, I., Lane, K., Mall, R. K., Mandal, S., Nori-Sarma, A., Prabhakaran, D., Rajiva, A., Tiwari, A. S., Wei, Y., Wellenius, G. A., Schwartz, J., Prabhakaran, P., and Ljungman, P. (2024). Ambient air pollution and daily mortality in ten cities of india: a causal modelling study. *The Lancet Planetary Health*, 8(7):e433–e440.
- de Chaisemartin, C. and D’Haultfoeuille, X. (2022). Two-way fixed effects and differences-in-differences with heterogeneous treatment effects: a survey. *The Econometrics Journal*, 26(3):C1–C30.
- Deshpande, R. (2022). Disaster management in india: are we fully equipped? *Journal of social and economic development*, 24(Suppl 1):242–281.
- Dive, R. (2021). San francisco restaurants blame sales declines on vaccine mandates and delta surge. Accessed: 2025-09-29.
- Doudchenko, N. and Imbens, G. W. (2016). Balancing, regression, difference-in-differences and synthetic control methods: A synthesis. *NBER Working Paper No. 22791*.
- Duffett, B. (2021). Sf officially mandates proof of vaccination for indoor dining. Accessed: 2025-09-29.

- Fan, J. and Gijbels, I. (1992). Variable bandwidth and local linear regression smoothers. *The Annals of Statistics*, 20(4):2008–2036.
- Ferman, B. and Pinto, C. (2021). Synthetic controls with imperfect pretreatment fit. *Quantitative Economics*, 12(4):1197–1221.
- Fortney, L. (2021). Nyc to require proof of vaccination to dine indoors starting august 16. Accessed: 2025-09-29.
- Friedson, A. I., McNichols, D., Sabia, J. J., and Dave, D. (2021). Shelter-in-place orders and public health: Evidence from california during the covid-19 pandemic. *Journal of Policy Analysis and Management*, 40(1):258–283.
- Fuller, S., Kazemian, S., Algara, C., and Simmons, D. J. (2022). Assessing the effectiveness of covid-19 vaccine lotteries: A cross-state synthetic control methods approach. *PLOS ONE*, 17(9):1–18.
- Ghassami, A., Shpitser, I., and Tchetgen Tchetgen, E. J. (2023). Partial identification of causal effects using proxy variables. arXiv:2304.04374.
- Gibson, J. and Sun, X. (2023). A synthetic control analysis of us state level covid-19 stay-at-home orders on new jobless claims. *Journal of Economics and Finance*, 47(1):1–14.
- Goel, A., Saxena, P., Sonwani, S., Rathi, S., Srivastava, A., Bharti, A. K., Jain, S., Singh, S., Shukla, A., and Srivastava, A. (2021). Health benefits due to reduction in respirable particulates during covid-19 lockdown in india. *Aerosol and Air Quality Research*, 21(5):200460.
- Gong, D., Shang, Z., Su, Y., Yan, A., and Zhang, Q. (2024). Economic impacts of china’s zero-covid policies. *China Economic Review*, 83:102101.

- Gupta, G. and Damani, S. (2021). Solving india’s air pollution can boost economy and business. here’s how. Accessed 2025-09-05.
- Hansen, N.-J. H. and Mano, R. C. (2023). Mask mandates save lives. *Journal of Health Economics*, 88:102721.
- Hawai’i Department of Business, Economic Development and Tourism (2025). Job count by industry (CES)—seasonally adjusted monthly data. Statewide Current Employment Statistics, Research and Economic Analysis Division. <https://dbedt.hawaii.gov/economic/labor/job-count-by-industry/>.
- Hawai’i Tourism Authority (2019). 2019 annual visitor research report: Statistics on hawai’i’s visitor industry. Prepared by Jennifer Chun, Minh-Chau Chun, Lawrence Liu, Ariana Kwan, and Joseph Patoskie; STR Inc. provided hotel occupancy and room rate statistics. Cover photo by Tor Johnson. Accessed via https://t.ly/a_r5Y.
- He, C., Hong, S., Mu, H., Yuan, Y., Chen, Y., Zhan, J., Jia, H., Chen, Z., Hao, X., Li, R., Cai, B., and Mu, Y. (2020). Diurnal and temporal changes in air pollution during covid-19 strict lockdown over different regions of india. *Environmental Pollution*, 266(3):115368.
- Heffernan, C., Koehler, K., Zamora, M. L., Buehler, C., Gentner, D. R., Peng, R. D., and Datta, A. (2025). A causal machine-learning framework for studying policy impact on air pollution: a case study in covid-19 lockdowns. *American Journal of Epidemiology*, 194(1):185–194.
- HNN Staff (2020). Sweeping ‘stay-at-home’ order in effect statewide in push to stem spread of coronavirus. Published March 25, 2020.

- Hollingsworth, A. and Wing, C. (2022). Tactics for design and inference in synthetic control studies: An applied example using high-dimensional data. Working paper.
- Hu, Z., Jin, Q., Ma, Y., Ji, Z., Zhu, X., and Dong, W. (2022). How does covid-19 lockdown impact air quality in india? *Remote Sensing*, 14(8).
- India State-Level Disease Burden Initiative Air Pollution Collaborators (2020). Health and economic impact of air pollution in the states of india: the global burden of disease study 2019. *The Lancet Planetary Health*, 5(1):e25–e38.
- Jaganathan, S., Stafoggia, M., Rajiva, A., Mandal, S., Dixit, S., de Bont, J., Wellenius, G. A., Lane, K. J., Nori-Sarma, A., Kloog, I., Prabhakaran, D., Prabhakaran, P., Schwartz, J., and Ljungman, P. (2024). Estimating the effect of annual pm₂₅ exposure on mortality in india: a difference-in-differences approach. *The Lancet Planetary Health*, 8(12):e987–e996.
- Jones, A. (2021a). Restaurants need workers. would an employee vaccine mandate bring them back? Accessed: 2025-09-29.
- Jones, A. (2021b). Restaurants Need Workers. Would an Employee Vaccine Mandate Bring Them Back? North Coast Journal. May 12, 2021. Accessed: 2025-09-29.
- Joshi, D. C., Negi, P., Devi, S., Lohani, H., Kumar, R., Gupta, M., and Ming, L. C. (2025). Fine particulate matter (pm_{2.5}, pm₁₀): A silent catalyst for chronic lung diseases in india; a comprehensive review. *Environmental Challenges*, 20:101215.
- Just, A. (2021a). Saints fans at superdome must show covid vaccine proof; here’s how it will work.
- Just, A. (2021b). Thousands sign up for la wallet app after new orleans announces vaccine

mandate.

- Kawano, A., Kelp, M., Qiu, M., Singh, K., Chaturvedi, E., Dahiya, S., Azevedo, I., and Burke, M. (2025). Improved daily pm2.5 estimates in india reveal inequalities in recent enhancement of air quality. *Science Advances*, 11(4):eadq1071.
- Ke, X. and Hsiao, C. (2022). Economic impact of the most drastic lockdown during covid-19 pandemic—the experience of hubei, china. *Journal of Applied Econometrics*, 37(1):187–209.
- Ke, X. and Hsiao, C. (2023). Data subject to multiple treatment effects — disentangle the impacts of global pandemic and a specific disease control policy. *The Singapore Economic Review*, 68(05):1507–1527.
- Khazanov, G. K., Stewart, R., Pieri, M. F., Huang, C., Robertson, C. T., Schaefer, K. A., Ko, H., and Fishman, J. (2023). The effectiveness of financial incentives for covid-19 vaccination: A systematic review. *Preventive Medicine*, 172:107538.
- Kumari, P. and Toshniwal, D. (2020). Impact of lockdown measures during covid-19 on air quality— a case study of india. *International Journal of Environmental Health Research*, 31(5):503–510.
- Kumari, P. and Toshniwal, D. (2021). Revisiting air quality during lockdown persuaded by second surge of covid-19 of megacity delhi, india. *Sustainable Cities and Society*, 72:103028.
- Lang, D., Esbenshade, L., and Willer, R. (2023). Did ohio’s vaccine lottery increase vaccination rates? a pre-registered, synthetic control study. *Journal of Experimental Political Science*, 10(2):242–260.
- Le Quéré, C., Jackson, R. B., Jones, M. W., Smith, A. J., Abernethy, S., Andrew, R. M.,

- De-Gol, A. J., Willis, D. R., Shan, Y., Canadell, J. G., et al. (2020). Temporary reduction in daily global co2 emissions during the covid-19 forced confinement. *Nature climate change*, 10(7):647–653.
- Lee, J. and Yang, H.-S. (2022). Pandemic and employment: evidence from covid-19 in south korea. *Journal of Asian economics*, 78:101432.
- Li, H., Li, Z., and Hsiao, C. (2023). Assessing the impacts of pandemic and the increase in minimum down payment rate on Shanghai housing prices. *Empirical Economics*.
- Li, K. and Shankar, V. (2023). Inference theory for quasi-experimental methods when data are non-stationary with unknown structure: Does the lifting of shelter-in-place keep covid-19 manageable? *Available at SSRN 4321847*.
- Lipsitch, M., Tchetgen Tchetgen, E., and Cohen, T. (2010). Negative controls: a tool for detecting confounding and bias in observational studies. *Epidemiology*, 21(3):383–388.
- Liu, A., Waldman, R. N., Deal, B., Duff, J., Batycki, J. N., Pretto, E. A., Saavedra, J., and Szapocznik, J. (2024a). Preparing for the next pandemic: Reflections and recommendations from Florida. *PLoS ONE*, 19(12):e0314570.
- Liu, J., Tchetgen Tchetgen, E., and Varjão, C. (2024b). Proximal causal inference for synthetic control with surrogates. In Dasgupta, S., Mandt, S., and Li, Y., editors, *Proceedings of The 27th International Conference on Artificial Intelligence and Statistics*, volume 238 of *Proceedings of Machine Learning Research*, pages 730–738. PMLR.
- Lorell, C. (2021). New orleans mandates proof of vaccination for indoor dining and drinking. Accessed: 2025-09-29.

- Mahato, S., Pal, S., and Ghosh, K. G. (2020). Effect of lockdown amid covid-19 pandemic on air quality of the megacity delhi, india. *Science of The Total Environment*, 730:139086.
- Maui Now Staff (2020). Hawaii governor issues emergency proclamation for covid-19. Published March 4, 2020. Updated March 6, 2020.
- Messenger, H. (2021). From mcdonald’s to goldman sachs, here are the companies mandating vaccines for all or some employees.
- Miao, W., Geng, Z., and Tchetgen Tchetgen, E. J. (2018). Identifying causal effects with proxy variables of an unmeasured confounder. *Biometrika*, 105(4):987–993.
- Milman, O. and Uteuova, A. (2025). New york city’s congestion pricing has cut pollution and traffic – but trump still wants to kill it. *The Guardian*, accessed 2025-09-05.
- Mishra, A. K., Rajput, P., Singh, A., Singh, C. K., and Mall, R. K. (2021). Effect of lockdown amid covid-19 on ambient air quality in 16 indian cities. *Frontiers in Sustainable cities*, 3:705051.
- Mitze, T., Kosfeld, R., Rode, J., and Wälde, K. (2020). Face masks considerably reduce covid-19 cases in germany. *Proceedings of the National Academy of Sciences*, 117(51):32294–32301.
- Nakaso, D. (2020). Gov. david ige calls for 30-day suspension of visitors to hawaii. Published March 17, 2020.
- Napolitano, J. M., Srikanth, S., Noorai, R. E., Wilson, S., Williams, K. E., Rosales-Garcia, R. A., et al. (2024). SARS-CoV-2 variant introduction following spring break travel and transmission mitigation strategies. *PLoS ONE*, 19:e0301225.

- Newman, R. and Noy, I. (2023). The global costs of extreme weather that are attributable to climate change. *Nature Communications*, 14(1):6103.
- Nigam, R., Pandya, K., Luis, A., and et al. (2021). Positive effects of covid-19 lockdown on air quality of industrial cities (ankleshwar and vapi) of western india. *Scientific Reports*, 11:4285.
- Okpani, A. I., Lockhart, K., Barker, S., Grant, J. M., and Yassi, A. (2024). Did the health care vaccine mandate work? an evaluation of the impact of the covid-19 vaccine mandate on vaccine uptake and infection risk in a large cohort of canadian health care workers. *American Journal of Infection Control*, 52(9):1065–1072.
- Oliu-Barton, M., Pradelski, B. S., Woloszko, N., Guetta-Jeanrenaud, L., Aghion, P., Artus, P., Fontanet, A., Martin, P., and Wolff, G. B. (2022). The effect of covid certificates on vaccine uptake, health outcomes, and the economy. *Nature Communications*, 13(1):3942.
- Osman, S. M. I., Islam, F., and Sakib, N. (2022). Economic resilience in times of public health shock: The case of the us states. *Research in Economics*, 76(4):277–289.
- Our Common Air Commission (2024). Clean air as an asset: Strengthening the economic case for action on clean air. Accessed 2025-09-05.
- Pal, S. C., Chowdhuri, I., Saha, A., Ghosh, M., Roy, P., Das, B., Chakrabortty, R., and Shit, M. (2022). Covid-19 strict lockdown impact on urban air quality and atmospheric temperature in four megacities of india. *Geoscience Frontiers*, 13(6):101368.
- Park, C. and Tchetgen Tchetgen, E. J. (2025). Single proxy synthetic control. Working paper.
- Pathakoti, M., Muppalla, A., Hazra, A., Dutta, M., Kalluri, R. O. R., Mahammad, S. D.,

- Kadiyam, S. K., Kallakuri, L. K., Morthala, V. R. R., Brirchi, Y., Verma, R. K., Madhavalatha, A., Mallikarjun, K., Dhubey, S. K., Kotalo, R. G., Jain, A. K., and Kadiyam, S. (2021). An assessment of the impact of a nation-wide lockdown on air pollution - a remote sensing perspective over india. *Atmospheric Chemistry and Physics*, 21(11):9047–9064.
- Porter, R. and McIlvaine-Newsad, H. (2022). ‘if I get Corona, I get Corona’: COVID-19 protocols and beach tourism in Pinellas county, Florida. *Leisure/Loisir*.
- Ranjan, A. K., Patra, A., and Gorai, A. (2020). Effect of lockdown due to sars covid-19 on aerosol optical depth (aod) over urban and mining regions in india. *Science of The Total Environment*, 745:141024.
- Reyes, E. A. (2021). L.a. to require proof of covid vaccination at indoor restaurants, salons, other venues. Accessed: 2025-09-29.
- Robertson, C., Schaefer, K. A., and Scheitrum, D. (2021). Are vaccine lotteries worth the money? *Economics Letters*, 209:110097.
- Roth, J., Sant’Anna, P. H., Bilinski, A., and Poe, J. (2023). What’s trending in difference-in-differences? a synthesis of the recent econometrics literature. *Journal of Econometrics*, 235(2):2218–2244.
- Saban, M., Myers, V., Ben Shetrit, S., and Wilf-Miron, R. (2021). Issues surrounding incentives and penalties for covid-19 vaccination: The israeli experience. *Preventive Medicine*, 153:106763.
- Saleem, F., Hina, S., Ullah, I., Habib, A., Hina, A., Ilyas, S., and Hamid, M. (2024). Impacts of irregular and strategic lockdown on air quality over indo-pak subcontinent: Pre-to-post

- covid-19 analysis. *Chaos, Solitons & Fractals*, 178:114255.
- San Francisco (2021). San francisco to require proof of vaccination for entry to certain indoor businesses and all large indoor events. Accessed: 2025-09-29.
- Shi, X., Li, K. Q., Yu, M., Miao, W., Kuchibhotla, A. K., Hu, M., and Tchetgen Tchetgen, E. J. (2026). Theory for identification and inference with synthetic controls: A proximal causal inference framework. *Journal of the American Statistical Association*.
- Shi, X., Miao, W., and Tchetgen Tchetgen, E. (2020). A selective review of negative control methods in epidemiology. *Current Epidemiology Reports*, 7:190–202.
- Silva, A. C. T., Branco, P. T. B. S., and Sousa, S. I. V. (2022). Impact of covid-19 pandemic on air quality: A systematic review. *International Journal of Environmental Research and Public Health*, 19(4).
- Singh, V., Singh, S., and Biswal, A. (2021). Exceedances and trends of particulate matter (pm_{2.5}) in five indian megacities. *Science of The Total Environment*, 750:141461.
- Sobiech Pellegrini, I. (2022). Untimely reopening? increase in the number of new covid-19 cases after reopening in one brazilian state. *OXford bulletin of economics and statistics*, 84(4):675–693.
- Solis, N. and Branson-Potts, H. (2021). Thousands protest covid-19 vaccination mandates as l.a.’s verification rules kick in.
- Srivastava, S., Kumar, A., Baudhh, K., Gautam, A. S., and Verma, S. (2020). 21-day lockdown in india dramatically reduced air pollution indices in lucknow and new delhi, india. *Bulletin of Environmental Contamination and Toxicology*, 105(1):9–15.

- Statista (2025). Air pollution in india - statistics & facts. Accessed 2025-09-05.
- Tchetgen Tchetgen, E. J., Ying, A., Cui, Y., Shi, X., and Miao, W. (2024). An introduction to proximal causal inference. *Statistical Science*, 39(3):375–390.
- Tinari, S. and Riva, C. (2021). Donuts, drugs, booze, and guns: what governments are offering people to take covid-19 vaccines. *BMJ*, 374.
- Venter, Z. S., Aunan, K., Chowdhury, S., and Lelieveld, J. (2020). Covid-19 lockdowns cause global air pollution declines. *Proceedings of the National Academy of Sciences*, 117(32):18984–18990.
- Walkey, A. J., Law, A., and Bosch, N. A. (2021). Lottery-based incentive in ohio and covid-19 vaccination rates. *JAMA*, 326(8):766–767.
- Wang, Y., Wen, Y., Wang, Y., Zhang, S., Zhang, K. M., Zheng, H., Xing, J., Wu, Y., and Hao, J. (2020). Four-month changes in air quality during and after the covid-19 lockdown in six megacities in china. *Environmental Science & Technology Letters*, 7(11):802–808.
- Watts, J. and Kommenda, N. (2020). Coronavirus pandemic leading to huge drop in air pollution. The Guardian, accessed 2025-09-05.
- Wiener-Bronner, D. (2021). Restaurant owners sue new york city over vaccine mandate.
- Witkowski, K., Yeo, J., Belligoni, S., Ganapati, N. E., Corbin, T., and Rivera, F. (2023). Florida as a COVID-19 epicenter: Exploring the role of institutions in the state’s response. *International Journal of Public Administration*, 46(7):484–498.
- Wu, J., Zhan, X., Xu, H., and Ma, C. (2023). The economic impacts of covid-19 and city lockdown: Early evidence from china. *Structural Change and Economic Dynamics*,

65:151–165.

- Yan, Y., Shah, M. I., Sharma, G. D., Chopra, R., Fareed, Z., and Shahzad, U. (2022). Can tourism sustain itself through the pandemic: nexus between tourism, covid-19 cases and air quality spread in the ‘pineapple state’ hawaii. *Current Issues in Tourism*, 25(3):421–440.
- Yang, X. (2021). Does city lockdown prevent the spread of covid-19? new evidence from the synthetic control method. *Global Health Research and Policy*, 6:1–14.
- Yuan, S. (2022). Zero covid in china: what next? *The Lancet*, 399(10338):1856–1857.
- Zhang, M., Katiyar, A., Zhu, S., Shen, J., Xia, M., Ma, J., Kota, S. H., Wang, P., and Zhang, H. (2020). Impact of reduced anthropogenic emissions during covid-19 on air quality in india. *Atmospheric Chemistry and Physics Discussions*, 2020:1–23.
- Zhang, X., Wang, W., and Zhang, X. (2022). Asymptotic properties of the synthetic control method. *arXiv preprint arXiv:2211.12095*.