

# Paradise Lost? Estimating the Economic Impact of Hawaii’s Mandatory Quarantine using Synthetic Historical Controls

JARED GREATHOUSE<sup>†</sup>

<sup>†</sup>*Department of Policy and Management, Georgia State University*  
E-mail: [jgreathouse3@student.gsu.edu](mailto:jgreathouse3@student.gsu.edu)

**Summary** This article estimates the causal impact of Hawaii’s March 2020 mandatory quarantine on its tourism industry and labor market using the Synthetic Historical Control method. Analyzing visitor days, hotel occupancy, accommodations employment, and the Coincident Economic Activity Index (all in annualized year-over-year growth rates), the study reveals substantial negative impacts from the quarantine. Observed outcomes fall significantly outside 90% conformal out-of-sample prediction intervals, indicating a clear treatment effect. Placebo tests further validate the model’s internal credibility. This research demonstrates the method’s utility and informs the broader debate on economic impacts of anti-COVID policy.

**Keywords:** *Synthetic Controls, Policy Evaluation, COVID-19*

## 1. INTRODUCTION

Many of the most pressing policy questions in economics concern one-time, systemic shocks: What would the United States’ economy have looked like absent the revolution? What would have happened to unemployment in Great Britain without Brexit? In these cases, credible untreated units often do not exist, making standard causal inference tools, like difference-in-differences (DiD) and synthetic control (SCM), difficult or impossible to apply.

This paper takes up one such case: Hawaii’s near-total quarantine in March 2020. While most regions experienced economic contraction from COVID-19 through voluntary distancing, fear, or uneven mandates, Hawaii implemented one of the few deliberate and centralized shutdowns in the U.S.—a policy aimed not merely at mitigating the virus through public compliance, but at actively suppressing travel and commerce. On March 26, 2020, the state imposed a 14-day quarantine on all arrivals, effectively sealing its borders. Hawaii’s economic and tourism collapse was thus, in part, an intended consequence. Within a month, a third of the state’s labor force was unemployed [Ivanova \(2020\)](#).

This paper contributes to the literature on the broader economic consequences of lockdown-style policies—beyond their direct effects on infection. Understanding the economic costs of aggressive containment measures is vital for future health policy, particularly in tourism-dependent regions where travel restrictions can have outsized effects. Hawaii’s case offers a uniquely compelling setting: its geographic isolation, centralized governance, and heavy reliance on tourism allowed it to pursue a form of near-total border control that was effectively impossible in mainland states. As such, Hawaii provides an empirical glimpse into a “zero-COVID”-like strategy implemented

within the U.S. federal system. Given how the borders were effectively closed, as in Wuhan, the strategy might be called “The Great Isolation with Aloha characteristics.”

Empirically, I find that Hawaii experienced substantial short-term economic losses across the tourism and labor market sectors, but that these losses were largely short-lived—easing by the end of 2020 and into early 2021. This paper proceeds as follows. Section~2 reviews related work on the economic effects of early COVID-19 interventions and highlights key methodological challenges in estimating their impacts. Section~3 presents the empirical strategy, introducing the Synthetic Historical Control (SHC) method and detailing its relevance for studying Hawaii’s unique shutdown. Section~4 describes the data sources used, including state-level tax revenues and sectoral employment. Section~5 reports the main findings, showing both the magnitude and duration of Hawaii’s economic losses. Section~6 interprets the results in the context of broader policy debates on early action and economic tradeoffs.

## 2. BACKGROUND

### 2.1. Previous Work

The onset of the COVID-19 pandemic spurred a global wave of non-pharmaceutical interventions (NPIs), with governments acting swiftly to mitigate viral transmission. Much of the early scholarly work emphasized the public health rationale behind these policies, particularly the role of early action in shaping outcomes. [Friedson et al. \(2021\)](#), for instance, study California’s shelter-in-place order (SIPO) and argue that its primary intent was not only to slow the spread of COVID-19 but also to delay its peak, giving the healthcare system time to prepare. They highlight California’s timing—implementing the policy while infection growth was still relatively modest—as a key to its potential effectiveness.

Other studies reinforce the theme that early and decisive interventions played a crucial role in shaping pandemic trajectories. [Yang \(2021\)](#) evaluate the lockdown of Wuhan, quantifying its causal effects on both population movement and COVID-19 transmission. They frame their analysis as providing timely policy guidance for other countries, emphasizing that rapid containment measures can have meaningful public health benefits. Similarly, [Cerqueti et al. \(2021\)](#), studying Italy, stress that the exponential nature of viral spread makes timing a fundamental determinant of policy success. They also raise an important political economy concern: if early policies succeed in flattening the curve, they may appear excessive in retrospect, prompting backlash over their economic cost. [Bayat et al. \(2020\)](#) quantify this idea directly, estimating that deaths in New York could have been reduced by over 80% had the lockdown been implemented just ten days earlier.

Outside of China and Europe, similar patterns are observed. [Lee and Yang \(2022\)](#) study South Korea and underscore the link between public health interventions and labor market outcomes. [Ke and Hsiao \(2022\)](#) show that while the economic toll of Hubei’s lockdown was sharp, it was also short-lived—suggesting that inaction might have caused even more persistent damage. [Sobiech Pellegrini \(2022\)](#) and [Li and Shankar \(2023\)](#) both find that premature lifting of lockdowns in Brazil and Texas, respectively, led to greater case counts.

Taken together, these studies make a consistent point: early action in the face of uncertainty can produce substantial public health benefits. But such actions are not costless, and much of the subsequent literature turned to the economic side of the equation.

The perceived tradeoff between public health and economic activity quickly became central to political discourse around COVID-19. [Osman et al. \(2022\)](#) document how opponents of lockdowns, particularly in conservative circles, criticized them as both ineffective and economically harmful. [Gibson and Sun \(2023\)](#) empirically assess the impact of statewide stay-at-home orders on labor markets, focusing on new jobless claims in the immediate aftermath. Their work reflects widespread concern that such restrictions, while potentially life-saving, might inflict deep labor market damage.

Other scholars have examined related interventions with broader welfare implications. [Cole et al. \(2020\)](#), using augmented synthetic control, study how the Wuhan lockdown affected air pollution and related mortality, arguing that large-scale behavioral shifts offer rare insight into the health impacts of everyday economic activity. [Oliu-Barton et al. \(2022\)](#) assess vaccine mandates (COVID certificates) in Europe and estimate that by increasing vaccine uptake, these policies both averted deaths and reduced GDP losses—highlighting cases where public health measures can yield direct economic benefits.

Across all of these studies, the central policy question remains: how should governments weigh the immediate economic costs of aggressive action against the uncertain but potentially vast benefits of preventing catastrophic health crises?

## 2.2. Methodological Concerns

Quantifying the economic impacts of such interventions poses serious challenges. Most of the empirical literature on COVID-19 interventions uses difference-in-differences (DiD) or synthetic control (SCM) frameworks. As noted by [Callaway and Li \(2023\)](#) and [Bilgel \(2022\)](#), these methods can be inappropriate in pandemic contexts where the assumptions they rely on—parallel trends, unaffected control units—are systematically violated. Of course, there are exceptions. Researchers have productively applied SCM and related panel methods to study the effects of NPIs in settings where some entities implemented policies while others did not (e.g., [Yang, 2021](#); [Ke and Hsiao, 2023](#)). These comparative case study designs are informative when credible cross-sectional variation exists.

Hawaii, however, presents a fundamentally different problem. The state is geographically and economically distinct. Every plausible control unit—whether another U.S. state or an island nation—faced concurrent COVID-related shocks. There is no unaffected or even approximately similar unit that could serve as a valid counterfactual. In such settings, cross-sectional comparisons break down.

This paper addresses that challenge by applying the Synthetic Historical Control (SHC) method, which constructs a counterfactual using only the treated unit’s own pre-intervention history. SHC represents a fundamental shift in the SCM paradigm: rather than borrowing strength across space, it borrows strength across time. This allows for credible causal inference when no control units exist—a situation that characterizes many of the most interesting and policy-relevant cases in economics.

## 3. METHODOLOGY

We know what happened given that Hawaii mandated two-week quarantines. We do not know what would have happened had they not done this. Given its novelty and radical deviation from SCM, this section summarizes the SHC method by [Chen et al. \(2024\)](#). For a comprehensive treatment of the method, including the theoretical proofs and asymptotic properties, readers are referred to [Chen et al. \(2024\)](#).

**Notation** Let  $\mathbb{R}$  denote the set of real numbers, and let  $\|\cdot\|_1$  denote the usual  $\ell_1$  vector norm. Throughout, I denote sets using calligraphic letters (e.g.,  $\mathcal{T}, \mathcal{S}, \mathcal{N}$ ) for clarity. Scalars are represented using plain lowercase letters (e.g.,  $h, t, n, m$ ), and matrices are denoted by bold uppercase letters (e.g.,  $\mathbf{X}, \mathbf{W}, \mathbf{Y}$ ). Given a vector  $\mathbf{x} = (x_1, x_2, \dots, x_T)^\top \in \mathbb{R}^T$  and an index set  $\mathcal{I} \subseteq \{1, \dots, T\}$ , I write  $\mathbf{x}_{\mathcal{I}} := (x_t)_{t \in \mathcal{I}} \in \mathbb{R}^{|\mathcal{I}|}$  to denote the subvector of  $\mathbf{x}$  corresponding to indices in  $\mathcal{I}$ , preserving their original order.

Let  $\mathcal{T} = \{1, 2, \dots, T\}$  index time at a monthly frequency. Define the pre-treatment period as  $\mathcal{T}_1 := \{t \in \mathcal{T} : t \leq T_0\}$  and the post-treatment period as  $\mathcal{T}_2 := \{t \in \mathcal{T} : t > T_0\}$ . The number of post-treatment periods is  $n = T - T_0$ . Let  $m \in \mathbb{N}$  denote the evaluation window length, i.e., the number of periods used to construct the SHC match. Define the evaluation period as the final  $m$  months of the pre-treatment period  $\mathcal{T}_{\text{eval}} := \{T_0 - m + 1, \dots, T_0\} \subset \mathcal{T}_1$ .

Let  $\mathbf{y} = (y_1, y_2, \dots, y_T)^\top \in \mathbb{R}^T$  denote the observed outcome vector for the treated unit. Define the pre-treatment outcome vector  $\mathbf{y}_{\text{pre}} := \mathbf{y}_{\mathcal{T}_1} \in \mathbb{R}^{T_0}$ , the evaluation subvector  $\mathbf{y}_{\text{eval}} := \mathbf{y}_{\mathcal{T}_{\text{eval}}} \in \mathbb{R}^m$ , and the post-treatment outcome vector  $\mathbf{y}_{\text{post}} := \mathbf{y}_{\mathcal{T}_2} \in \mathbb{R}^n$ .

The evaluation subvector  $\mathbf{y}_{\text{eval}}$  is the target pattern to be reconstructed using convex combinations of earlier segments from the pre-treatment period. To reconstruct the evaluation window, we compare it against earlier segments of the treated unit's own pre-treatment trajectory.

Define the donor matrix  $\tilde{\mathbf{Y}}_{\text{pre}} \in \mathbb{R}^{m \times N}$  as a collection of  $N$  overlapping length- $m$  subvectors extracted from the smoothed pre-treatment series. Each column is a contiguous subvector of the form  $\tilde{\mathbf{y}}_{[i, i+m-1]}$ , with  $i$  ranging over eligible start points in  $\mathcal{T}_1$  that precede the evaluation window. The precise construction is given in [Section 3.2](#).

## 3.1. Synthetic Historical Control Method

**Potential Outcomes** Following the Rubin potential outcomes framework, let  $y_t^1$  and  $y_t^0$  denote the potential outcomes for the treated unit at time  $t \in \mathcal{T}$ , corresponding to the outcomes under treatment and no treatment, respectively. The observed outcome at time  $t$  is  $y_t = d_t y_t^1 + (1 - d_t) y_t^0$ , where  $d_t = \mathbf{1}\{t > T_0\}$ .

The fundamental problem of causal inference is that for  $t > T_0$ , we observe only the treated potential outcome  $y_t^1 = y_t$ , while the untreated potential outcome  $y_t^0$  remains unobserved. Let the SHC estimator for  $y_t^0$  be denoted  $\hat{y}_t^0 := y_t^{\text{SHC}}$  for all  $t \in \mathcal{T}_2$ . Define the observed post-treatment outcome vector as  $\mathbf{y}_{\text{post}} := (y_t^1)_{t \in \mathcal{T}_2}$  and the estimated counterfactual vector as  $\mathbf{y}_{\text{post}}^0 := (\hat{y}_t^0)_{t \in \mathcal{T}_2}$ . The out-of-sample treatment effect is

$$\alpha_t := y_t^1 - y_t^0, \quad \text{for } t \in \mathcal{T}_2, \quad (3.1)$$

with estimator  $\hat{\alpha}_t := y_t - \hat{y}_t^0$ . Stacking over  $t \in \mathcal{T}_2$  yields:

$$\hat{\alpha} := \mathbf{y}_{\text{post}} - \mathbf{y}_{\text{post}}^0. \quad (3.2)$$

The average treatment effect on the treated (ATT) is then given by:

$$\widehat{\text{ATT}} := \frac{1}{n} \sum_{t=T_0+1}^{T_0+n} (y_t - \hat{y}_t^0) = \frac{1}{n} \mathbf{1}^\top \hat{\alpha}, \quad (3.3)$$

where  $\mathbf{1} \in \mathbb{R}^n$  is a vector of ones.

### 3.2. The Model

The SHC method introduced by [Chen et al. \(2024\)](#) constructs a counterfactual using the treated unit's own pre-treatment time series.<sup>1</sup> SHC requires a few assumptions: [Assumption 3.1](#), [Assumption 3.2](#), and [Assumption 3.3](#).

**ASSUMPTION 3.1. (ERROR PROPERTIES)** *The error term satisfies  $\mathbb{E}[\varepsilon_t] = 0$  and  $\mathbb{E}[\varepsilon_t^2] < \infty$  for all  $t \in \mathcal{T}$ .*

[Assumption 3.1](#) says unmodeled fluctuations or noise in the outcome are random.

**ASSUMPTION 3.2. (LATENT COMPONENT SMOOTHNESS AND MATCHING CONDITION)** *The latent trend  $\tilde{y}(t)$  is  $(H + 1)$  times differentiable for some integer  $H \geq 3$ , with derivative bounded as  $|\tilde{y}^{(H+1)}(t)| < \varepsilon$  uniformly over  $t \in \mathcal{T}_1$ . Additionally,  $\exists \mathbf{w}^* \in \Delta^{N-1}$  such that  $\mathbf{y}_{\text{eval}} = \tilde{\mathbf{Y}}_{\text{pre}} \mathbf{w}^*$ .*

[Assumption 3.2](#) allows us to approximate the post-treatment latent component using a local linear extrapolation of pre-treatment values and posits that perfect match exists within the donor pool for the evaluation window, much as [Abadie et al. \(2010\)](#) assume about donor pool units matching to the treated unit.

**ASSUMPTION 3.3. (FEASIBILITY)** *There exists a weight vector  $\mathbf{w} \in \Delta^{N-1}$  such that  $\hat{\mathbf{y}}_{\text{eval}} = \tilde{\mathbf{Y}}_{\text{pre}} \mathbf{w}$ , with  $\|\mathbf{w}\|_0 \leq s \ll N$ . Let  $\tilde{\mathbf{Y}}_{\text{pre}}^{(s)}$  denote the submatrix formed by the active columns of  $\tilde{\mathbf{Y}}_{\text{pre}}$  indexed by  $\text{supp}(\mathbf{w})$ ; then its Gram matrix satisfies  $\left(\tilde{\mathbf{Y}}_{\text{pre}}^{(s)}\right)^\top \tilde{\mathbf{Y}}_{\text{pre}}^{(s)} \succ 0$ .*

[Chen et al. \(2024\)](#) note that [Assumption 3.3](#) is the sample-variant of [Assumption 3.2](#). It

<sup>1</sup>SCM estimates the counterfactual by solving:

$$\mathbf{w}^* = \underset{\mathbf{w} \in \Delta^{N_0}}{\text{argmin}} \|\mathbf{y}_1 - \mathbf{Y}_0 \mathbf{w}\|_2^2, \quad \forall t \in \mathcal{T}_1, \quad (3.4)$$

where  $\mathbf{Y}_0 \in \mathbb{R}^{T_0 \times N_0}$  is the donor matrix of unexposed units and  $\Delta^{N_0} = \{\mathbf{w} \in \mathbb{R}_{\geq 0}^{N_0} : \|\mathbf{w}\|_1 = 1\}$  is the probability simplex.

says that the SHC weights are identifiable, sparse, and numerically stable, similar to the assumptions described by [Abadie \(2021\)](#).

The first step of the SHC method smooths  $\mathbf{y}_{\text{pre}}$  using local linear regression, where we fit the regression model:

$$y_j \approx a_t + b_t(j - t), \quad \text{for } j = 1, \dots, T_0, \quad (3.5)$$

with local intercept  $a_t \in \mathbb{R}$  and slope  $b_t \in \mathbb{R}$ . Define the centered time vector  $\mathbf{s}_t = (1 - t, 2 - t, \dots, T_0 - t)^\top \in \mathbb{R}^{T_0}$ , and the design matrix  $\mathbf{X}_t = [\mathbf{1} \ \mathbf{s}_t] \in \mathbb{R}^{T_0 \times 2}$ . The diagonal kernel weight matrix is:

$$\mathbf{W}_t = \text{diag}(w_{1t}, w_{2t}, \dots, w_{T_0t}), \quad w_{jt} \propto \exp\left(-\frac{1}{2} \left(\frac{j - t}{h}\right)^2\right), \quad (3.6)$$

where  $(h > 0)$  is a bandwidth parameter governing smoothness. The local coefficients  $\boldsymbol{\beta}_t = (a_t, b_t)^\top$  solve quadratic programming problem that is strictly convex

$$\boldsymbol{\beta}_t = \underset{\boldsymbol{\beta} \in \mathbb{R}^2}{\text{argmin}} \left\| \mathbf{W}_t^{1/2} (\mathbf{y}_{\text{pre}} - \mathbf{X}_t \boldsymbol{\beta}) \right\|_2^2. \quad (3.7)$$

Fortunately, this optimization admits a closed-form solution for the coefficients:

$$\boldsymbol{\beta}_t = (\mathbf{X}_t^\top \mathbf{W}_t \mathbf{X}_t)^{-1} \mathbf{X}_t^\top \mathbf{W}_t \mathbf{y}_{\text{pre}}. \quad (3.8)$$

The smoothed outcome at time  $t$  is  $\tilde{y}_t = a_t = \mathbf{e}_1^\top \boldsymbol{\beta}_t$ , where  $\mathbf{e}_1 = (1, 0)^\top$ . Let  $\tilde{\mathbf{y}} = (\tilde{y}_1, \dots, \tilde{y}_{T_0})^\top \in \mathbb{R}^{T_0}$  denote the full smoothed series. After we have smoothed over the pre-treatment period before the evaluation period, we build our donor matrix of historical control segments/units.

In the second step, SHC constructs  $N = T_0 - m - n + 1$  overlapping temporal segments of length  $m$ . Each segment  $\tilde{\mathbf{y}}_{[i, i+m-1]}$  is a contiguous subvector of the smoothed pre-treatment series. These comprise the donor matrix  $\tilde{\mathbf{Y}}_{\text{pre}} := [\tilde{\mathbf{y}}_{[1, m]} \ \tilde{\mathbf{y}}_{[2, m+1]} \ \dots \ \tilde{\mathbf{y}}_{[N, N+m-1]}] \in \mathbb{R}^{m \times N}$ .

**EXAMPLE 3.1. (CONSTRUCTING  $\tilde{\mathbf{Y}}_{\text{PRE}}$ )** Suppose we have  $T_0 = 60$  months of pre-treatment data and we wish to construct donor segments of length  $m = 12$  months in order to predict a post-treatment horizon of  $n = 6$  months. The number of eligible segments is:  $N = T_0 - m - n + 1 = 60 - 12 - 6 + 1 = 43$ . Each donor segment is a vector of 12 consecutive smoothed pre-treatment outcomes.

In the third step, we solve the following program to obtain the optimal weights:

$$\begin{aligned} \mathbf{w}^* = \underset{\mathbf{w} \in \mathbb{R}^N}{\text{argmin}} \quad & \left\| \tilde{\mathbf{y}}_{\text{eval}} - \tilde{\mathbf{Y}}_{\text{pre}} \mathbf{w} \right\|_2^2 + \varsigma \left\| \mathbf{C}_0^\top \mathbf{w} \right\|_2^2 \\ \text{subject to} \quad & \mathbf{w} \geq \mathbf{0}, \quad \mathbf{1}^\top \mathbf{w} = 1 \end{aligned} \quad (3.9)$$

[Chen et al. \(2024\)](#) note that in the general case when we have more donor units than we do the length of the segments ( $N > m$ ), the optimization becomes an ill-posed problem. The regularization term stabilizes the optimization by penalizing weight allocations along directions of low variance in the donor pool, reducing sensitivity to collinear segments. Formally,  $\mathbf{C}_0$  contains the eigenvectors of the Grammian matrix  $\mathbf{G} \in \mathbb{R}^{N \times N}$ .  $\mathbf{C}_0$

corresponds to the eigenvalues less than  $\tau > 0$ . These directions capture low-variance combinations of donor segments, and the penalty shrinks weights in those directions.<sup>2</sup> The corresponding weight vector  $\mathbf{w}^*$  is used to compute the counterfactual by applying it to the shifted donor segments:

$$\tilde{\mathbf{Y}}_{\text{post}} = [\tilde{\mathbf{y}}_{[1+m, m+n]} \quad \tilde{\mathbf{y}}_{[2+m, m+n+1]} \quad \cdots \quad \tilde{\mathbf{y}}_{[N+m, N+m+n-1]}] \in \mathbb{R}^{n \times N}. \quad (3.10)$$

This yields in-sample and out-of-sample predictions for the SHC estimator. The full counterfactual vector is then:  $\mathbf{y}^0 = \tilde{\mathbf{Y}}\mathbf{w}^* \in \mathbb{R}^{m+n}$ .

**Bandwidth Selection via Cross-Validation** To select the bandwidth parameter  $h$  used in local linear smoothing of  $\mathbf{y}_{\text{pre}}$ , I implement a *leave-one-out cross-validation* (LOOCV) procedure over a grid of candidate values  $h \in [h_{\min}, h_{\max}] \subset \mathbb{R}_{>0}$ . This procedure chooses the bandwidth that minimizes out-of-sample prediction error for the smoothed pre-treatment series  $\tilde{\mathbf{y}} \in \mathbb{R}^{T_0}$ , as defined in Equation~(3.8).

For each candidate  $h$ , and for each time point  $t \in \mathcal{T}_1$ , define Gaussian kernel weights

$$w_{\tau t}^{(h)} = \exp\left(-\frac{1}{2} \left(\frac{\tau - t}{h}\right)^2\right), \quad \tau \in \mathcal{T}_1 \setminus \{t\}, \quad (3.11)$$

used to construct the diagonal matrix  $\mathbf{W}_t^{(-t)} \in \mathbb{R}^{(T_0-1) \times (T_0-1)}$ , omitting the index  $t$ . Let  $\mathbf{X}_t^{(-t)} \in \mathbb{R}^{(T_0-1) \times 2}$  and  $\mathbf{y}_{\text{pre}}^{(-t)} \in \mathbb{R}^{T_0-1}$  denote the design matrix and response vector with time  $t$  excluded.

Then the local linear coefficients  $\beta_t^{(-t)} = (a_t^{(-t)}, b_t^{(-t)})^\top$  are obtained by solving Equation~(3.7). The leave-one-out smoothed value at time  $t$  is then:

$$\tilde{y}_t^{(-t)} = a_t^{(-t)} = \mathbf{e}_1^\top \beta_t^{(-t)}, \quad \text{where } \mathbf{e}_1 = (1, 0)^\top. \quad (3.12)$$

To pick the ideal bandwidth, we minimize:

$$h^* = \underset{h \in [h_{\min}, h_{\max}]}{\operatorname{argmin}} \left\{ \text{CV}(h) := \frac{1}{T_0} \sum_{t=1}^{T_0} \left( y_t - \tilde{y}_t^{(-t)} \right)^2 \right\}. \quad (3.13)$$

Once selected, the optimal bandwidth  $h^*$  is used to construct the full smoothed pre-treatment series.

**Control Group Selection** Chen et al. (2024) note that we should have a parsimonious approach to which historical controls we use, rather than choosing all of them. In their examples, SHC typically requires less than 28 controls. They advocate for a forward selection procedure to choose the optimal control group. I implement a slightly modified variant of this method to choose the ideal number of historical controls. Let  $\mathcal{N} = \{1, 2, \dots, N\}$  index the full set of donor segments in the pre-treatment matrix  $\tilde{\mathbf{Y}}_{\text{pre}} \in \mathbb{R}^{m \times N}$ . At each step the algorithm selects the donor index  $i_j \in \mathcal{N} \setminus \mathcal{S}_{j-1}$  that, when added to the current donor set  $\mathcal{S}_{j-1}$ , minimizes the in-sample mean squared error

<sup>2</sup>The  $\varsigma$  (sigma) term is a small positive constant used to stabilize the solution in cases of near multicollinearity.

(MSE) between the evaluation window and in-sample fit:

$$\mathcal{S}_j = \mathcal{S}_{j-1} \cup \left\{ \underset{i \in \mathcal{N} \setminus \mathcal{S}_{j-1}}{\operatorname{argmin}} \left\| \tilde{\mathbf{y}}_{\text{eval}} - \tilde{\mathbf{Y}}_{\text{pre}}^{(\mathcal{S}_{j-1} \cup \{i\})} \mathbf{w}^{(j)} \right\|_2^2 + \varsigma \left\| \mathbf{C}_0^\top \mathbf{w} \right\|_2^2 \right\}, \quad \mathcal{S}_0 = \emptyset, \quad (3.14)$$

where  $\tilde{\mathbf{Y}}_{\text{pre}}^{(\mathcal{S})}$  denotes the submatrix of  $\tilde{\mathbf{Y}}_{\text{pre}}$  with columns indexed by  $\mathcal{S}$ , and the weights  $\mathbf{w}^{(j)} \in \Delta^j$  are the optimal simplex-constrained weights for the selected donor set  $\mathcal{S}_j$ . This procedure greedily expands the donor set one segment at a time, always choosing the segment that yields the greatest marginal reduction in in-sample error.

To avoid overfitting, [Chen et al. \(2024\)](#) manually choose the number of donors the forward selection procedure should choose. In their case, they experiment with the single best historical segment (which rarely does well), the best 15, 20, and 27 historical control units. However, this is a subjective decision that will yield different results depending on what we select. To mitigate researcher degrees of freedom, I view the number of historical segments as a tuning parameter. I employ a forward selection-based algorithm to choose the number of controls. For my implementation of SHC, I employ an early stopping rule based on a modified Bayesian information criterion (BIC), borrowing inspiration from [Shi and Huang \(2023\)](#). For each donor set size  $j$ , I compute:

$$\text{BIC}(j) = m \cdot \log(\text{MSE}_j) + \lambda j, \quad (3.15)$$

where  $\text{MSE}_j$  is the in-sample error using donor set  $\mathcal{S}_j$ , and  $\lambda = \log(m)$  penalizes the inclusion of too many donors. The selection procedure terminates when  $\text{BIC}(j)$  increases for two consecutive steps. The final donor set  $\mathcal{S}^*$  is then taken to be  $\mathcal{S}_{j^*}$ , where:

$$j^* = \underset{j}{\operatorname{argmin}} \text{BIC}(j). \quad (3.16)$$

This approach adaptively selects a sparse donor subset that closely matches the treated unit's pre-intervention trajectory. The selected donor set  $\mathcal{S}^*$  is used to compute the SHC model predictions. I would like to mention that this forward selection method holds much in common with recently proposed methods from the literature. For example, [Shi and Huang \(2023\)](#) propose the forward selected panel data approach which selects the comparison group for the panel data approach by [Hsiao et al. \(2012\)](#) using an aggressive forward selection procedure. [Li \(2024\)](#), correcting for overfitting in the former, proposed the forward Difference-in-Differences estimator, where a forward selection approach is used to select the donor pool for the DiD model. Similar advances have been made SCM, as developed by [Cerulli \(2024\)](#), where a similar selection mechanism is used to estimate the treatment effect. It would be interesting, but is far outside the scope of this work, to compare these estimators to one another and see when their in-sample fits and point estimates agree with one another.

**Inference** Quantifying uncertainty around the counterfactual predictions is essential for valid inference. [Chen et al. \(2024\)](#) do not provide any inferential method to use aside from in-time placebos. Following recent work in post-modeling inference by [Cattaneo et al. \(2025\)](#), I construct agnostic conformal prediction intervals around the counterfactual series. These intervals are distribution-free, require no parametric assumptions, and are valid under exchangeability of residuals between the pre- and post-treatment periods (or, good in-sample fit).

Let  $\mathbf{y}_{\text{eval}}^0 \in \mathbb{R}^m$  denote the SHC-fitted values during the evaluation period and define the residuals as:

$$\boldsymbol{\varepsilon}_{\text{eval}} := \mathbf{y}_{\text{eval}} - \mathbf{y}_{\text{eval}}^0. \quad (3.17)$$

I then compute the sample mean  $\bar{\varepsilon}$  and variance  $\hat{\sigma}^2$  of these residuals:

$$\bar{\varepsilon} = m^{-1} \sum_{t \in \mathcal{T}_{\text{eval}}} \varepsilon_t, \quad \hat{\sigma}^2 = (m-1)^{-1} \sum_{t \in \mathcal{T}_{\text{eval}}} (\varepsilon_t - \bar{\varepsilon})^2. \quad (3.18)$$

To construct prediction intervals for the post-treatment counterfactual  $\mathbf{y}_{\text{post}}^0 \in \mathbb{R}^n$ , I assume the residuals are sub-Gaussian and apply a conformal bound based on Hoeffding-type concentration inequalities. For a miscoverage rate  $\alpha \in (0, 1)$ , the interval half-width is given by:

$$c_\alpha = \sqrt{2\hat{\sigma}^2 \log\left(\frac{2}{\alpha}\right)}. \quad (3.19)$$

Then, the agnostic  $(1 - \alpha)$  prediction interval for each  $t \in \mathcal{T}_2$  is given by:

$$[\hat{y}_t^0 + \bar{\varepsilon} - c_\alpha, \hat{y}_t^0 + \bar{\varepsilon} + c_\alpha]. \quad (3.20)$$

Stacking these yields lower and upper bound vectors:

$$\mathbf{L} := \mathbf{y}_{\text{post}}^0 + \bar{\varepsilon} \cdot \mathbf{1} - c_\alpha \cdot \mathbf{1}, \quad \mathbf{U} := \mathbf{y}_{\text{post}}^0 + \bar{\varepsilon} \cdot \mathbf{1} + c_\alpha \cdot \mathbf{1}, \quad (3.21)$$

which define the conformal confidence band around the counterfactual trajectory.

This approach allows finite-sample valid coverage under the assumption that the pre- and post-treatment residuals are exchangeable. Intuitively, if the SHC model fits the pre-treatment series well (i.e., residuals are small and centered), then its extrapolations into the post-treatment period are expected to carry similarly bounded uncertainty. In this sense, the residual distribution learned from the evaluation window is used to calibrate uncertainty for the out-of-sample outcomes.

These prediction intervals reflect uncertainty in the estimated counterfactual. Hence, if the observed outcome falls far outside the prediction interval, this is interpreted as evidence of a treatment effect. Unlike traditional confidence intervals around treatment effect estimates (e.g., using standard errors), the conformal intervals here do not assume a data-generating model. I use a default coverage level of 90% (i.e.,  $\alpha = 0.1$ ) in all applications.

#### 4. DATA

To evaluate the impact of Hawaii's COVID-19 border closure on its tourism economy, I use monthly time series spanning multiple economic dimensions, focusing on four primary outcomes: two related to tourism demand and two capturing broader economic impacts. All outcomes are transformed into annualized year-over-year (YoY) growth rates. This transformation mitigates nonstationary trends—crucial for the validity of the SHC estimator [Chen et al. \(2024\)](#)—and allows treatment effects to be interpreted directly as percentage point changes.

The first two outcomes capture the demand side of Hawaii’s tourism sector. Visitor Days measures the total number of person-days spent by visitors in a given month, integrating both the volume of arrivals and their length of stay. Occupancy refers to the statewide hotel occupancy rate, defined as the percentage of available hotel rooms that are occupied. Together, these indicators provide complementary measures of tourism demand: the former captures total visitor presence, while the latter reflects realized capacity utilization across the hospitality sector. Both series are sourced from Hawaii’s Department of Business, Economic Development & Tourism (DBEDT) and span from January 1990 to December 2020.

The remaining outcomes capture economic effects, one sector-specific and one broader in scope. *Accommodation Employment* tracks the number of employees in Hawaii’s accommodation industry, a labor-intensive sector directly tied to tourism activity. *The Coincident Economic Activity Index* summarizes statewide macroeconomic conditions based on four indicators: nonfarm payroll employment, the unemployment rate, average weekly hours worked in manufacturing, and wage and salary disbursements. This index is benchmarked to match the trend of gross state product, making it a comprehensive indicator of the state’s economic health. Both economic indicators are sourced from the Federal Reserve Economic Data (FRED) system and cover January 1991 to December 2020.


Hawaii’s mandatory quarantine policy took effect in late March 2020. I define the intervention to begin in March 2020, corresponding to time  $t = 362$ , and formalize the treatment indicator as:

$$d_t := \mathbf{1}\{t \geq 362\},$$


so that  $d_t = 1$  for all months after the policy’s implementation and  $d_t = 0$  otherwise. The final pre-treatment period is February 2020, yielding a pre-treatment sample of  $T_0 = 361$  months.

## 5. RESULTS

In Figure we visualize the outcomes of interest over time from January of 1991 to December 2020. We can see that the growth rate trends for all outcomes were relatively smooth across the full pre-intervention series, exhibiting seasonality and relatively low variance. We also see that the outcomes of interest precipitously dropped when the mandatory quarantine policy was implemented. The plot also suggests that [Assumption 3.1](#) and [Assumption 3.2](#) may be valid. For example, for [Assumption 3.1](#) to be violated, we would need a situation where the outcome had extremely high variance. While we see seasonality across all outcomes, the time series seemed smooth enough to be approximated by pre-treatment observations.



Here are the initial SHC results. For these results,  $m = 60$  (5 years of evaluation data), but the main results remain robust up until 8 years of evaluation data, with slight deviations in uncertainty.




Across all four primary outcomes, I find large and precisely estimated negative effects of Hawaii's border closure. Tourism demand was severely impacted. Visitor Days declined by an average of 85.4 percentage points, with the SHC model achieving a tight pre-treatment fit ( $\text{RMSE} = 1.17$ ). Hotel Occupancy dropped by 75.8 percentage points ( $\text{RMSE} = 1.097$ ). These effects represent near-complete collapses in aggregate visitor presence and hotel room utilization, consistent with the near-total shutdown of tourism activity. The closeness of pre-treatment fit suggests [Assumption 3.2](#) is plausible, given that the match is good.

Labor market outcomes also show pronounced effects. Accommodation Employment fell by an average of 60.7 percentage points ( $\text{ATT} = -60.691$ ;  $\text{RMSE} = 0.739$ ), while the broader Economic Activity Index dropped by 27.7 percentage points ( $\text{ATT} = -27.706$ ;  $\text{RMSE} = 0.303$ ). These results reflect both the direct impact on tourism-facing sectors and the wider economic slowdown induced by the closure. Even accounting for finite-sample prediction uncertainty, the observed collapse in tourism is extreme relative to the synthetic historical distribution. The observed values fall far outside of the 90% conformal prediction intervals, suggesting that the results are far outside of what we'd expect under a no-quarantine scenario.

The sharp decline in overall employment, especially in accommodations, is plausibly mediated by the collapse in tourism. Given Hawaii's economic structure, the vast majority of hotel/hospitality employment is directly demand-driven. Thus, the drop in visitor-days and occupancy rates leads, mechanically, to reduced staffing needs across the sector. This channel provides a natural explanation for the labor market contraction observed in the data.

To assess the credibility of our identification strategy, we conduct a placebo test by

assigning a false treatment date of September 2019 and evaluating outcomes through February 2020. Because this window predates the onset of the COVID-19 pandemic, a well-specified model should yield no substantial treatment effects.




The placebo results are reassuring. Across all four outcomes, the placebo ATTs are small relative to the true post-COVID estimates and comparable in magnitude to their pre-treatment RMSEs. For example, the placebo ATT for Visitor Days is  $-1.701$ , with an RMSE of  $1.433$ , which are orders of magnitude smaller than the estimated ATT. Similarly, for Hotel Occupancy, the placebo ATT is  $2.645$  with a  $\text{RMSE} = 1.312$ , compared to the estimated treatment effect of  $75$ . The labor market outcomes exhibit the same pattern: the placebo ATT for Accommodation Employment is  $1.121$  (versus  $-60.691$  post-COVID), and for the Economic Activity Index, it is  $0.147$  (versus  $-27.706$ ). In all cases, the true treatment effects dwarf any signal that emerges under a false intervention.

These results offer two key takeaways. First, the SHC model does not erroneously detect treatment effects when none exist—it exhibits strong temporal specificity. Second, even in cases where placebo ATTs are slightly nonzero, they fall within normal forecast uncertainty. This is confirmed by the conformal prediction intervals, which entirely cover the observed outcomes during the placebo window. That is, any apparent deviation between predicted and realized values is consistent with random variation under no treatment. These placebo checks strengthen the case for a causal interpretation of the post-March 2020 findings: the observed deviations are not artifacts of overfitting or model instability, but genuine responses to Hawaii’s border closure.

Now for additional outcomes. To estimate the impact of the mandatory quarantine on

fiscal losses, I collect data on taxes generated from Hawaii’s General Excise tax for the hotel industry and overall transient accommodations, using quarterly revenue data from 2010 to 2019 (with the pre-treatment period being from 1991 to the end of 2009).

I estimate that Hawaii’s General Excise Tax revenue from hotels fell by 88.3 percentage points after the March 2020 quarantine. Similarly, Transient Accommodations Tax receipts dropped by 91.7 percentage points. While quarterly data exhibit greater volatility and correspondingly higher prediction error (RMSEs = 6.85 and 8.6, respectively), the sheer magnitude of these effects supports the conclusion that the policy triggered a sharp fiscal contraction. These results provide a direct estimate of the policy’s revenue impact on the state budget.




## 6. DISCUSSION

Now, I discuss the broader interpretations of these results. One central question is the causal mechanism at work — some may reasonably ask: How do we know these treatment effects aren’t simply the result of people traveling less during the pandemic more generally?

This is a fair concern. While some portion of Hawaii’s economic contraction likely reflects reduced demand from risk-averse travelers, the magnitude and persistence of the effect suggest that policy played a primary role. Comparison to Florida helps isolate the role of policy design. In Q2 of 2020, Hawaii’s accommodation and food services earnings fell by 60% year-over-year, compared to a 38% decline in Florida, a state that avoided strict lockdowns and travel restrictions altogether.

Both Hawaii and Florida are tourism-heavy states. But while Hawaii shut down inbound travel with a strict quarantine policy, Florida remained open. The fact that Florida did not experience a comparable collapse — despite a similar economic structure — implies that quarantine policy, not just falling demand, drove the bulk of Hawaii’s losses. This interpretation is further supported by the recovery pattern: in Q1 of 2021, Hawaii’s year-over-year growth rate in earnings reached approximately 117%, compared to Florida’s 91%. This rebound suggests that the structure of Hawaii’s tourism economy remained intact — the losses were sharp, but temporary.



A recent paper by [Bond-Smith and Fuleky \(2022\)](#) further contextualizes these findings. In a detailed descriptive analysis of Hawaii’s economic contraction and sluggish employment recovery, they report that, as of 2022, non-farm employment remained nearly 10% below pre-pandemic levels, raising concerns of long-run “scarring” through persistent unemployment, out-migration, and educational disruption. However, they also note that “Hawaii experienced very few cases of COVID-19 relative to the rest of the US until the Delta and Omicron waves” in late 2021 and 2022. This reflects the protective success of early interventions such as the quarantine policy, as is similarly noted by [Ke and Hsiao \(2022\)](#).

The economic differences are even more striking in light of the public health outcomes. In 2020, Hawaii recorded just 16.8 COVID-19 deaths per 100,000 residents, the second-lowest in the country, while Florida recorded 56.4, ranking 11th highest. In 2021, Hawaii maintained its second-place position with 36.5 deaths per 100,000, while Florida’s rate more than doubled to 111.7, placing 33rd overall. These numbers highlight the tradeoff at the heart of aggressive containment policy: Hawaii’s border quarantine imposed a heavy short-run economic cost, but delivered substantial reductions in mortality. Contrary to claims that lockdowns would lead to long-term economic ruin, the experience of Hawaii suggests otherwise. Despite absorbing a dramatic short-term economic shock in 2020, Hawaii’s economy rebounded quickly after reopening. In contrast, Florida’s “open economy” strategy yielded only a modest short-run benefit, and by many metrics, including accommodation employment and gross state product, Florida lagged behind Hawaii in YoY recovery well into 2021 and beyond. In this light, the supposed economic tradeoff used to justify remaining open may have delivered neither the short-term protection nor the long-term gains that its advocates promised.

The present paper provides causal estimates of the short-run economic costs of Hawaii’s quarantine. While the estimated ATTs reveal a sharp and deliberate collapse in tourism-related activity, that collapse was not simply a pandemic inevitability, it was the direct consequence of policy designed to mitigate the spread of the virus. But if the policy prevented thousands of deaths by delaying widespread viral transmission, then that cost is not just measurable, it is arguably justifiable. Economic recovery, while uneven, is possible. The loss of life is not reversible.

## REFERENCES

- Abadie, A. (2021, 06). Using synthetic controls: Feasibility, data requirements, and methodological aspects. *Journal of Economic Literature* 59(2), 391–425.
- Abadie, A., A. Diamond, and J. Hainmueller (2010). Synthetic control methods for comparative case studies: Estimating the effect of california’s tobacco control program. *Journal of the American Statistical Association* 105(490), 493–505.
- Bayat, N., C. Morrin, Y. Wang, and V. Misra (2020). Synthetic control, synthetic interventions, and covid-19 spread: Exploring the impact of lockdown measures and herd immunity.
- Bilgel, F. (2022, 05). Effects of covid-19 lockdowns on social distancing in turkey. *The Econometrics Journal* 25(3), 781–805.
- Bond-Smith, S. and P. Fuleky (2022). The effects of the pandemic on the economy of hawaii. *Current Issues in Tourism* 26(23), 3846–3852.
- Callaway, B. and T. Li (2023). Policy evaluation during a pandemic. *Journal of Econometrics* 236(1), 105454.
- Cattaneo, M. D., Y. Feng, F. Palomba, and R. Titiunik (2025). scpi: Uncertainty quantification for synthetic control methods. *Journal of Statistical Software* 113(1), 1–38.
- Cerqueti, R., R. Coppier, A. Girardi, and M. Ventura (2021). The sooner the better: lives saved by the lockdown during the covid-19 outbreak. the case of italy. *The Econometrics Journal* 25(1), 46–70.
- Cerulli, G. (2024). Optimal initial donor selection for the synthetic control method. *Economics Letters* 244, 111976.
- Chen, Y.-T., J.-C. Yang, and T.-T. Yang (2024, September). Synthetic historical control for policy evaluation. Available at SSRN: <https://ssrn.com/abstract=4995085>.
- Cole, M. A., R. J. Elliott, and B. Liu (2020). The impact of the wuhan covid-19 lockdown on air pollution and health: a machine learning and augmented synthetic control approach. *Environmental and Resource Economics* 76(4), 553–580.
- Friedson, A. I., D. McNichols, J. J. Sabia, and D. Dave (2021). Shelter-in-place orders and public health: Evidence from california during the covid-19 pandemic. *Journal of Policy Analysis and Management* 40(1), 258–283.
- Gibson, J. and X. Sun (2023). A synthetic control analysis of us state level covid-19 stay-at-home orders on new jobless claims. *Journal of Economics and Finance* 47(1), 1–14.
- Hsiao, C., H. Steve Ching, and S. Ki Wan (2012). A panel data approach for program evaluation: Measuring the benefits of political and economic integration of hong kong with mainland china. *Journal of Applied Econometrics* 27(5), 705–740.
- Ivanova, I. (2020, May). Nearly 30% of hawaii’s workforce filed jobless claims. here’s a state-by-state breakdown of unemployment woes. *CBS News*. Accessed: 2025-07-14.
- Ke, X. and C. Hsiao (2022). Economic impact of the most drastic lockdown during covid-19 pandemic—the experience of hubei, china. *Journal of Applied Econometrics* 37(1), 187–209.
- Ke, X. and C. Hsiao (2023). Data subject to multiple treatment effects — disentangle the impacts of global pandemic and a specific disease control policy. *The Singapore Economic Review* 68(05), 1507–1527.
- Lee, J. and H.-S. Yang (2022). Pandemic and employment: evidence from covid-19 in south korea. *Journal of Asian economics* 78, 101432.
- Li, K. and V. Shankar (2023). Inference theory for quasi-experimental methods when data are non-stationary with unknown structure: Does the lifting of shelter-in-place keep covid-19 manageable? Available at SSRN 4321847.

- Li, K. T. (2024). Frontiers: A simple forward difference-in-differences method. *Marketing Science* 43(2), 239–468.
- Oliu-Barton, M., B. S. Pradelski, N. Woloszko, L. Guetta-Jeanrenaud, P. Aghion, P. Artus, A. Fontanet, P. Martin, and G. B. Wolff (2022). The effect of covid certificates on vaccine uptake, health outcomes, and the economy. *Nature Communications* 13(1), 3942.
- Osman, S. M. I., F. Islam, and N. Sakib (2022). Economic resilience in times of public health shock: The case of the us states. *Research in Economics* 76(4), 277–289.
- Shi, Z. and J. Huang (2023). Forward-selected panel data approach for program evaluation. *Journal of Econometrics* 234(2), 512–535.
- Sobiech Pellegrini, I. (2022). Untimely reopening? increase in the number of new covid-19 cases after reopening in one brazilian state. *OXford bulletin of economics and statistics* 84(4), 675–693.
- Yang, X. (2021). Does city lockdown prevent the spread of covid-19? new evidence from the synthetic control method. *Global Health Research and Policy* 6, 1–14.