

Paradise Lost? Estimating the Economic Impact of Hawaii’s Mandatory Quarantine using Synthetic Historical Controls

JARED GREATHOUSE[†]

[†]*Department of Policy and Management, Georgia State University*
E-mail: `jgreathouse3@student.gsu.edu`

Summary This article estimates the causal impact of Hawaii’s March 2020 mandatory quarantine on its tourism industry and labor market using the Synthetic Historical Control method. Analyzing visitor days, hotel occupancy, accommodations employment, and the Coincident Economic Activity Index (all in annualized year-over-year growth rates), the study reveals substantial negative impacts from the quarantine. Observed outcomes fall significantly outside 90% conformal out-of-sample prediction intervals, indicating a clear treatment effect. Placebo tests further validate the model’s internal credibility. This research demonstrates the method’s utility and informs the broader debate on economic impacts of anti-COVID policy.

Keywords: *Synthetic Controls, Policy Evaluation, COVID-19*

1. INTRODUCTION

This paper studies Hawaii’s mandatory quarantine to prevent the spread of COVID-19 (Yan et al., 2022). On March 26, 2020, the state imposed a mandatory 14-day quarantine on all arrivals, effectively sealing its borders. Within a month, Ivanova (2020) documents a third of Hawaii’s labor force was unemployed. Unlike most regions that experienced economic contraction from COVID-19 through voluntary distancing, fear, or uneven mandates, Hawaii implemented one of the few deliberate and centralized shutdowns in the United States. This policy was aimed not merely at mitigating the virus, but at suppressing travel and commerce altogether.

At the heart of this paper is a core policy question: What are the short-term economic costs of decisive early action in the face of an uncertain crisis? While much has been written about the epidemiological and social effects of pandemic interventions (say, on COVID-19 cases), less is known about their economic toll. This is especially true in cases where policymakers opted for maximal containment. The most famous case was Wuhan’s January 2020 lockdown, as studied by Ke and Hsiao (2022). However, work outside of this context is rare. Understanding these costs is essential not just for pandemic preparedness, but for future crises, such as climate shocks or geopolitical disruptions, where early intervention may again carry significant short-run economic risks in exchange for long-term societal benefits.

Hawaii offers a compelling empirical case. Its geographic isolation and economic dependence on tourism allowed it to pursue a form of near-total border control that was logistically and politically infeasible in most mainland states. As such, Hawaii provides a rare glimpse into a **zero-COVID** strategy implemented within a U.S.

federal context. Given how thoroughly it shut itself off from the outside world, much like Wuhan in early 2020, the strategy might be called The Great Isolation with Aloha characteristics.’’

Despite its policy distinctiveness and the strength of its intervention relative to the mainland, Hawaii has been largely excluded from existing empirical evaluations. This is partly due to the context of the policy, given the lack of suitable control units to draw comparisons from. Many of the most pressing policy questions in public policy and economics concern one-time, systemic shocks: What would the United States’ economy have looked like absent the American revolution? What would have happened to unemployment in Great Britain without Brexit? In these cases, credible untreated units often do not exist, making standard causal inference tools, like difference-in-differences (DiD) and synthetic control (SCM), difficult or impossible to apply. To compensate for this, I employ the Synthetic Historical Control (SHC) method by [Chen et al. \(2024\)](#). This method constructs a counterfactual using the treated unit’s own pre-intervention history using non-parametric regression based methods. Empirically, I find that Hawaii experienced substantial—but mostly short-lived—economic losses across the tourism and labor sectors, easing by early 2021.


The remainder of the paper proceeds as follows. Section 2 reviews prior research on early pandemic interventions and outlines key methodological challenges in studying one-off, jurisdiction-wide policies. Section 3 presents the SHC framework and its application in this setting. Section 4 describes the data. Section 5 reports the empirical results. Section 6 discusses broader implications for public policy and crisis response.

2. BACKGROUND

Much of the early policy discourse around COVID-19 focused on stopping the spread of the virus through non-pharmaceutical interventions (NPIs). A consistent theme emerged: the earlier and more decisive the intervention, the greater its potential to delay transmission and buy time for public health officials. Yet with these aggressive measures came a central concern—how to balance short-term economic pain caused by these interventions against long-term public health gain. This perceived tradeoff (economic ruin versus viral containment) framed many of the political and policy debates in 2020. But how real is this tradeoff? Are large-scale economic losses an unavoidable consequence of early containment, or can they be short-lived and ultimately outweighed by public health benefits?

Few comparisons illustrate these stakes more clearly than Hawaii and Florida. Both states rely heavily on tourism, yet took radically different approaches at the outset of the pandemic. Unlike Hawaii, Florida, kept its economy largely open and imposed minimal restrictions. As I show below, the consequences of these divergent strategies offer a window into the real costs—and possible benefits—of early, aggressive containment.

The economic effects of the mandatory quarantine/pandemic were immediate. In Q2 of 2020, Hawaii’s accommodation and food services earnings fell by 60% year-over-year. Florida’s declined by only 38%. But the gap soon narrowed. By Q1 of 2021, Hawaii’s year-over-year growth rate had surged to 117%, compared to Florida’s 91%, suggesting a rapid and robust rebound. Meanwhile, the public health effects diverged even more sharply. Hawaii recorded 16.8 COVID-19 deaths per 100,000 residents in 2020, the second-lowest



2.1. Literature Review

Collectively, these studies underscore a key lesson: early interventions can mitigate public health disasters, but often at economic cost. This tradeoff between health and the econ-

omy became central to both policymaking and public discourse. [Osman et al. \(2022\)](#) document how opponents of lockdowns, particularly in conservative circles, criticized them as economically inefficient and ineffective. [Gibson and Sun \(2023\)](#) evaluate statewide stay-at-home orders using synthetic control, focusing on new jobless claims as an economic indicator of interest. [Lee and Yang \(2022\)](#) examine the labor market impacts of COVID-19 in South Korea, highlighting the economic consequences of NPIs. [Ke and Hsiao \(2022\)](#) study the economic impact of Hubei’s lockdown. They show that while Hubei’s lockdown caused a sharp downturn in the short term, the effects were short-lived—suggesting that delay or inaction might have led to more lasting harm.

Other work has assessed broader welfare outcomes of interventions. [Cole et al. \(2020\)](#) show how the Wuhan lockdown reduced air pollution and associated mortality using augmented SCM. [Oliu-Barton et al. \(2022\)](#) study vaccine mandates in Europe and estimate that increased vaccine uptake not only averted deaths but also prevented GDP losses—suggesting that public health policy can have broader economic upside when well-designed. Despite the breadth of this literature, Hawaii’s case remains notably absent. Yet it represents perhaps the clearest U.S. example of an early, and most aggressive intervention. This paper contributes by bringing Hawaii into the empirical record and by offering credible estimates of the short-run economic costs of a “zero-COVID” style border ([Yuan, 2022](#)).

While much has been learned about the effects of pandemic interventions, measuring their economic impacts remains challenging. Most studies rely on DiD or SCM. Where cross-sectional variation in policy exists—e.g., some counties or states implementing lockdowns while others do not—SCM and panel methods have been used productively (e.g., [Yang, 2021](#); [Ke and Hsiao, 2023](#); [Gong et al., 2024](#)). But these designs rely on the availability of cross-sectional units to serve as the comparison group. They are ill-suited for settings like Hawaii, where no plausible comparison unit exists. However, both designs assume the existence of unaffected or minimally affected comparison units. As noted by [Callaway and Li \(2023\)](#) and [Bilgel \(2022\)](#), this and other critical assumptions are often violated in the context of a global pandemic. Every region experienced some form of shock, making it difficult to identify clean control units.

Hawaii is economically and geographically distinct. Every other U.S. state, as well as comparable island economies, experienced their own COVID-related shocks. There is no unaffected or even reasonably similar unit that can serve as a counterfactual. In such cases, cross-sectional methods fail. To address these methodological challenges, this study employs the SHC method by [Chen et al. \(2024\)](#), which constructs a counterfactual using the treated unit’s own pre-intervention history. Rather than borrowing information from across space, SHC borrows strength across time. This allows for credible inference in settings where traditional comparative approaches are unusable—precisely the situation posed by Hawaii’s quarantine.

3. METHODOLOGY

We know what happened given that Hawaii mandated two-week quarantines. We do not know what would have happened had they not done this. Given its novelty and radical deviation from SCM, this section summarizes the SHC method by [Chen et al.](#)

(2024). For a comprehensive treatment of the method, including the theoretical proofs and asymptotic properties, readers are referred to [Chen et al. \(2024\)](#).

Notation Let $\mathbb{R}$ denote the set of real numbers, and let $\|\cdot\|_1$ denote the usual ℓ_1 vector norm. Throughout, I denote sets using calligraphic letters (e.g., $\mathcal{T}, \mathcal{S}, \mathcal{N}$) for clarity. Scalars are represented using plain lowercase letters (e.g., h, t, n, m), and matrices are denoted by bold uppercase letters (e.g., $\mathbf{X}, \mathbf{W}, \mathbf{Y}$). Given a vector $\mathbf{x} = (x_1, x_2, \dots, x_T)^\top \in \mathbb{R}^T$ and an index set $\mathcal{I} \subseteq \{1, \dots, T\}$, I write $\mathbf{x}_{\mathcal{I}} := (x_t)_{t \in \mathcal{I}} \in \mathbb{R}^{|\mathcal{I}|}$ to denote the subvector of $\mathbf{x}$ corresponding to indices in $\mathcal{I}$, preserving their original order.

Let $\mathcal{T} = \{1, 2, \dots, T\}$ index time at a monthly frequency. Define the pre-treatment period as $\mathcal{T}_1 := \{t \in \mathcal{T} : t \leq T_0\}$ and the post-treatment period as $\mathcal{T}_2 := \{t \in \mathcal{T} : t > T_0\}$. The number of post-treatment periods is $n = T - T_0$. Let $m \in \mathbb{N}$ denote the evaluation window length, i.e., the number of periods used to construct the SHC match. Define the evaluation period as the final m months of the pre-treatment period $\mathcal{T}_{\text{eval}} := \{T_0 - m + 1, \dots, T_0\} \subset \mathcal{T}_1$.

Let $\mathbf{y} = (y_1, y_2, \dots, y_T)^\top \in \mathbb{R}^T$ denote the observed outcome vector for the treated unit. Define the pre-treatment outcome vector $\mathbf{y}_{\text{pre}} := \mathbf{y}_{\mathcal{T}_1} \in \mathbb{R}^{T_0}$, the evaluation subvector $\mathbf{y}_{\text{eval}} := \mathbf{y}_{\mathcal{T}_{\text{eval}}} \in \mathbb{R}^m$, and the post-treatment outcome vector $\mathbf{y}_{\text{post}} := \mathbf{y}_{\mathcal{T}_2} \in \mathbb{R}^n$.

The evaluation subvector $\mathbf{y}_{\text{eval}}$ is the target pattern to be reconstructed using convex combinations of earlier segments from the pre-treatment period. To reconstruct the evaluation window, we compare it against earlier segments of the treated unit's own pre-treatment trajectory.

Define the donor matrix $\tilde{\mathbf{Y}}_{\text{pre}} \in \mathbb{R}^{m \times N}$ as a collection of N overlapping length- m subvectors extracted from the smoothed pre-treatment series. Each column is a contiguous subvector of the form $\tilde{\mathbf{y}}_{[i, i+m-1]}$, with i ranging over eligible start points in $\mathcal{T}_1$ that precede the evaluation window. The precise construction is given in [Section 3.2](#).

3.1. Synthetic Historical Control Method

Potential Outcomes Following the Rubin potential outcomes framework, let y_t^1 and y_t^0 denote the potential outcomes for the treated unit at time $t \in \mathcal{T}$, corresponding to the outcomes under treatment and no treatment, respectively. The observed outcome at time t is $y_t = d_t y_t^1 + (1 - d_t) y_t^0$, where $d_t = \mathbf{1}\{t > T_0\}$.

The fundamental problem of causal inference is that for $t > T_0$, we observe only the treated potential outcome $y_t^1 = y_t$, while the untreated potential outcome y_t^0 remains unobserved. Let the SHC estimator for y_t^0 be denoted $\hat{y}_t^0 := y_t^{\text{SHC}}$ for all $t \in \mathcal{T}_2$. Define the observed post-treatment outcome vector as $\mathbf{y}_{\text{post}} := (y_t^1)_{t \in \mathcal{T}_2}$ and the estimated counterfactual vector as $\mathbf{y}_{\text{post}}^0 := (\hat{y}_t^0)_{t \in \mathcal{T}_2}$. The out-of-sample treatment effect is

$$\alpha_t := y_t^1 - y_t^0, \quad \text{for } t \in \mathcal{T}_2, \quad (3.1)$$

with estimator $\hat{\alpha}_t := y_t - \hat{y}_t^0$. Stacking over $t \in \mathcal{T}_2$ yields:

$$\hat{\boldsymbol{\alpha}} := \mathbf{y}_{\text{post}} - \mathbf{y}_{\text{post}}^0. \quad (3.2)$$

The average treatment effect on the treated (ATT) is then given by:

$$\widehat{\text{ATT}} := \frac{1}{n} \sum_{t=T_0+1}^{T_0+n} (y_t - \hat{y}_t^0) = \frac{1}{n} \mathbf{1}^\top \hat{\boldsymbol{\alpha}}, \quad (3.3)$$

where $\mathbf{1} \in \mathbb{R}^n$ is a vector of ones.

3.2. The Model

The SHC method introduced by [Chen et al. \(2024\)](#) constructs a counterfactual using the treated unit's own pre-treatment time series.¹ SHC requires a few assumptions: [Assumption 3.1](#), [Assumption 3.2](#), and [Assumption 3.3](#).

ASSUMPTION 3.1. (ERROR PROPERTIES) *The error term satisfies $\mathbb{E}[\varepsilon_t] = 0$ and $\mathbb{E}[\varepsilon_t^2] < \infty$ for all $t \in \mathcal{T}$.*

[Assumption 3.1](#) says unmodeled fluctuations or noise in the outcome are random.

ASSUMPTION 3.2. (LATENT COMPONENT SMOOTHNESS AND MATCHING CONDITION) *The latent trend $\tilde{y}(t)$ is $(H + 1)$ times differentiable for some integer $H \geq 3$, with derivative bounded as $|\tilde{y}^{(H+1)}(t)| < \varepsilon$ uniformly over $t \in \mathcal{T}_1$. Additionally, $\exists \mathbf{w}^* \in \Delta^{N-1}$ such that $\mathbf{y}_{eval} = \tilde{\mathbf{Y}}_{pre} \mathbf{w}^*$.*

[Assumption 3.2](#) allows us to approximate the post-treatment latent component using a local linear extrapolation of pre-treatment values and posits that perfect match exists within the donor pool for the evaluation window, much as [Abadie et al. \(2010\)](#) assume about donor pool units matching to the treated unit.

ASSUMPTION 3.3. (FEASIBILITY) *There exists a weight vector $\mathbf{w} \in \Delta^{N-1}$ such that $\hat{\mathbf{y}}_{eval} = \tilde{\mathbf{Y}}_{pre} \mathbf{w}$, with $\|\mathbf{w}\|_0 \leq s \ll N$. Let $\tilde{\mathbf{Y}}_{pre}^{(s)}$ denote the submatrix formed by the active columns of $\tilde{\mathbf{Y}}_{pre}$ indexed by $\text{supp}(\mathbf{w})$; then its Gram matrix satisfies $(\tilde{\mathbf{Y}}_{pre}^{(s)})^\top \tilde{\mathbf{Y}}_{pre}^{(s)} \succ 0$.*

[Chen et al. \(2024\)](#) note that [Assumption 3.3](#) is the sample-variant of [Assumption 3.2](#). It says that the SHC weights are identifiable, sparse, and numerically stable, similar to the assumptions described by [Abadie \(2021\)](#).

¹SCM estimates the counterfactual by solving:

$$\mathbf{w}^* = \underset{\mathbf{w} \in \Delta^{N_0}}{\text{argmin}} \|\mathbf{y}_1 - \mathbf{Y}_0 \mathbf{w}\|_2^2, \quad \forall t \in \mathcal{T}_1, \quad (3.4)$$

where $\mathbf{Y}_0 \in \mathbb{R}^{T_0 \times N_0}$ is the donor matrix of unexposed units and $\Delta^{N_0} = \{\mathbf{w} \in \mathbb{R}_{\geq 0}^{N_0} : \|\mathbf{w}\|_1 = 1\}$ is the probability simplex.

The first step of the SHC method smooths $\mathbf{y}_{\text{pre}}$ using local linear regression, where we fit the regression model:

$$y_j \approx a_t + b_t(j - t), \quad \text{for } j = 1, \dots, T_0, \quad (3.5)$$

with local intercept $a_t \in \mathbb{R}$ and slope $b_t \in \mathbb{R}$. Define the centered time vector $\mathbf{s}_t = (1 - t, 2 - t, \dots, T_0 - t)^\top \in \mathbb{R}^{T_0}$, and the design matrix $\mathbf{X}_t = [\mathbf{1} \ \mathbf{s}_t] \in \mathbb{R}^{T_0 \times 2}$. The diagonal kernel weight matrix is:

$$\mathbf{W}_t = \text{diag}(w_{1t}, w_{2t}, \dots, w_{T_0t}), \quad w_{jt} \propto \exp\left(-\frac{1}{2} \left(\frac{j - t}{h}\right)^2\right), \quad (3.6)$$

where $(h > 0)$ is a bandwidth parameter governing smoothness. The local coefficients $\beta_t = (a_t, b_t)^\top$ solve quadratic programming problem that is strictly convex

$$\beta_t = \underset{\beta \in \mathbb{R}^2}{\text{argmin}} \left\| \mathbf{W}_t^{1/2} (\mathbf{y}_{\text{pre}} - \mathbf{X}_t \beta) \right\|_2^2. \quad (3.7)$$

Fortunately, this optimization admits a closed-form solution for the coefficients:

$$\beta_t = (\mathbf{X}_t^\top \mathbf{W}_t \mathbf{X}_t)^{-1} \mathbf{X}_t^\top \mathbf{W}_t \mathbf{y}_{\text{pre}}. \quad (3.8)$$

The smoothed outcome at time t is $\tilde{y}_t = a_t = \mathbf{e}_1^\top \beta_t$, where $\mathbf{e}_1 = (1, 0)^\top$. Let $\tilde{\mathbf{y}} = (\tilde{y}_1, \dots, \tilde{y}_{T_0})^\top \in \mathbb{R}^{T_0}$ denote the full smoothed series. After we have smoothed over the pre-treatment period before the evaluation period, we build our donor matrix of historical control segments/units.

In the second step, SHC constructs $N = T_0 - m - n + 1$ overlapping temporal segments of length m . Each segment $\tilde{\mathbf{y}}_{[i, i+m-1]}$ is a contiguous subvector of the smoothed pre-treatment series. These comprise the donor matrix $\tilde{\mathbf{Y}}_{\text{pre}} := [\tilde{\mathbf{y}}_{[1, m]} \ \tilde{\mathbf{y}}_{[2, m+1]} \ \dots \ \tilde{\mathbf{y}}_{[N, N+m-1]}] \in \mathbb{R}^{m \times N}$.

EXAMPLE 3.1. (CONSTRUCTING $\tilde{\mathbf{Y}}_{\text{PRE}}$) Suppose we have $T_0 = 60$ months of pre-treatment data and we wish to construct donor segments of length $m = 12$ months in order to predict a post-treatment horizon of $n = 6$ months. The number of eligible segments is: $N = T_0 - m - n + 1 = 60 - 12 - 6 + 1 = 43$. Each donor segment is a vector of 12 consecutive smoothed pre-treatment outcomes.

In the third step, we solve the following program to obtain the optimal weights:

$$\begin{aligned} \mathbf{w}^* = \underset{\mathbf{w} \in \mathbb{R}^N}{\text{argmin}} \quad & \left\| \tilde{\mathbf{y}}_{\text{eval}} - \tilde{\mathbf{Y}}_{\text{pre}} \mathbf{w} \right\|_2^2 + \varsigma \left\| \mathbf{C}_0^\top \mathbf{w} \right\|_2^2 \\ \text{subject to} \quad & \mathbf{w} \geq \mathbf{0}, \quad \mathbf{1}^\top \mathbf{w} = 1 \end{aligned} \quad (3.9)$$

Chen et al. (2024) note that in the general case when we have more donor units than we do the length of the segments ($N > m$), the optimization becomes an ill-posed problem. The regularization term stabilizes the optimization by penalizing weight allocations along directions of low variance in the donor pool, reducing sensitivity to collinear segments. Formally, $\mathbf{C}_0$ contains the eigenvectors of the Grammian matrix $\mathbf{G} \in \mathbb{R}^{N \times N}$. $\mathbf{C}_0$ corresponds to the eigenvalues less than $\tau > 0$. These directions capture low-variance

combinations of donor segments, and the penalty shrinks weights in those directions.². The corresponding weight vector $\mathbf{w}^*$ is used to compute the counterfactual by applying it to the shifted donor segments:

$$\tilde{\mathbf{Y}}_{\text{post}} = [\tilde{\mathbf{y}}_{[1+m, m+n]} \quad \tilde{\mathbf{y}}_{[2+m, m+n+1]} \quad \cdots \quad \tilde{\mathbf{y}}_{[N+m, N+m+n-1]}] \in \mathbb{R}^{n \times N}. \quad (3.10)$$

This yields in-sample and out-of-sample predictions for the SHC estimator. The full counterfactual vector is then: $\mathbf{y}^0 = \tilde{\mathbf{Y}}\mathbf{w}^* \in \mathbb{R}^{m+n}$.

Bandwidth Selection via Cross-Validation To select the bandwidth parameter h used in local linear smoothing of $\mathbf{y}_{\text{pre}}$, I implement a *leave-one-out cross-validation* (LOOCV) procedure over a grid of candidate values $h \in [h_{\min}, h_{\max}] \subset \mathbb{R}_{>0}$. This procedure chooses the bandwidth that minimizes out-of-sample prediction error for the smoothed pre-treatment series $\tilde{\mathbf{y}} \in \mathbb{R}^{T_0}$, as defined in Equation~(3.8).

For each candidate h , and for each time point $t \in \mathcal{T}_1$, define Gaussian kernel weights

$$w_{\tau t}^{(h)} = \exp\left(-\frac{1}{2}\left(\frac{\tau - t}{h}\right)^2\right), \quad \tau \in \mathcal{T}_1 \setminus \{t\}, \quad (3.11)$$

used to construct the diagonal matrix $\mathbf{W}_t^{(-t)} \in \mathbb{R}^{(T_0-1) \times (T_0-1)}$, omitting the index t . Let $\mathbf{X}_t^{(-t)} \in \mathbb{R}^{(T_0-1) \times 2}$ and $\mathbf{y}_{\text{pre}}^{(-t)} \in \mathbb{R}^{T_0-1}$ denote the design matrix and response vector with time t excluded.

Then the local linear coefficients $\boldsymbol{\beta}_t^{(-t)} = (a_t^{(-t)}, b_t^{(-t)})^\top$ are obtained by solving Equation~(3.7). The leave-one-out smoothed value at time t is then:

$$\tilde{y}_t^{(-t)} = a_t^{(-t)} = \mathbf{e}_1^\top \boldsymbol{\beta}_t^{(-t)}, \quad \text{where } \mathbf{e}_1 = (1, 0)^\top. \quad (3.12)$$

To pick the ideal bandwidth, we minimize:

$$h^* = \underset{h \in [h_{\min}, h_{\max}]}{\operatorname{argmin}} \left\{ \text{CV}(h) := \frac{1}{T_0} \sum_{t=1}^{T_0} \left(y_t - \tilde{y}_t^{(-t)} \right)^2 \right\}. \quad (3.13)$$

Once selected, the optimal bandwidth h^* is used to construct the full smoothed pre-treatment series.

Control Group Selection Chen et al. (2024) note that we should have a parsimonious approach to which historical controls we use, rather than choosing all of them. In their examples, SHC typically requires less than 28 controls. They advocate for a forward selection procedure to choose the optimal control group. I implement a slightly modified variant of this method to choose the ideal number of historical controls. Let $\mathcal{N} = \{1, 2, \dots, N\}$ index the full set of donor segments in the pre-treatment matrix $\tilde{\mathbf{Y}}_{\text{pre}} \in \mathbb{R}^{m \times N}$. At each step the algorithm selects the donor index $i_j \in \mathcal{N} \setminus \mathcal{S}_{j-1}$ that, when added to the current donor set $\mathcal{S}_{j-1}$, minimizes the in-sample mean squared error (MSE) between the evaluation window and in-sample fit:

$$\mathcal{S}_j = \mathcal{S}_{j-1} \cup \left\{ \underset{i \in \mathcal{N} \setminus \mathcal{S}_{j-1}}{\operatorname{argmin}} \left\| \tilde{\mathbf{y}}_{\text{eval}} - \tilde{\mathbf{Y}}_{\text{pre}}^{(\mathcal{S}_{j-1} \cup \{i\})} \mathbf{w}^{(j)} \right\|_2^2 + \varsigma \left\| \mathbf{C}_0^\top \mathbf{w} \right\|_2^2 \right\}, \quad \mathcal{S}_0 = \emptyset, \quad (3.14)$$

²The ς (sigma) term is a small positive constant used to stabilize the solution in cases of near multicollinearity.

where $\tilde{\mathbf{Y}}_{\text{pre}}^{(\mathcal{S})}$ denotes the submatrix of $\tilde{\mathbf{Y}}_{\text{pre}}$ with columns indexed by $\mathcal{S}$, and the weights $\mathbf{w}^{(j)} \in \Delta^j$ are the optimal simplex-constrained weights for the selected donor set $\mathcal{S}_j$. This procedure greedily expands the donor set one segment at a time, always choosing the segment that yields the greatest marginal reduction in in-sample error.

To avoid overfitting, [Chen et al. \(2024\)](#) manually choose the number of donors the forward selection procedure should choose. In their case, they experiment with the single best historical segment (which rarely does well), the best 15, 20, and 27 historical control units. However, this is a subjective decision that will yield different results depending on what we select. To mitigate researcher degrees of freedom, I view the number of historical segments as a tuning parameter. I employ a forward selection-based algorithm to choose the number of controls. For my implementation of SHC, I employ an early stopping rule based on a modified Bayesian information criterion (BIC), borrowing inspiration from [Shi and Huang \(2023\)](#). For each donor set size j , I compute:

$$\text{BIC}(j) = m \cdot \log(\text{MSE}_j) + \lambda j, \quad (3.15)$$

where MSE_j is the in-sample error using donor set $\mathcal{S}_j$, and $\lambda = \log(m)$ penalizes the inclusion of too many donors. The selection procedure terminates when $\text{BIC}(j)$ increases for two consecutive steps. The final donor set $\mathcal{S}^*$ is then taken to be $\mathcal{S}_{j^*}$, where:

$$j^* = \underset{j}{\operatorname{argmin}} \text{BIC}(j). \quad (3.16)$$

This approach adaptively selects a sparse donor subset that closely matches the treated unit's pre-intervention trajectory. The selected donor set $\mathcal{S}^*$ is used to compute the SHC model predictions. I would like to mention that this forward selection method holds much in common with recently proposed methods from the literature. For example, [Shi and Huang \(2023\)](#) propose the forward selected panel data approach which selects the comparison group for the panel data approach by [Hsiao et al. \(2012\)](#) using an aggressive forward selection procedure. [Li \(2024\)](#), correcting for overfitting in the former, proposed the forward Difference-in-Differences estimator, where a forward selection approach is used to select the donor pool for the DiD model. Similar advances have been made SCM, as developed by [Cerulli \(2024\)](#), where a similar selection mechanism is used to estimate the treatment effect. It would be interesting, but is far outside the scope of this work, to compare these estimators to one another and see when their in-sample fits and point estimates agree with one another.

Inference Quantifying uncertainty around the counterfactual predictions is essential for valid inference. [Chen et al. \(2024\)](#) do not provide any inferential method to use aside from in-time placebos. Following recent work in post-modeling inference by [Cattaneo et al. \(2025\)](#), I construct agnostic conformal prediction intervals around the counterfactual series. These intervals are distribution-free, require no parametric assumptions, and are valid under exchangeability of residuals between the pre- and post-treatment periods (or, good in-sample fit).

Let $\mathbf{y}_{\text{eval}}^0 \in \mathbb{R}^m$ denote the SHC-fitted values during the evaluation period and define the residuals as:

$$\boldsymbol{\varepsilon}_{\text{eval}} := \mathbf{y}_{\text{eval}} - \mathbf{y}_{\text{eval}}^0. \quad (3.17)$$

I then compute the sample mean $\bar{\varepsilon}$ and variance $\hat{\sigma}^2$ of these residuals:

$$\bar{\varepsilon} = m^{-1} \sum_{t \in \mathcal{T}_{\text{eval}}} \varepsilon_t, \quad \hat{\sigma}^2 = (m-1)^{-1} \sum_{t \in \mathcal{T}_{\text{eval}}} (\varepsilon_t - \bar{\varepsilon})^2. \quad (3.18)$$

To construct prediction intervals for the post-treatment counterfactual $\mathbf{y}_{\text{post}}^0 \in \mathbb{R}^n$, I assume the residuals are sub-Gaussian and apply a conformal bound based on Hoeffding-type concentration inequalities. For a miscoverage rate $\alpha \in (0, 1)$, the interval half-width is given by:

$$c_\alpha = \sqrt{2\hat{\sigma}^2 \log\left(\frac{2}{\alpha}\right)}. \quad (3.19)$$

Then, the agnostic $(1 - \alpha)$ prediction interval for each $t \in \mathcal{T}_2$ is given by:

$$[\hat{y}_t^0 + \bar{\varepsilon} - c_\alpha, \hat{y}_t^0 + \bar{\varepsilon} + c_\alpha]. \quad (3.20)$$

Stacking these yields lower and upper bound vectors:

$$\mathbf{L} := \mathbf{y}_{\text{post}}^0 + \bar{\varepsilon} \cdot \mathbf{1} - c_\alpha \cdot \mathbf{1}, \quad \mathbf{U} := \mathbf{y}_{\text{post}}^0 + \bar{\varepsilon} \cdot \mathbf{1} + c_\alpha \cdot \mathbf{1}, \quad (3.21)$$

which define the conformal confidence band around the counterfactual trajectory.

This approach allows finite-sample valid coverage under the assumption that the pre- and post-treatment residuals are exchangeable. Intuitively, if the SHC model fits the pre-treatment series well (i.e., residuals are small and centered), then its extrapolations into the post-treatment period are expected to carry similarly bounded uncertainty. In this sense, the residual distribution learned from the evaluation window is used to calibrate uncertainty for the out-of-sample outcomes.

These prediction intervals reflect uncertainty in the estimated counterfactual. Hence, if the observed outcome falls far outside the prediction interval, this is interpreted as evidence of a treatment effect. Unlike traditional confidence intervals around treatment effect estimates (e.g., using standard errors), the conformal intervals here do not assume a data-generating model. I use a default coverage level of 90% (i.e., $\alpha = 0.1$) in all applications.

4. DATA

To evaluate the impact of Hawaii’s COVID-19 border closure on its tourism economy, I use monthly time series spanning multiple economic dimensions, focusing on four primary outcomes: two related to tourism demand and two capturing broader economic impacts. All outcomes are transformed into annualized year-over-year (YoY) growth rates. This transformation mitigates nonstationary trends—crucial for the validity of the SHC estimator [Chen et al. \(2024\)](#)—and allows treatment effects to be interpreted directly as percentage point changes.

The first two outcomes capture the demand side of Hawaii’s tourism sector. Visitor Days measures the total number of person-days spent by visitors in a given month, integrating both the volume of arrivals and their length of stay. Occupancy refers to the statewide

hotel occupancy rate, defined as the percentage of available hotel rooms that are occupied. Together, these indicators provide complementary measures of tourism demand: the former captures total visitor presence, while the latter reflects realized capacity utilization across the hospitality sector. Both series are sourced from Hawaii’s Department of Business, Economic Development & Tourism (DBEDT) and span from January 1990 to December 2020.

The remaining outcomes capture economic effects, one sector-specific and one broader in scope. *Accommodation Employment* tracks the number of employees in Hawaii’s accommodation industry, a labor-intensive sector directly tied to tourism activity. *The Coincident Economic Activity Index* summarizes statewide macroeconomic conditions based on four indicators: nonfarm payroll employment, the unemployment rate, average weekly hours worked in manufacturing, and wage and salary disbursements. This index is benchmarked to match the trend of gross state product, making it a comprehensive indicator of the state’s economic health. Both economic indicators are sourced from the Federal Reserve Economic Data (FRED) system and cover January 1991 to December 2020.


Hawaii’s mandatory quarantine policy took effect in late March 2020. I define the intervention to begin in March 2020, corresponding to time $t = 362$, and formalize the treatment indicator as:

$$d_t := \mathbf{1}\{t \geq 362\},$$


so that $d_t = 1$ for all months after the policy’s implementation and $d_t = 0$ otherwise. The final pre-treatment period is February 2020, yielding a pre-treatment sample of $T_0 = 361$ months.

5. RESULTS

In Figure we visualize the outcomes of interest over time from January of 1991 to December 2020. We can see that the growth rate trends for all outcomes were relatively smooth across the full pre-intervention series, exhibiting seasonality and relatively low variance. We also see that the outcomes of interest precipitously dropped when the mandatory quarantine policy was implemented. The plot also suggests that [Assumption 3.1](#) and [Assumption 3.2](#) may be valid. For example, for [Assumption 3.1](#) to be violated, we would need a situation where the outcome had extremely high variance. While we see seasonality across all outcomes, the time series seemed smooth enough to be approximated by pre-treatment observations.



Here are the initial SHC results. For these results, $m = 60$ (5 years of evaluation data), but the main results remain robust up until 8 years of evaluation data, with slight deviations in uncertainty.




Across all four primary outcomes, I find large and precisely estimated negative effects of Hawaii's border closure. Tourism demand was severely impacted. Visitor Days declined by an average of 85.4 percentage points, with the SHC model achieving a tight pre-treatment fit ($\text{RMSE} = 1.17$). Hotel Occupancy dropped by 75.8 percentage points ($\text{RMSE} = 1.097$). These effects represent near-complete collapses in aggregate visitor presence and hotel room utilization, consistent with the near-total shutdown of tourism activity. The closeness of pre-treatment fit suggests [Assumption 3.2](#) is plausible, given that the match is good.

Labor market outcomes also show pronounced effects. Accommodation Employment fell by an average of 60.7 percentage points ($\text{ATT} = -60.691$; $\text{RMSE} = 0.739$), while the broader Economic Activity Index dropped by 27.7 percentage points ($\text{ATT} = -27.706$; $\text{RMSE} = 0.303$). These results reflect both the direct impact on tourism-facing sectors and the wider economic slowdown induced by the closure. Even accounting for finite-sample prediction uncertainty, the observed collapse in tourism is extreme relative to the synthetic historical distribution. The observed values fall far outside of the 90% conformal prediction intervals, suggesting that the results are far outside of what we'd expect under a no-quarantine scenario.

The sharp decline in overall employment, especially in accommodations, is plausibly mediated by the collapse in tourism. Given Hawaii's economic structure, the vast majority of hotel/hospitality employment is directly demand-driven. Thus, the drop in visitor-days and occupancy rates leads, mechanically, to reduced staffing needs across the sector. This channel provides a natural explanation for the labor market contraction observed in the data.

To assess the credibility of our identification strategy, we conduct a placebo test by

assigning a false treatment date of September 2019 and evaluating outcomes through February 2020. Because this window predates the onset of the COVID-19 pandemic, a well-specified model should yield no substantial treatment effects.




The placebo results are reassuring. Across all four outcomes, the placebo ATTs are small relative to the true post-COVID estimates and comparable in magnitude to their pre-treatment RMSEs. For example, the placebo ATT for Visitor Days is -1.701 , with an RMSE of 1.433 , which are orders of magnitude smaller than the estimated ATT. Similarly, for Hotel Occupancy, the placebo ATT is 2.645 with a $\text{RMSE} = 1.312$, compared to the estimated treatment effect of 75 . The labor market outcomes exhibit the same pattern: the placebo ATT for Accommodation Employment is 1.121 (versus -60.691 post-COVID), and for the Economic Activity Index, it is 0.147 (versus -27.706). In all cases, the true treatment effects dwarf any signal that emerges under a false intervention.

These results offer two key takeaways. First, the SHC model does not erroneously detect treatment effects when none exist—it exhibits strong temporal specificity. Second, even in cases where placebo ATTs are slightly nonzero, they fall within normal forecast uncertainty. This is confirmed by the conformal prediction intervals, which entirely cover the observed outcomes during the placebo window. That is, any apparent deviation between predicted and realized values is consistent with random variation under no treatment. These placebo checks strengthen the case for a causal interpretation of the post-March 2020 findings: the observed deviations are not artifacts of overfitting or model instability, but genuine responses to Hawaii’s border closure.

Now for additional outcomes. To estimate the impact of the mandatory quarantine on

fiscal losses, I collect data on taxes generated from Hawaii’s General Excise tax for the hotel industry and overall transient accommodations, using quarterly revenue data from 2010 to 2019 (with the pre-treatment period being from 1991 to the end of 2009).

I estimate that Hawaii’s General Excise Tax revenue from hotels fell by 88.3 percentage points after the March 2020 quarantine. Similarly, Transient Accommodations Tax receipts dropped by 91.7 percentage points. While quarterly data exhibit greater volatility and correspondingly higher prediction error (RMSEs = 6.85 and 8.6, respectively), the sheer magnitude of these effects supports the conclusion that the policy triggered a sharp fiscal contraction. These results provide a direct estimate of the policy’s revenue impact on the state budget.



6. DISCUSSION

A recent paper by [Bond-Smith and Fuleky \(2022\)](#) further contextualizes these findings. In a detailed descriptive analysis of Hawaii’s economic contraction and sluggish employment recovery, they report that, as of 2022, non-farm employment remained nearly 10% below pre-pandemic levels, raising concerns of long-run “scarring” through persistent unemployment, out-migration, and educational disruption. However, they also note that “Hawaii experienced very few cases of COVID-19 relative to the rest of the US until the Delta and Omicron waves” in late 2021 and 2022. This reflects the protective success of early interventions such as the quarantine policy, as is similarly noted by [Ke and Hsiao \(2022\)](#).

Returning to the above Florida comparison, these numbers highlight the tradeoff at the heart of aggressive containment policy: Hawaii’s border quarantine imposed a heavy short-run economic cost, but delivered substantial reductions in mortality. Contrary to claims that lockdowns would lead to long-term economic ruin, the experience of Hawaii suggests otherwise. Despite absorbing a dramatic short-term economic shock in 2020, Hawaii’s economy rebounded quickly after reopening. In contrast, Florida’s “open economy” strategy yielded only a modest short-run benefit, and by many metrics, including accommodation employment and gross state product, Florida lagged behind Hawaii in YoY recovery well into 2021 and beyond. In this light, the supposed economic tradeoff used to justify remaining open may have delivered neither the short-term protection nor the long-term gains that its advocates promised.

The present paper provides causal estimates of the short-run economic costs of Hawaii’s quarantine. While the estimated ATTs reveal a sharp and deliberate collapse in tourism-related activity, that collapse was not simply a pandemic inevitability, it was the direct consequence of policy designed to mitigate the spread of the virus. But if the policy prevented thousands of deaths by delaying widespread viral transmission, then that cost is not just measurable, it is arguably justifiable. Economic recovery, while uneven, is possible. The loss of life is not reversible.

REFERENCES

- Abadie, A. (2021, 06). Using synthetic controls: Feasibility, data requirements, and methodological aspects. *Journal of Economic Literature* 59(2), 391–425.
- Abadie, A., A. Diamond, and J. Hainmueller (2010). Synthetic control methods for comparative case studies: Estimating the effect of california’s tobacco control program. *Journal of the American Statistical Association* 105(490), 493–505.
- Bayat, N., C. Morrin, Y. Wang, and V. Misra (2020). Synthetic control, synthetic interventions, and covid-19 spread: Exploring the impact of lockdown measures and herd immunity.
- Bilgel, F. (2022, 05). Effects of covid-19 lockdowns on social distancing in turkey. *The Econometrics Journal* 25(3), 781–805.
- Bond-Smith, S. and P. Fuleky (2022). The effects of the pandemic on the economy of hawaii. *Current Issues in Tourism* 26(23), 3846–3852.
- Callaway, B. and T. Li (2023). Policy evaluation during a pandemic. *Journal of Econometrics* 236(1), 105454.
- Cattaneo, M. D., Y. Feng, F. Palomba, and R. Titiunik (2025). scpi: Uncertainty quantification for synthetic control methods. *Journal of Statistical Software* 113(1), 1–38.
- Cerqueti, R., R. Coppier, A. Girardi, and M. Ventura (2021). The sooner the better: lives saved by the lockdown during the covid-19 outbreak. the case of italy. *The Econometrics Journal* 25(1), 46–70.
- Cerulli, G. (2024). Optimal initial donor selection for the synthetic control method. *Economics Letters* 244, 111976.
- Chen, Y.-T., J.-C. Yang, and T.-T. Yang (2024, September). Synthetic historical control for policy evaluation. Available at SSRN: <https://ssrn.com/abstract=4995085>.
- Cole, M. A., R. J. Elliott, and B. Liu (2020). The impact of the wuhan covid-19 lockdown on air pollution and health: a machine learning and augmented synthetic control approach. *Environmental and Resource Economics* 76(4), 553–580.
- Friedson, A. I., D. McNichols, J. J. Sabia, and D. Dave (2021). Shelter-in-place orders and public health: Evidence from california during the covid-19 pandemic. *Journal of Policy Analysis and Management* 40(1), 258–283.
- Gibson, J. and X. Sun (2023). A synthetic control analysis of us state level covid-19 stay-at-home orders on new jobless claims. *Journal of Economics and Finance* 47(1), 1–14.
- Gong, D., Z. Shang, Y. Su, A. Yan, and Q. Zhang (2024). Economic impacts of china’s zero-covid policies. *China Economic Review* 83, 102101.
- Hsiao, C., H. Steve Ching, and S. Ki Wan (2012). A panel data approach for program evaluation: Measuring the benefits of political and economic integration of hong kong with mainland china. *Journal of Applied Econometrics* 27(5), 705–740.
- Ivanova, I. (2020, May). Nearly 30% of hawaii’s workforce filed jobless claims. here’s a state-by-state breakdown of unemployment woes. *CBS News*. Accessed: 2025-07-14.

- Ke, X. and C. Hsiao (2022). Economic impact of the most drastic lockdown during covid-19 pandemic—the experience of hubei, china. *Journal of Applied Econometrics* 37(1), 187–209.
- Ke, X. and C. Hsiao (2023). Data subject to multiple treatment effects — disentangle the impacts of global pandemic and a specific disease control policy. *The Singapore Economic Review* 68(05), 1507–1527.
- Lee, J. and H.-S. Yang (2022). Pandemic and employment: evidence from covid-19 in south korea. *Journal of Asian economics* 78, 101432.
- Li, K. and V. Shankar (2023). Inference theory for quasi-experimental methods when data are non-stationary with unknown structure: Does the lifting of shelter-in-place keep covid-19 manageable? *Available at SSRN* 4321847.
- Li, K. T. (2024). Frontiers: A simple forward difference-in-differences method. *Marketing Science* 43(2), 239–468.
- National Center for Health Statistics (2023, February). Covid-19 mortality by state. https://www.cdc.gov/nchs/pressroom/sosmap/covid19_mortality_final/COVID19.htm. Accessed: 2025-07-23.
- Oliu-Barton, M., B. S. Pradelski, N. Woloszko, L. Guetta-Jeanrenaud, P. Aghion, P. Artus, A. Fontanet, P. Martin, and G. B. Wolff (2022). The effect of covid certificates on vaccine uptake, health outcomes, and the economy. *Nature Communications* 13(1), 3942.
- Osman, S. M. I., F. Islam, and N. Sakib (2022). Economic resilience in times of public health shock: The case of the us states. *Research in Economics* 76(4), 277–289.
- Shi, Z. and J. Huang (2023). Forward-selected panel data approach for program evaluation. *Journal of Econometrics* 234(2), 512–535.
- Sobiech Pellegrini, I. (2022). Untimely reopening? increase in the number of new covid-19 cases after reopening in one brazilian state. *OXford bulletin of economics and statistics* 84(4), 675–693.
- Yan, Y., M. I. Shah, G. D. Sharma, R. Chopra, Z. Fareed, and U. Shahzad (2022). Can tourism sustain itself through the pandemic: nexus between tourism, covid-19 cases and air quality spread in the ‘pineapple state’ hawaii. *Current Issues in Tourism* 25(3), 421–440.
- Yang, X. (2021). Does city lockdown prevent the spread of covid-19? new evidence from the synthetic control method. *Global Health Research and Policy* 6, 1–14.
- Yuan, S. (2022). Zero covid in china: what next? *The Lancet* 399(10338), 1856–1857.