

The Augmented Synthetic Historical Control Method

1 Introduction

Causal inference in economic and policy research often compares a treated set of units to a control or donor pool. The synthetic control method (SCM) of [Abadie et al. \[2010\]](#) is one of the most widely used approaches in panel data econometrics, second only to difference-in-differences. The idea of SC is straightforward: the untreated potential outcomes of a treated unit are approximated by a weighted average of donor units. Difference-in-differences relies on the assumption of parallel pre-intervention trends [\[Baker et al., 2025\]](#).

In many settings, however, no suitable donor units are available—for example, during global events such as pandemics or financial crises, or when no valid donor units exist. When no valid untreated donor units are available, researchers often turn to the interrupted time series (ITS) design. In its canonical form, ITS models the pre-treatment outcome path as a parametric linear trend and tests for changes in level or slope following an intervention. In fact, even comparative interrupted time series designs have been used by policy analysts [\[Coopersmith et al., 2022\]](#). Even so, ITS can be fragile in many realistic applications. Its validity depends critically on the assumption that an OLS line fitted to the pre-treatment period provides a reasonable counterfactual projection. In practice, many outcome series feature nonlinear dynamics, seasonality, or structural breaks that violate this assumption. Moreover, ITS typically imposes a homogeneous effect structure, assuming a single jump in level and/or slope that applies equally across all post-treatment periods. These restrictions can lead to biased inference when the true effects are time-varying or when pre-treatment is poor.

To address this shortcoming, [Chen et al. \[2024\]](#) proposed the synthetic historical control method (SHC). Rather than relying on external donors, SHC constructs a donor pool from the treated unit’s own pre-treatment history, splitting the pre-treatment period into overlapping segments. The SHC also, unlike ITS, uses kernel smoothing to account for non-linearities in the pre-treatment time series data which ITS does not naturally account for. A quadratic program then estimates weights over these segments to predict counterfactual outcomes, allowing for effect size heterogeneity found in SCM and increasingly DID but not ITS. However, like [Abadie et al. \[2010, 2015\]](#), SHC restricts weights to the simplex. This guarantees interpretability but can be too restrictive when noise or complex dynamics make exact pre-treatment fit unattainable.

This paper introduces the Augmented Synthetic Historical Control method (ASHC), which extends SHC in two key ways. First, ASHC relaxes the simplex constraint to an affine constraint, requiring only that weights sum to one. This permits negative weights, allowing controlled extrapolation beyond the convex hull of observed segments in exchange for better fit. Second, ASHC adds a quadratic proximity penalty that discourages large deviations from the SHC solution, in the spirit of the augmented SCM [\[Ben-Michael et al., 2021\]](#). This encourages the estimator to remain close to SHC unless extrapolation is needed. ASHC retains the stabilizing Gram penalty from [Chen et al. \[2024\]](#), which regularizes low-variance directions in the donor matrix. Together, these three components yield a new estimator that generalizes SHC to settings where exact pre-treatment fit is otherwise impossible.

The paper makes three contributions. Methodologically, it introduces the ASHC estimator, which generalizes SHC by permitting limited, penalized extrapolation while maintaining the stabilizing Gram penalty. Theoretically, it derives finite-sample error bounds that decompose ASHC prediction error into four interpretable parts: the SHC residual error, a regularization-controlled deviation term, smoothness-induced bias, and noise. It further establishes asymptotic consistency and a convergence rate of $O_p(m^{-H/(2H+3)})$, which approaches the near-parametric rate $O_p(m^{-1/2})$ as the smoothing order H increases. Practically, this means that when

the pre-treatment length m is large (e.g. $m = 120$) and H grows, the convergence rate becomes practically close to the parametric $O_p(m^{-1/2})$, so ASHC achieves efficient performance in realistic sample sizes.

While the finite-sample analysis provides guarantees on the accuracy of ASHC for realistic data sizes, asymptotic results are equally important. The finite-sample bounds show how ASHC predictions may differ from SHC depending on regularization strength, operator norm, smoothness bias, and noise. The asymptotic results complement this by demonstrating that, as the amount of pre-treatment data grows, ASHC’s prediction error vanishes at a near-parametric rate under mild conditions. Together, these perspectives ensure that ASHC is both practically reliable in finite samples and theoretically sound in large-sample limits.

The discussion begins with a careful review of SHC, its assumptions, and its implementation, before introducing the ASHC framework. We then turn to the theoretical analysis, establishing finite-sample error bounds and asymptotic properties. The empirical performance of ASHC is explored through calibrated simulations, followed by an application that illustrates its usefulness in practice. Proofs are collected in the appendix, and all computations are carried out using Python’s `mlsynth` library (<https://mlsynth.readthedocs.io/en/latest/>).

For policy analysts, the message is straightforward: ASHC provides a principled way to construct reliable counterfactuals even in environments where no external donor units are available. Furthermore, it allows policy analysts to do this even when methods such as interrupted time series method or the [Chen et al. \[2024\]](#) would not do well.

The remainder of the paper proceeds as follows. Section 1.1 situates our contribution within the existing literature on causal inference in policy evaluation. We then review the SHC method in Section 2.1 before introducing the augmented SHC (ASHC) in Section 2.2. Section 3 presents our theoretical results, including finite-sample error bounds and asymptotic properties. We then turn to an empirical application in Section 4, replicating the analysis of Barcelona’s 2015 hotel moratorium from [Bel et al. \[2021\]](#), a setting that allows us to compare ITS, SCM, SHC, and ASHC on equal footing. Finally, Section 5 concludes with implications for both methodological research and policy analysis.

1.1 Literature Review

2 Methodology

2.1 The Synthetic Historical Control Method

Let $\mathcal{T} = \{1, \dots, T\}$ index time, with pre-treatment period $\mathcal{T}_1 = \{t \in \mathcal{T} : t \leq T_0\}$ and post-treatment period $\mathcal{T}_2 = \{t \in \mathcal{T} : t > T_0\}$. The number of post-treatment periods is $n = T - T_0$, and an evaluation window of length m is defined as the final m pre-treatment periods $\mathcal{T}_{\text{eval}} = \{T_0 - m + 1, \dots, T_0\}$. Let $\mathbf{y} = (y_1, \dots, y_T)^\top$ denote the observed outcome series. The evaluation subvector is $\mathbf{y}_{\text{eval}} = \mathbf{y}_{\mathcal{T}_{\text{eval}}} \in \mathbb{R}^m$, and the post-treatment outcome vector is $\mathbf{y}_{\text{post}} = \mathbf{y}_{\mathcal{T}_2} \in \mathbb{R}^n$.

For each $t \in \mathcal{T}$, let y_t^1 and y_t^0 denote the potential outcomes under treatment and no treatment, respectively. The observed outcome is

$$y_t = d_t y_t^1 + (1 - d_t) y_t^0, \quad d_t = \mathbf{1}\{t > T_0\}.$$

For $t > T_0$, only y_t^1 is observed. SHC provides an estimate $\hat{y}_t^0$ for the missing counterfactual, and the treatment effect is estimated as $\hat{\alpha}_t = y_t - \hat{y}_t^0$, $t \in \mathcal{T}_2$, with the average treatment effect on the treated (ATT) defined as $\widehat{\text{ATT}} = n^{-1} \sum_{t \in \mathcal{T}_2} \hat{\alpha}_t$. The SHC framework is grounded in the following assumptions.

Assumption 2.1 (Error Properties). *The error term satisfies $\mathbb{E}[\varepsilon_t] = 0$ and $\mathbb{E}[\varepsilon_t^2] < \infty$ for all $t \in \mathcal{T}$.*

Assumption 2.2 (Latent Smoothness and Matching). *The latent trend $\tilde{y}(t)$ is $(H + 1)$ times differentiable for some integer $H \geq 3$, with uniformly bounded derivative. Furthermore, there exists $\mathbf{w}^* \in \Delta^{N-1}$ such that $\mathbf{y}_{\text{eval}} = \tilde{\mathbf{Y}}_{\text{pre}} \mathbf{w}^*$, where Δ^{N-1} denotes the probability simplex.*

Assumption 2.3 (Feasibility and Stability). *There exists a sparse weight vector $\mathbf{w} \in \Delta^{N-1}$ such that $\mathbf{y}_{\text{eval}} \approx \tilde{\mathbf{Y}}_{\text{pre}} \mathbf{w}$ and the corresponding Gram matrix is positive definite.*

These assumptions ensure that noise is mean-zero and bounded, that the latent component is smooth enough to be locally approximated, and that the evaluation window can be matched stably from the donor pool of pre-treatment segments.

The implementation of SHC begins by smoothing the pre-treatment series $\mathbf{y}_{\text{pre}}$ using local linear regression with a Gaussian kernel, producing a smoothed trajectory $\tilde{\mathbf{y}}$. From this smoothed series, N overlapping segments of length m are extracted to form the donor matrix $\tilde{\mathbf{Y}}_{\text{pre}} \in \mathbb{R}^{m \times N}$. The evaluation vector $\mathbf{y}_{\text{eval}}$ is reconstructed as a convex combination of the donor segments by solving the quadratic program

$$\mathbf{w}^* = \arg \min_{\mathbf{w} \in \mathbb{R}^N} \left\| \mathbf{y}_{\text{eval}} - \tilde{\mathbf{Y}}_{\text{pre}} \mathbf{w} \right\|_2^2 + \varsigma \left\| \mathbf{C}_0^\top \mathbf{w} \right\|_2^2, \quad \text{subject to } \mathbf{w} \geq \mathbf{0}, \mathbf{1}^\top \mathbf{w} = 1,$$

where $\mathbf{C}_0$ contains the eigenvectors of the donor Gram matrix associated with eigenvalues below a threshold τ . The additional penalty shrinks weights along low-variance directions, mitigating instability from collinearity. Applying $\mathbf{w}^*$ to the aligned post-treatment donor matrix yields $\hat{\mathbf{y}}_{\text{post}}^0$, which serves as the SHC estimate of the untreated counterfactual trajectory.

2.2 The Augmented Synthetic Historical Control Method

The augmented synthetic historical control (ASHC) method extends SHC by permitting limited extrapolation beyond the convex hull of historical segments, while stabilizing the solution through a proximity penalty. This modification is motivated by cases where the simplex constraint of SHC is too restrictive to achieve adequate pre-treatment fit. By allowing negative weights under an affine constraint and penalizing deviations from the SHC solution, ASHC balances improved fit with regularization.

Let $\mathbf{w}^{\text{SHC}}$ denote the SHC weight vector. The ASHC estimator solves

$$\mathbf{w}^{\text{ASHC}} = \arg \min_{\mathbf{w} \in \mathbb{R}^N} \left\| \mathbf{y}_{\text{eval}} - \tilde{\mathbf{Y}}_{\text{pre}} \mathbf{w} \right\|_2^2 + \varsigma \left\| \mathbf{C}_0^\top \mathbf{w} \right\|_2^2 + \lambda \left\| \mathbf{w} - \mathbf{w}^{\text{SHC}} \right\|_2^2, \quad \text{subject to } \mathbf{1}^\top \mathbf{w} = 1.$$

As we can see, the first two terms are essentially unchanged. We have a fit term to learn the treated unit's outcomes in the evaluation period and the stability penalty term from [Chen et al. \[2024\]](#). The third term in the ASHC objective,

$$\lambda \left\| \mathbf{w} - \mathbf{w}^{\text{SHC}} \right\|_2^2,$$

is a proximity penalty. This term discourages the ASHC solution from deviating too far from the original SHC weights. The penalty plays a dual role. Conceptually, it regularizes the affine extension of SHC by shrinking the solution back toward the purely convex combination used in SHC. This ensures that the additional flexibility provided by negative weights is only exercised when it leads to a substantial reduction in pre-treatment error. Mathematically, the penalty stabilizes the optimization problem by making the objective function strongly convex. The tuning parameter $\lambda \geq 0$ controls the degree of shrinkage: when λ is very large, the ASHC solution converges to $\mathbf{w}^{\text{SHC}}$, while when $\lambda = 0$ the penalty vanishes and the estimator reduces to an unpenalized affine extension of SHC. Thus, the proximity penalty creates a continuum between SHC and its affine relaxation, allowing researchers to balance interpretability and stability against the potential gains in fit from extrapolation.

The ASHC counterfactual is obtained as $\hat{\mathbf{y}}_{\text{post}}^{0, \text{ASHC}} = \tilde{\mathbf{Y}}_{\text{post}} \mathbf{w}^{\text{ASHC}}$, with treatment effects estimated analogously to SHC. In addition to the assumptions underlying SHC, ASHC requires bounds on the operator norms of the donor matrices.

2.3 Bandwidth and Regularization Parameter Selection

The performance of ASHC depends critically on two hyperparameters: the smoothing bandwidth h and the proximity penalty parameter λ . Both are chosen via cross-validation.

First I discuss the bandwidth for the kernel smoothing. Let $\mathbf{y}_{\text{pre}} = (y_1, \dots, y_{T_0})^\top$ denote the pre-treatment series.

For a candidate bandwidth $h > 0$, define the local linear smoother estimate $\hat{y}_i^{(-i)}(h)$ as the fitted intercept obtained by regressing $\{y_j : j \neq i\}$ on $\{j - i : j \neq i\}$ with kernel weights

$$K_h(j - i) = \exp\left(-\frac{1}{2} \left(\frac{j-i}{h}\right)^2\right),$$

normalized so that $\sum_{j \neq i} K_h(j - i) = 1$. Formally, for each $i \in \{1, \dots, T_0\}$,

$$\hat{\beta}(h)^{(-i)} = \arg \min_{\beta_0, \beta_1} \sum_{j \neq i} K_h(j - i) \left(y_j - \beta_0 - \beta_1(j - i)\right)^2,$$

and the leave-one-out prediction is

$$\hat{y}_i^{(-i)}(h) = \hat{\beta}_0(h)^{(-i)}.$$

The LOOCV error is

$$\text{CV}(h) = \frac{1}{T_0} \sum_{i=1}^{T_0} \left(y_i - \hat{y}_i^{(-i)}(h)\right)^2.$$

The optimal bandwidth is then chosen as

$$h^* = \arg \min_{h \in \mathcal{H}} \text{CV}(h),$$

where $\mathcal{H}$ is a prespecified candidate grid. This is part of the normal SHC method.

Next we discuss lambda, exclusive to the ASHC. Recall the ASHC objective

$$Q(\mathbf{w}; \lambda) = \left\| \mathbf{y}_{\text{train}} - \tilde{\mathbf{Y}}_{\text{train}} \mathbf{w} \right\|_2^2 + \varsigma \left\| \mathbf{C}_0^\top \mathbf{w} \right\|_2^2 + \lambda \left\| \mathbf{w} - \mathbf{w}^{\text{SHC}} \right\|_2^2, \quad \text{subject to } \mathbf{1}^\top \mathbf{w} = 1.$$

Let the evaluation window $\mathcal{T}_{\text{eval}}$ be split into a training set $\mathcal{T}_{\text{train}}$ and a validation set $\mathcal{T}_{\text{val}}$, with corresponding subvectors $\mathbf{y}_{\text{train}}, \mathbf{y}_{\text{val}}$ and donor matrices $\tilde{\mathbf{Y}}_{\text{train}}, \tilde{\mathbf{Y}}_{\text{val}}$.

For each $\lambda > 0$,

$$\hat{\mathbf{w}}(\lambda) = \arg \min_{\mathbf{w} \in \mathbb{R}^N} Q(\mathbf{w}; \lambda),$$

and the validation error is

$$\text{Val}(\lambda) = \frac{1}{|\mathcal{T}_{\text{val}}|} \left\| \mathbf{y}_{\text{val}} - \tilde{\mathbf{Y}}_{\text{val}} \hat{\mathbf{w}}(\lambda) \right\|_2^2.$$

The optimal penalty parameter is selected as

$$\lambda^* = \arg \min_{\lambda \in \Lambda} \text{Val}(\lambda),$$

where Λ is a candidate grid of λ values (e.g. logarithmically spaced).

Together, h^* and λ^* balance the dual bias–variance tradeoffs of ASHC: smoothing bias versus noise, and extrapolation bias versus instability.

3 Theoretical Analysis of ASHC

This section develops the finite-sample error bounds and asymptotic properties of the Augmented Synthetic Historical Control (ASHC) estimator. Complete proofs are deferred to Appendix~???. Here we focus on explaining what these results mean for applied researchers.

3.1 Weight deviation from SHC

Let $\Delta := \mathbf{w}^{ASHC} - \mathbf{w}^{SHC}$ denote the deviation of ASHC weights from the SHC solution.

Lemma 3.1 (Deviation from SHC). *Under Assumptions 2.1–3.4,*

$$\|\Delta\| \leq \frac{\gamma}{2\lambda + 2\varsigma\|\mathbf{C}_0\|^2} \epsilon_{SHC},$$

where $\epsilon_{SHC} := \|\tilde{\mathbf{y}} - \tilde{\mathbf{Y}}\mathbf{w}^{SHC}\|$ is the SHC in-sample fit error.

This result shows that ASHC weights (and therefore the predicted outcomes) depart from SHC only when SHC leaves meaningful residual error.

Using this Lemma, we provide an explicit bound on the error for ASHC with respect to the SHC method.

Proposition 3.1 (Pre-treatment Error Bound). *Under Assumptions 2.1–2.3,*

$$\|\tilde{\mathbf{y}} - \tilde{\mathbf{Y}}\mathbf{w}^{ASHC}\| \leq \epsilon_{SHC} \left(1 + \frac{\gamma}{2\lambda + 2\varsigma\|\mathbf{C}_0\|^2}\right) + b_\epsilon(H, k) + \delta.$$

The bound ties ASHC’s pre-treatment error to three ingredients: the SHC fit error ϵ_{SHC} , a smoothness bias term $b_\epsilon(H, k)$ from kernel smoothing, and the noise level δ of the outcome variable. In other words, ASHC cannot worsen the in sample fit, and will be at least as good as the SHC. When SHC’s fit leaves room for improvement, ASHC reduces error without letting variance get out of control. The ASHC method balances gains in fit against the risk of overfitting, using the cross validation method above.

We may also provide out of sample bounds in the prediction error for the ATT.

Proposition 3.2 (Out-of-sample Error Bound). *Under Assumptions 2.1–3.5,*

$$\|\mathbf{y}_{post} - \mathbf{Y}_{post}\mathbf{w}^{ASHC}\| \leq b_{\epsilon, post}(H, k) + \delta_{post} + \gamma_{post} \left(\eta + \frac{\epsilon_{SHC}}{2\lambda + 2\varsigma\|\mathbf{C}_0\|^2}\right),$$

where η bounds the difference $\|\mathbf{w}^* - \mathbf{w}^{SHC}\|$.

This inequality clarifies what governs the reliability of ASHC’s counterfactual predictions. Smoothness bias and noise matter, as does the quality of the SHC fit, measured by ϵ_{SHC} . The term η reflects how far SHC is from the infeasible “oracle” weights. The regularization parameter λ limits how much ASHC departs from SHC in seeking a better fit. For applied use, the upshot is that ASHC’s post-treatment predictions improve on SHC when SHC fits poorly, yet the scope of extrapolation is bounded so that counterfactuals do not swing wildly.

We must introduce two new assumptions for this:

Assumption 3.4 (Operator Norm Bound). *The pre-treatment donor matrix satisfies $\|\tilde{\mathbf{Y}}\| \leq \gamma$ for some constant $\gamma > 0$.*

Assumption 3.5 (Post-Treatment Operator Norm). *There exists $\gamma_{post} > 0$ such that $\|\tilde{\mathbf{Y}}_{post}\| \leq \gamma_{post}$.*

Assumption~3.4 ensures that the smoothed pre-treatment donor matrix does not admit arbitrarily large operator norm growth, which would destabilize the estimation of weights under affine relaxation. Assumption~3.5 provides a corresponding bound for the post-treatment donor matrix, guaranteeing that extrapolation via affine weights cannot yield counterfactual predictions with unbounded variance.

Together with the base SHC assumptions, these conditions provide stability and control of the ASHC estimator, enabling its theoretical properties to be derived and ensuring its performance in finite samples.

3.2 Asymptotic properties

I now establish the large sample properties of the ASHC. Below I show its consistency.

Theorem 3.1 (Consistency of ASHC). *Under Assumptions 2.1–3.5, as $m \rightarrow \infty$ and $H \rightarrow \infty$,*

$$\|\tilde{\mathbf{y}} - \tilde{\mathbf{Y}}w^{ASHC}\| \longrightarrow \delta.$$

If $\boldsymbol{\eta}$ is sub-Gaussian, then

$$\|\tilde{\mathbf{y}} - \tilde{\mathbf{Y}}w^{ASHC}\| \xrightarrow{P} 0.$$

Theorem 3.1 establishes that with a sufficiently long evaluation window and smooth latent trends, the ASHC prediction error converges to the noise level, and under standard probabilistic assumptions it converges in probability to zero. For empirical work, this means that with enough data, ASHC recovers the true counterfactual trajectory and ensures consistency of the estimated treatment effect.

To summarize, these results show that ASHC remains close to SHC when SHC already performs well, improves fit when SHC leaves residual error, provides bounded error guarantees in finite samples, and delivers consistent estimates asymptotically. For applied researchers, this means ASHC can be adopted as a robust extension of SHC, especially in settings where SHC struggles to achieve a satisfactory pre-treatment fit.

4 Empirical Application

We replicate the setting of Bel et al. [2021], who study the effects of Barcelona’s 2015 hotel moratorium using data on hotel prices from 2010 to 2017. This application provides a compelling case to benchmark ITS, SCM, SHC, and ASHC against one another.

5 Conclusion

References

- Alberto Abadie, Alexis Diamond, and Jens Hainmueller. Synthetic control methods for comparative case studies: Estimating the effect of california’s tobacco control program. *Journal of the American Statistical Association*, 105(490):493–505, 2010. URL <https://doi.org/10.1198/jasa.2009.ap08746>.
- Alberto Abadie, Alexis Diamond, and Jens Hainmueller. Comparative politics and the synthetic control method. *American Journal of Political Science*, 59(2):495–510, 2015. doi: <https://doi.org/10.1111/ajps.12116>.
- Andrew Baker, Brantly Callaway, Scott Cunningham, Andrew Goodman-Bacon, and Pedro H. C. Sant’Anna. Difference-in-differences designs: A practitioner’s guide, 2025. URL <https://arxiv.org/abs/2503.13323>.
- Germà Bel, Stephan Joseph, and Ferran A. Mazaira-Font. The effects of moratoriums on hotel building: An anti-tourism measure, or rather protection for local incumbents? *Applied Economics Letters*, 29(14): 1319–1324, 2021. doi: 10.1080/13504851.2021.1927958.
- Eli Ben-Michael, Avi Feller, and Jesse Rothstein. The augmented synthetic control method. *Journal of the American Statistical Association*, 116(536):1789–1803, dec 2021. doi: 10.1080/01621459.2021.1929245.
- Yi-Ting Chen, Jui-Chung Yang, and Tzu-Ting Yang. Synthetic historical control for policy evaluation. Available at SSRN: <https://ssrn.com/abstract=4995085>, September 2024. URL <http://dx.doi.org/10.2139/ssrn.4995085>.
- Jared Coopersmith, Thomas D. Cook, Jelena Zurovac, Duncan Chaplin, and Lauren V. Forrow. Internal and external validity of the comparative interrupted time-series design: A meta-analysis. *Journal of Policy Analysis and Management*, 41(1):252–277, 2022. doi: <https://doi.org/10.1002/pam.22361>. URL <https://onlinelibrary.wiley.com/doi/abs/10.1002/pam.22361>.